

# Towards Implicit Parallel Programming for Systems

Sebastian Ertel

Dr.-Ing. Defense  
Chair for Compiler Construction  
TU Dresden, Germany

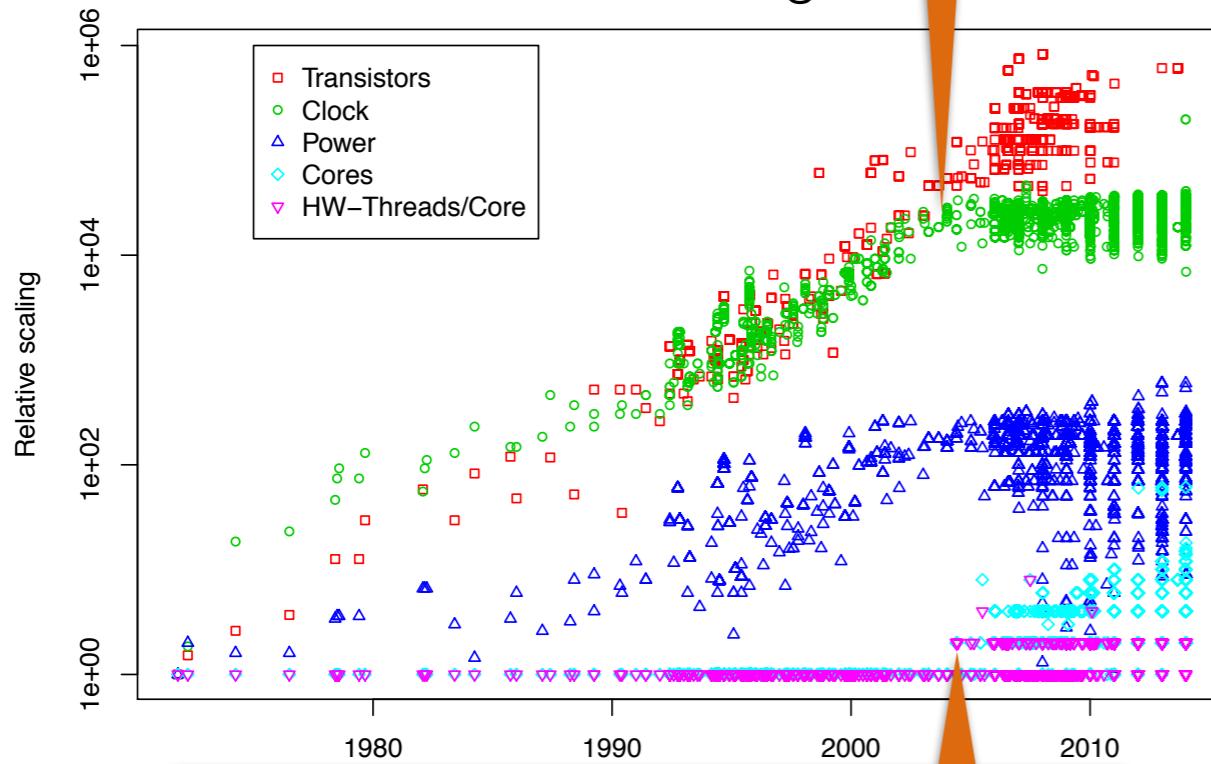
Supervisor: Prof. Dr.-Ing. Jeronimo Castrillon  
Fachreferent : Prof. Dr.-Ing. Hermann Härtig

4.9.2019

# The Shape Of Things To Run

The end of the single-core era

Processor scaling trends



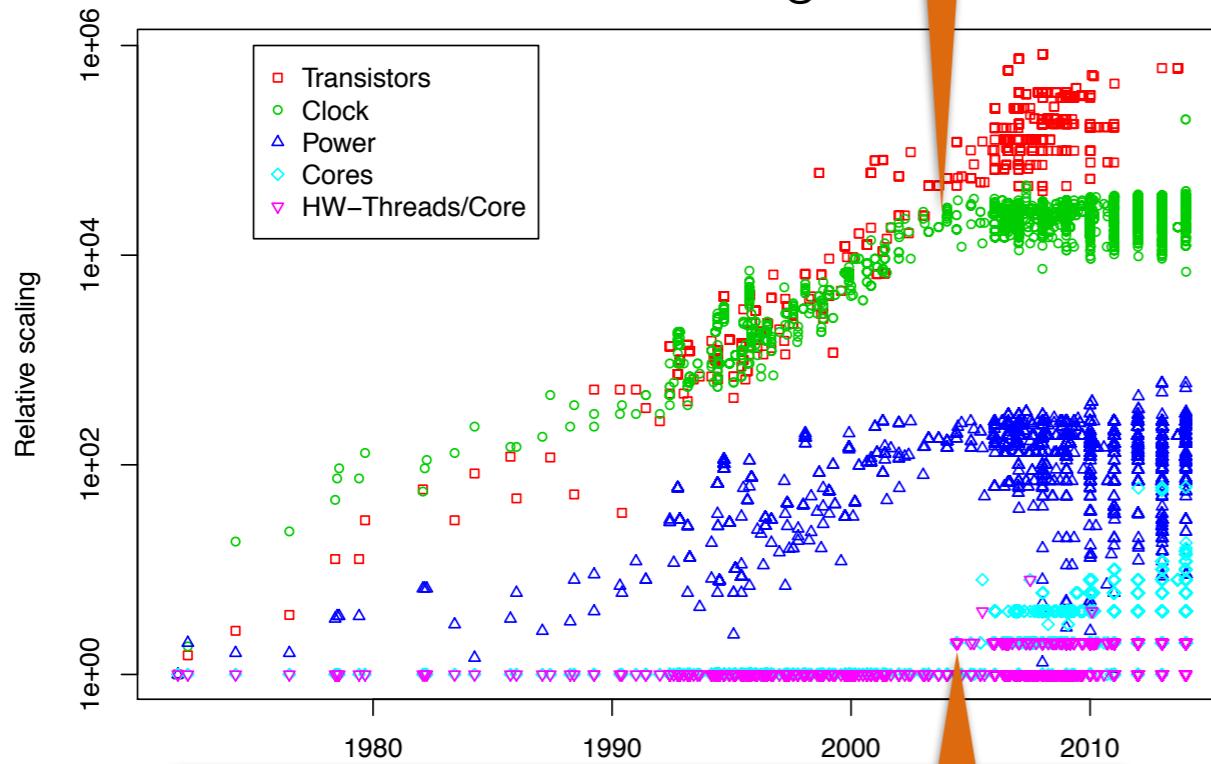
The dawn of the multi-core era

Herb Sutter. 2005. The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software. Dr. Dobbs Journal.

# The Shape Of Things To Run

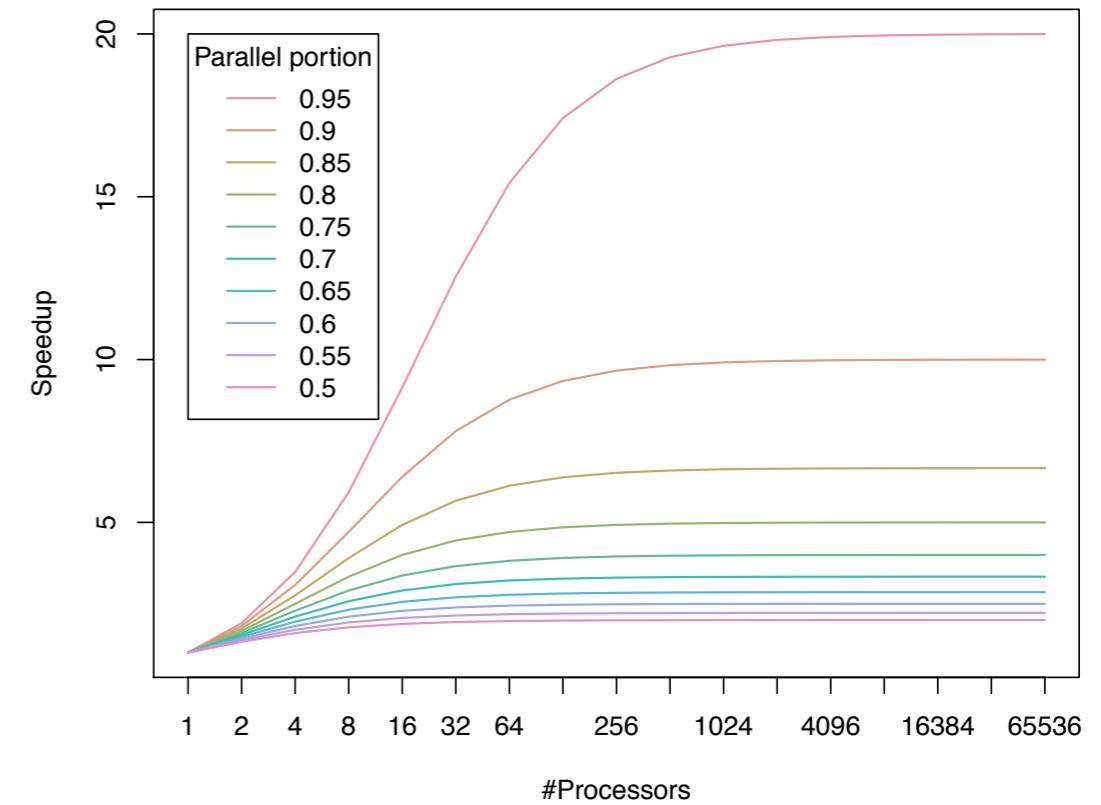
The end of the single-core era

Processor scaling trends



The dawn of the multi-core era

Amdahl's Law



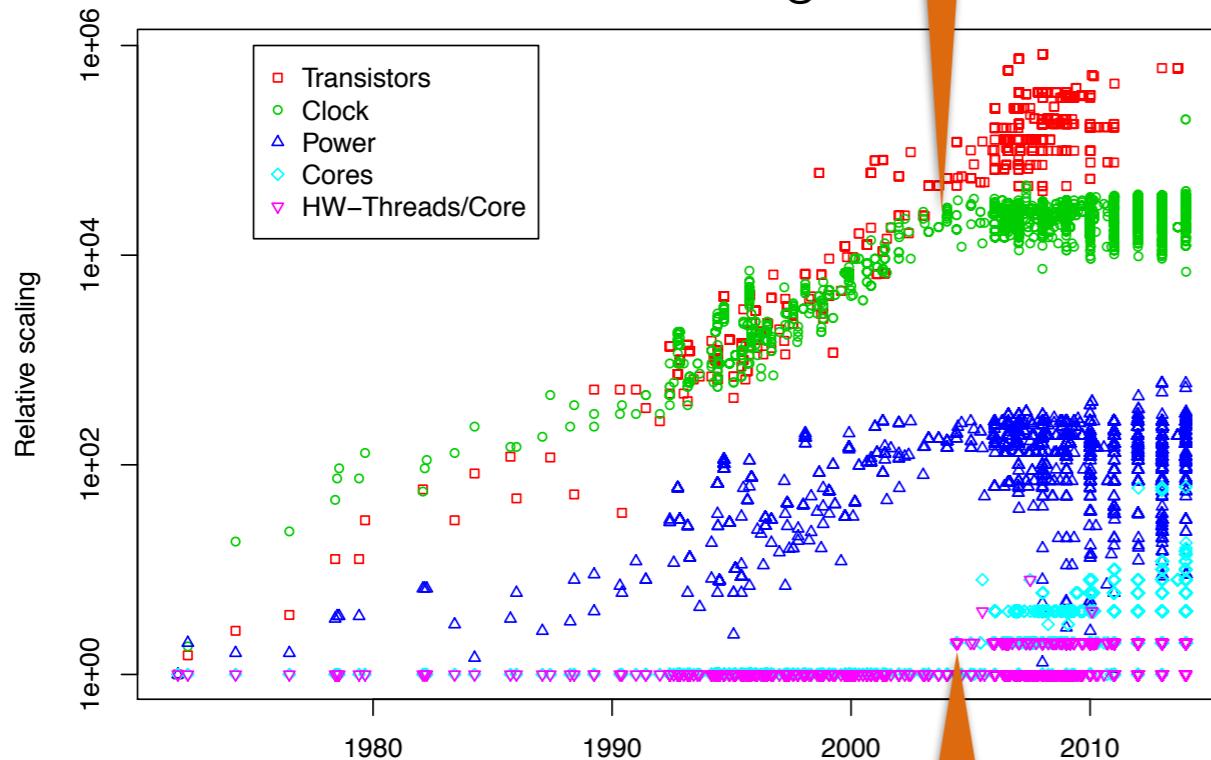
Herb Sutter. 2005. The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software. Dr. Dobbs Journal.

G. M. Amdahl. 1967. Validity of the single processor approach to achieving large scale computing capabilities. AFIPS '67. ACM.

# The Shape Of Things To Run

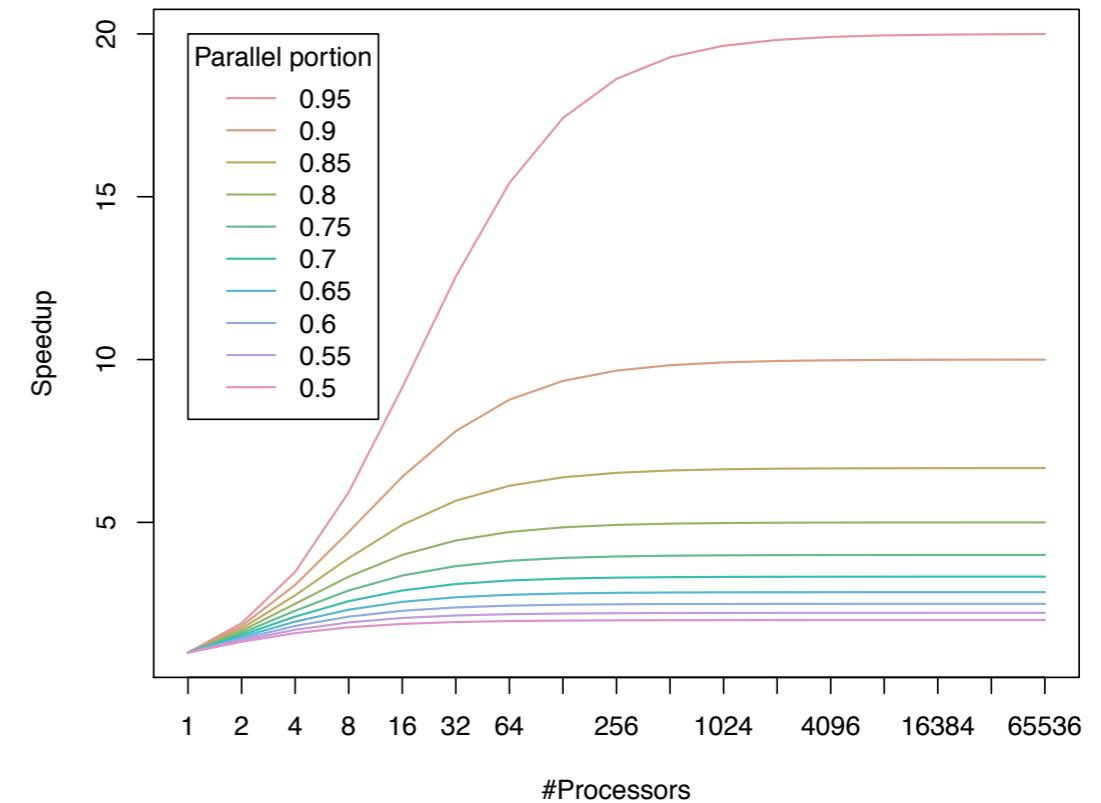
The end of the single-core era

Processor scaling trends



The dawn of the multi-core era

Amdahl's Law



Challenge	Near-Term	Long-Term
1,000-fold software parallelism	Data parallel languages and "mapping" of operators, library and tool-based approaches	New high-level languages, compositional and deterministic frameworks

Herb Sutter. 2005. The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software. Dr. Dobbs Journal.

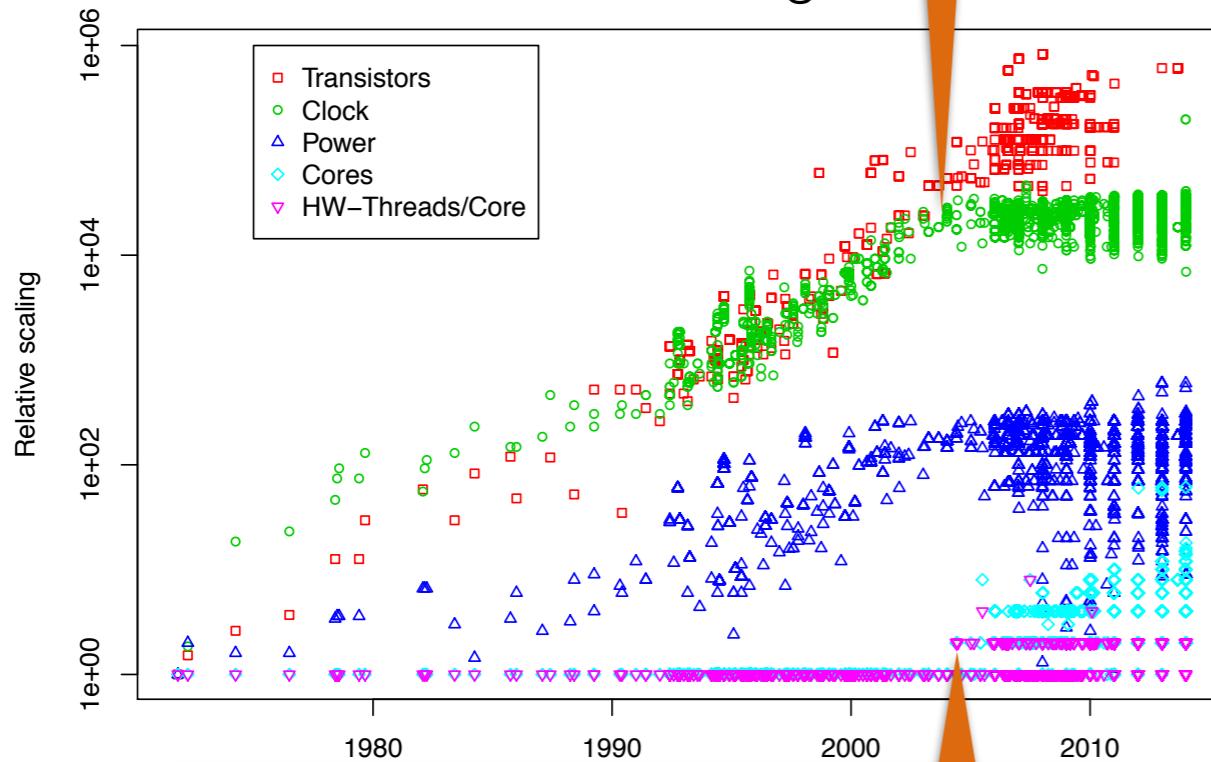
G. M. Amdahl. 1967. Validity of the single processor approach to achieving large scale computing capabilities. AFIPS '67. ACM.

Shekhar Borkar and Andrew A. Chien. 2011. The future of microprocessors. Commun. ACM.

# The Shape Of Things To Run

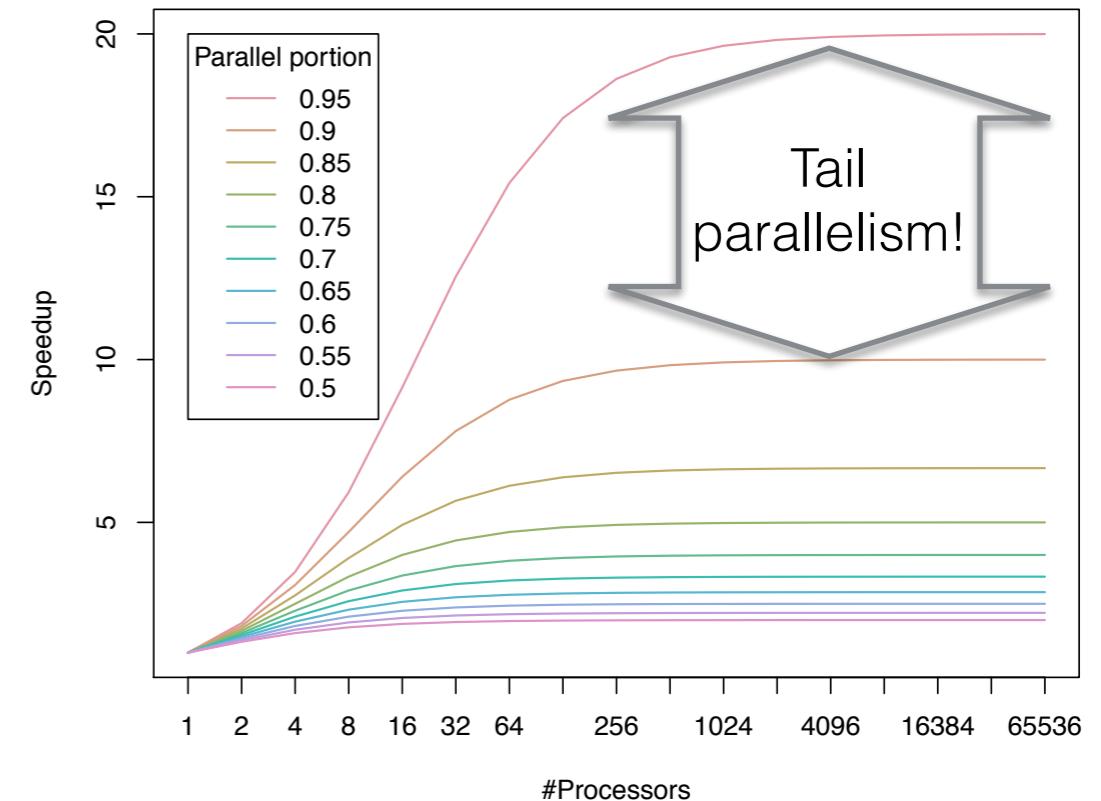
The end of the single-core era

Processor scaling trends



The dawn of the multi-core era

Amdahl's Law



Challenge	Near-Term	Long-Term
1,000-fold software parallelism	Data parallel languages and "mapping" of operators, library and tool-based approaches	New high-level languages, compositional and deterministic frameworks

Herb Sutter. 2005. The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software. Dr. Dobbs Journal.

G. M. Amdahl. 1967. Validity of the single processor approach to achieving large scale computing capabilities. AFIPS '67. ACM.

Shekhar Borkar and Andrew A. Chien. 2011. The future of microprocessors. Commun. ACM.

# Programming Is Simple!

Programming

Variables  
Functions  
*Function Calls*

Conditionals  
Loops

Classes  
Objects  
Methods

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compare-  
and-swap  
Threads

Operating System

LVars  
MVars  
IVars

Fork/Join

Tasks

Work-Stealing

Coroutines

Spawn

Futures

Events

Promises

Actors

Fibers

Events

OpenMP (loops)

Map/Reduce

Custom/optimized

Parallel  
Haskell (pH)  
Reactive  
Programming  
Streams  
*Dataflow*

SQL

High level

Low level

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compa  
and-sw

LVars  
MVars

Threads

Tasks

Futures

Events

Algorithmic  
DSLs

Map/Reduce

SQL

Parallel  
Haskell (pH)  
Reactive  
Programming

Streams

Dataflow

Composition and state are key for systems!

Work-Stealing

Custom/optimized

Low level

High level

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compare-  
and-swap  
Threads

Operating System

MVCC  
Fork/Join  
Tasks

LVars  
IVars  
Spaces  
Futures  
Events

State  
Composition

Procedures  
Fibers

Events

Algorithmic  
Skeletons  
OpenMP (loops)  
Map/Reduce

Custom/optimized

SQL  
Parallel  
Haskell (pH)  
Reactive  
Programming  
Streams  
*Dataflow*

Low level

High level

Lee EA. The problem with threads. Computer. 2006.

Van Renesse R. Goal-oriented programming, or composition using events, or threads considered harmful. In ACM SIGOPS European Workshop 1998.

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compare-  
and-swap  
Threads

Operating System

LVars

MV

IVa

Fork/Join

Tasks

Work

✓ State  
✓ Composition

outines

ibers

Algorithmic  
Skeletons

OpenMP (loops)

Map/Reduce

Custom/optimized

SQL

Parallel  
Haskell (pH)

Reactive  
Programming

Streams

Dataflow

Low level

High level

Herlihy M, Moss JE. Transactional memory: Architectural support for lock-free data structures. ACM; 1993 Jun 1.

Shavit N, Touitou D. Software transactional memory. Distributed Computing. 1997 Feb 1;10(2):99-116.

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes  
Threads  
Operating System

LVars  
MVars  
IVars  
Fork/Join  
Threads  
Tasks  
Work Queue

Routines  
Fibers  
Custom/optimized

Algorithmic  
Skeletons  
OpenMP (loops)  
Map/Reduce  
*Dataflow*

SQL  
Parallel  
Haskell (pH)  
Reactive  
Programming  
Streams  
*Dataflow*

Low level

High level

✓ State  
✓ Composition  
✗ Composition

J. Swalens, S. Marr, J. De Koster, and T. Van Cutsem. 2014. Towards Composable Concurrency Abstractions. In Proceedings of PLACES'14.

J. Swalens, J. De Koster, and W. De Meuter. 2016. Transactional Tasks: Parallelism in Software Transactions. In Proceedings of ECOOP'16.

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compare-  
and-swap  
Threads

Operating System

Fork/Join

LVars  
MVars  
IVars

Work-Stealing

Spa  
Tasks

? State  
? Composition

oroutines  
Fibers  
ts

Algorithmic  
Skeletons  
OpenMP (loops)  
Map/Reduce

Custom/optimized

SQL

Parallel  
Haskell (pH)  
Reactive  
Programming

Streams  
*Dataflow*

Low level

High level

Mou ZG, Hudak P. An algebraic model for divide-and-conquer and its parallelism. *The Journal of Supercomputing*. 1988.

Lämmel R. Google's MapReduce programming model—Revisited. *Science of computer programming*. 2008.

Cole MI. Algorithmic skeletons: structured management of parallel computation. London: Pitman; 1989.

Buono D, Danelutto M, Lametti S. Map, reduce and mapreduce, the skeleton way. *Procedia computer science*. 2010

Dagum L, Menon R. OpenMP: An industry-standard API for shared-memory programming. *Computing in Science & Engineering*. 1998.

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes  
Threads  
Operating System

compare-  
and-swap  
Fork/Join  
Tasks

MV  
IVar  
Space  
Work

State  
Composition  
Systems

Algorithmic  
Skeletons  
OpenMP (loops)  
Map/Reduce  
Custom/optimized

SQL  
Parallel  
Haskell (pH)  
Reactive  
Programming  
Streams  
**Dataflow**

Low level

High level

Chamberlin DD, Boyce RF. SEQUEL: A structured English query language. In Proceedings of the 1974 ACM SIGFIDET (now SIGMOD) workshop on Data description, access and control. ACM.

Nikhil RS. Implicit parallel programming in pH. Morgan Kaufmann; 2001.

Chambers C, Raniwala A, Perry F, Adams S, Henry RR, Bradshaw R, Weizenbaum N. FlumeJava: easy, efficient data-parallel pipelines. In ACM Sigplan Notices 2010. ACM.

Coutts D, Leshchinskiy R, Stewart D. Stream fusion: From lists to streams to nothing at all. In ACM SIGPLAN Notices 2007. ACM.

Thies W, Karczmarek M, Amarasinghe S. StreamIt: A language for streaming applications. In International Conference on Compiler Construction 2002. Springer.

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

LVars  
MVars

Composition and state are key for systems!

Algorithmic  
Composition

compar  
and-sw

Tasks  
Futures  
Events

Map/Reduce

Operating System

Work-Stealing

Custom/optimized

Low level

High level

Explicit

Implicit

Concurrency Bugs

Optimizations  
(Sequential  $\Rightarrow$  Parallel)

# Parallel Programming Is Hard!

Programming

Variables  
Functions  
Function Calls

Conditionals  
Loops

Classes  
Objects  
Methods



Concurrent/Parallel Programming

Transactional  
Memory  
Locks  
Message-  
Passing  
Processes

compa  
and-sw

Threads

LVars  
MVars

Tasks  
Futures

Events

Algorithmic  
Composition

Map/Reduce

SQL  
Parallel  
Haskell (pH)  
Reactive  
Programming  
Streams  
*Dataflow*

Work-Stealing

Custom/optimized

Low level

High level

Explicit

Implicit

Concurrency Bugs

PL/compiler co-design!

Optimizations  
(Sequential  $\Rightarrow$  Parallel)

# Ingredients for Implicit Parallel Programming

## Programming Model/ Language

- No concurrency abstractions.
- Practicality:
  - Integration into existing programming models (e.g., OOP)/languages.
  - Gradual switch/reuse of existing code base.

# Ingredients for Implicit Parallel Programming

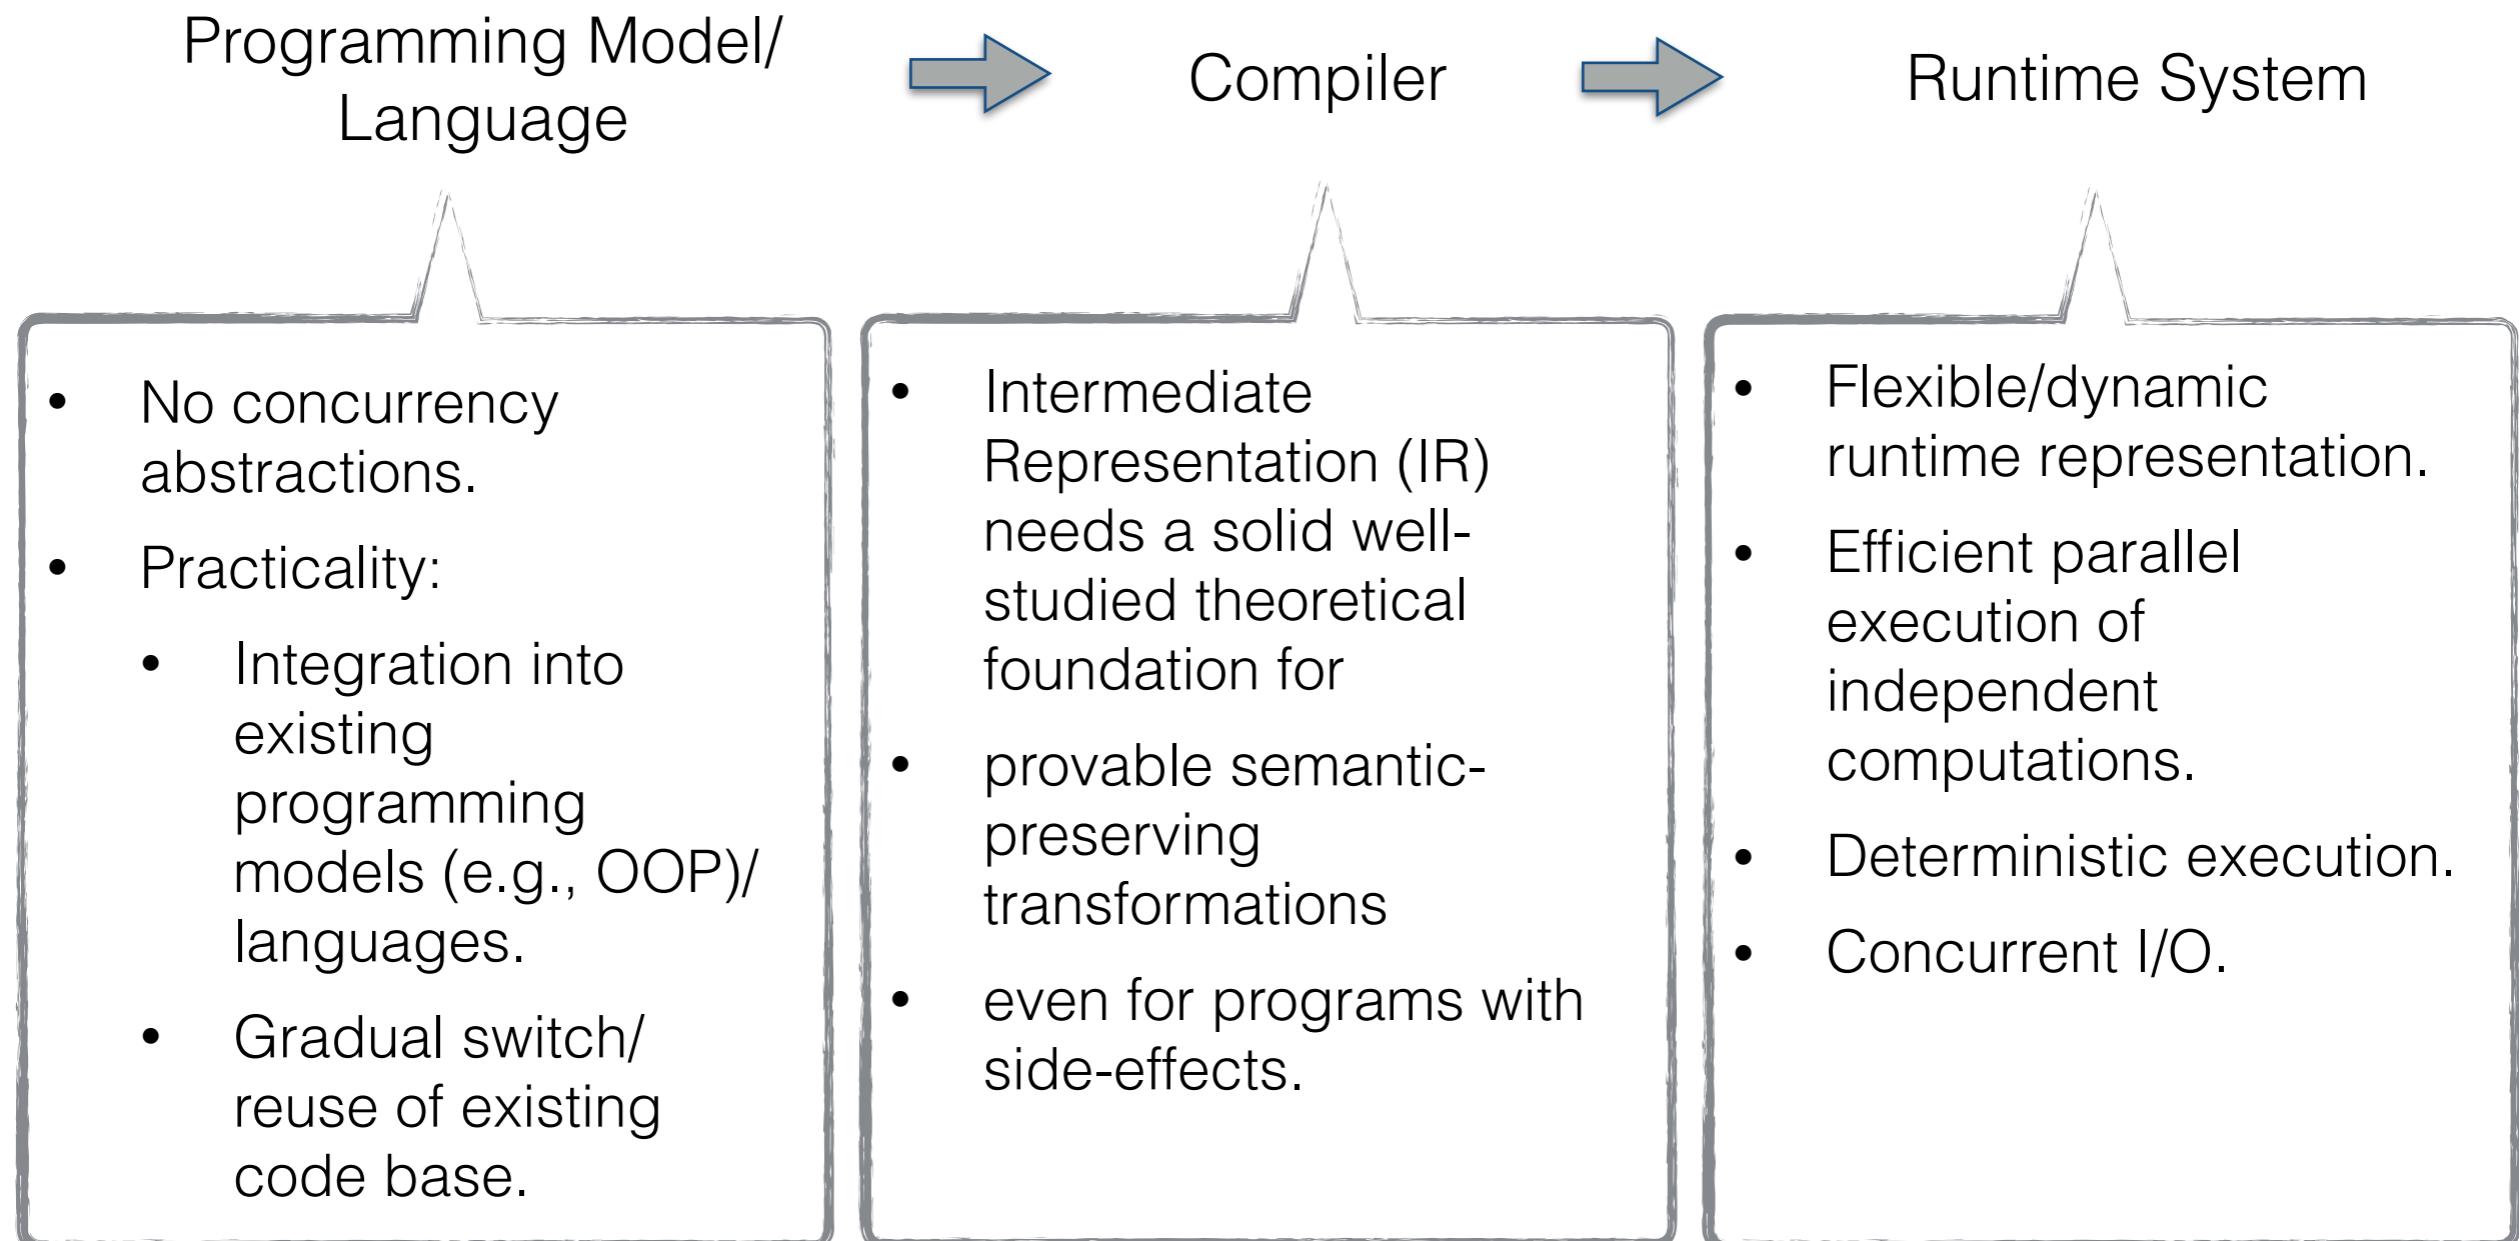
Programming Model/  
Language



Compiler

- No concurrency abstractions.
- Practicality:
  - Integration into existing programming models (e.g., OOP)/languages.
  - Gradual switch/reuse of existing code base.
- Intermediate Representation (IR) needs a solid well-studied theoretical foundation for
  - provable semantic-preserving transformations
  - even for programs with side-effects.

# Ingredients for Implicit Parallel Programming



Programming Model/  
Language



Compiler



Runtime System

Programming Model/  
Language



Compiler



Runtime System

Programming Model/  
Language

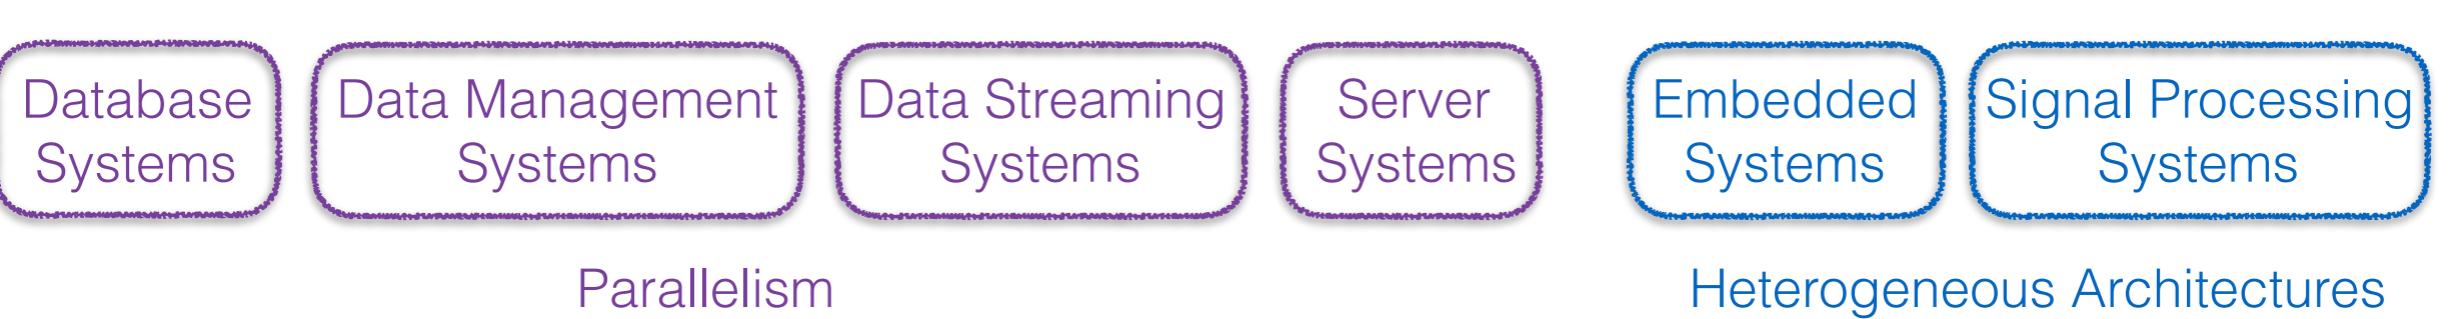


Compiler

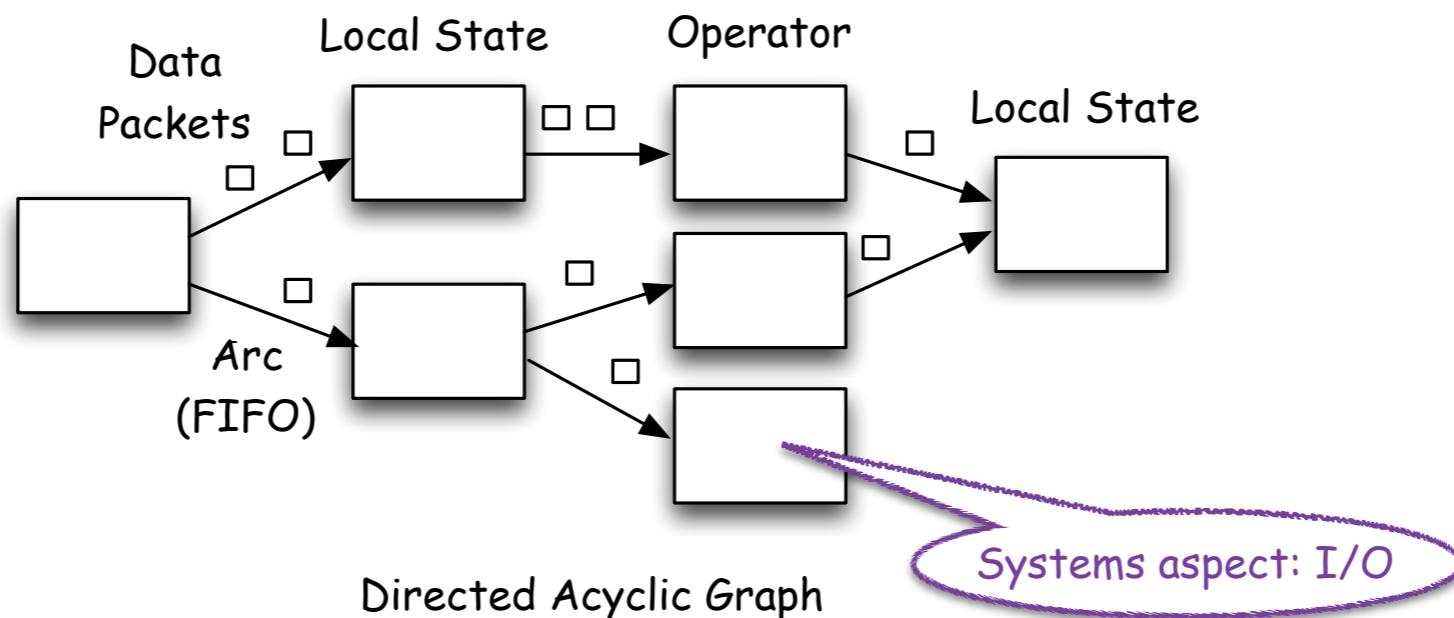


Runtime System

- Parallelism
- Mapping/Scheduling
-



# Dataflow



# Database Systems

# Data Management Systems

# Data Streaming Systems

# Server Systems

# Embedded Systems

# Signal Processing Systems

## Parallelism

## Heterogeneous Architectures

J. B. Dennis. Data flow supercomputers. Computer, 13(11):48–56, Nov. 1980.

J. P. Morrison. Flow-Based Programming. Nostrand Reinhold, 1994

Aryind and D. E. Culler. Annual review of computer science vol. I, 1986. chapter Dataflow architectures. Annual Reviews Inc.

Programming Model/  
Language



Compiler



Runtime System

- Parallelism
- Mapping/Scheduling
-

Programming Model/  
Language



Compiler



Runtime System

### Dataflow

- Parallelism
- Mapping/Scheduling
- Dynamic Software Evolution

Sebastian Ertel and Pascal Felber. 2014. A framework for the dynamic evolution of highly-available dataflow programs. In Proceedings of the 15th International Middleware Conference (Middleware '14). ACM.

Sebastian Ertel, Christof Fetzer, and Pascal Felber. 2015. Ohua: Implicit Dataflow Programming for Concurrent Systems. In Proceedings of the Principles and Practices of Programming on The Java Platform (PPPJ '15). ACM.

Programming Model/  
Language



Compiler



Runtime System

### Dataflow

- Parallelism
- Mapping/Scheduling
- Dynamic Software Evolution

Programming Model/  
Language



Compiler



Runtime System  
**Dataflow**

- Flexible/dynamic runtime representation.
- Efficient parallel execution of independent computations.
- Deterministic execution.
- Concurrent I/O.

Programming Model/  
Language



Compiler



Runtime System  
**Dataflow**

**Dataflow:**

- Not concise
- Limited scalability
- Concurrency abstraction

Programming Model/  
Language



Compiler

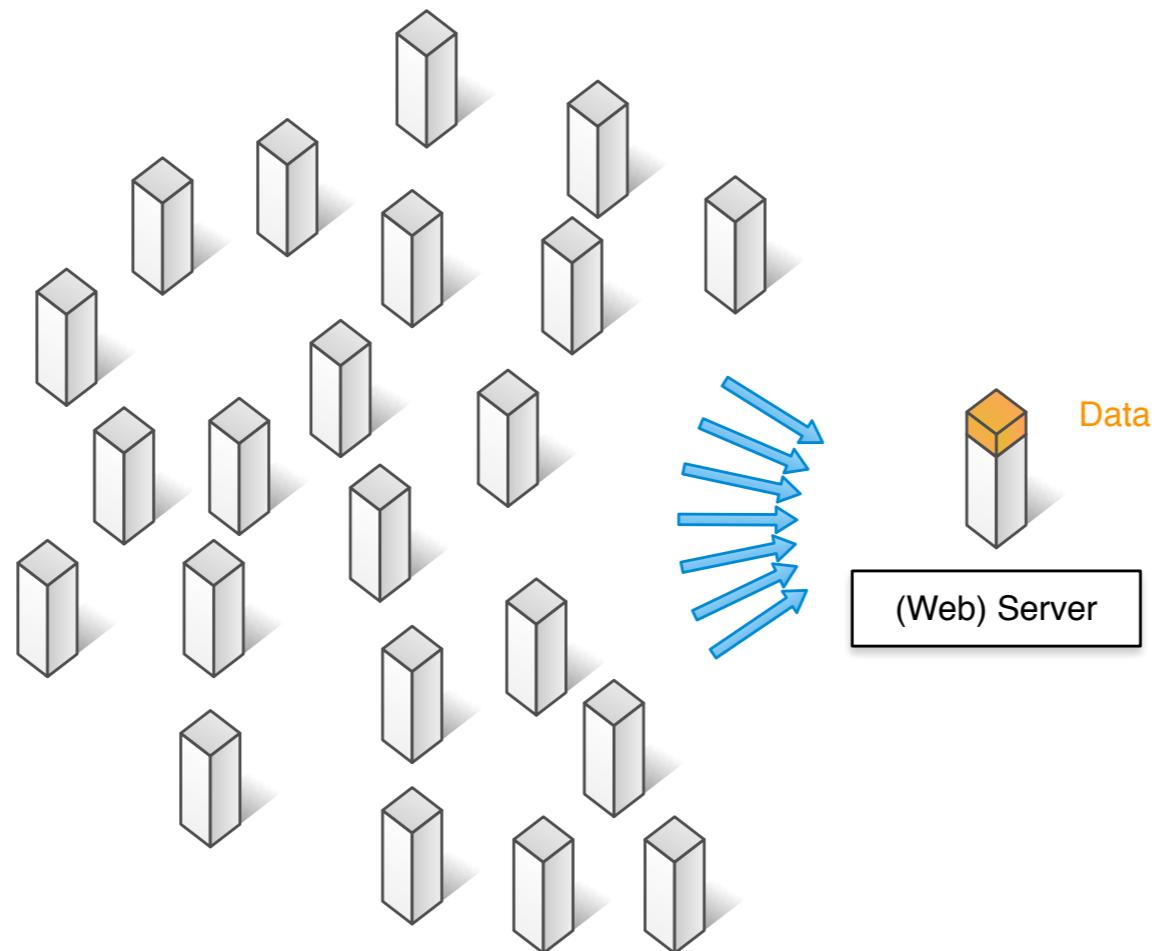


Runtime System  
**Dataflow**

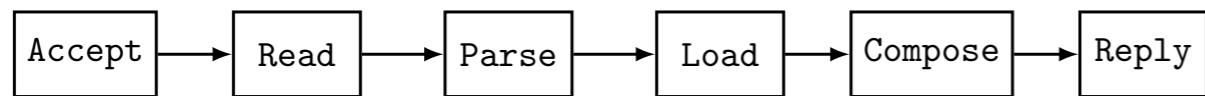
SFP:

- Variables
- Functions
- State

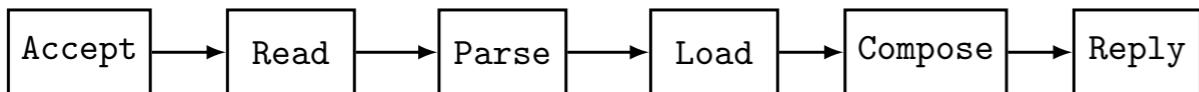
# Server Architecture



# Explicit Dataflow Construction



# Explicit Dataflow Construction

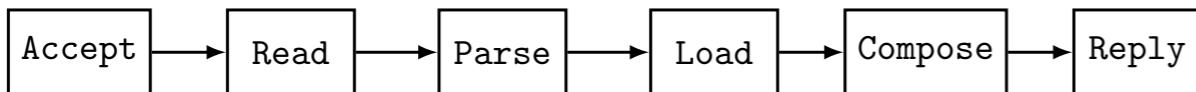


```
public class WebServer extends StreamFlexGraph {
    private Filter a, r, p, l, c, rep;

    public WebServer() {
        a = makeFilter(Accept.class);
        r = makeFilter(Read.class);
        p = makeFilter(Parse.class);
        l = makeFilter(Compose.class);
        rep = makeFilter(Reply.class);
        connect(a, r);
        connect(r, p);
        connect(p, l);
        connect(l, c);
        connect(c, rep);
        validate();
    }

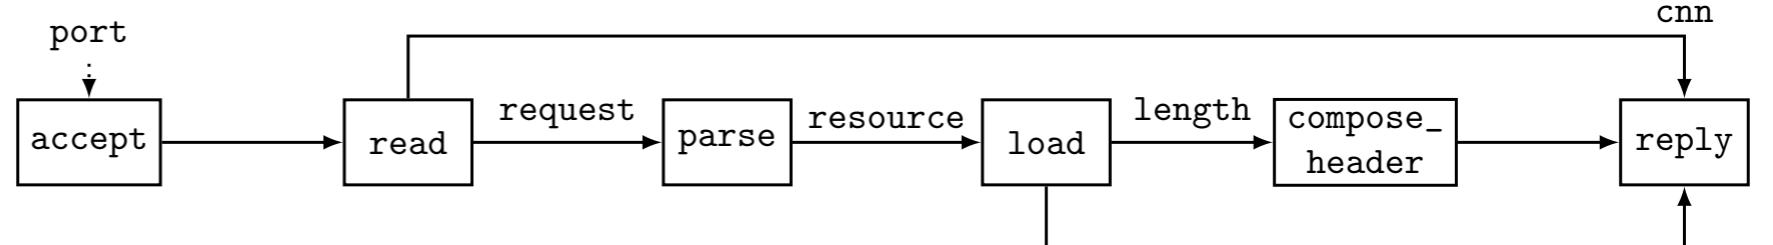
    public void start() {
        new Synthesizer(a).start();
        super.start();
    }
}
```

# Implicit Dataflow Construction



```
public class WebServer extends StreamFlexGraph {  
    private Filter a, r, p, l, c, rep;  
  
    public WebServer() {  
        a = makeFilter(Accept.class);  
        r = makeFilter(Read.class);  
        p = makeFilter(Parse.class);  
        l = makeFilter(Compose.class);  
        rep = makeFilter(Reply.class);  
        connect(a, r);  
        connect(r, p);  
        connect(p, l);  
        connect(l, c);  
        connect(c, rep);  
        validate();  
    }  
  
    public void start() {  
        new Synthesizer(a).start();  
        super.start();  
    }  
}
```

```
fn server(port) {  
    let (cnn, request) = read(accept(port));  
    let resource = parse(request);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```



# Classic Operators

```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
}
```



Explicit dataflow

# Classic Operators

```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



Explicit dataflow



Computation

# Stateful Functions

```
class FileLoad extends Filter {  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



The diagram illustrates the state flow of a stateful function. It starts with a large upward-pointing arrow labeled "Computation" pointing from the code block to the top of the class definition. Below the class definition is another upward-pointing arrow, indicating the transition from computation to storage or memory.

# Stateful Functions

```
class FileLoad extends Filter {  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}  
  
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}
```



The diagram illustrates a computation process. On the left, there is a block of Java code representing stateful computation. An arrow points from this code to the right, where the word "Computation" is written above a large upward-pointing arrow. This visual metaphor suggests that the computation described in the code results in a value or state that can be retrieved.

Computation

Computation is (almost) a function:

- ✗ Some loss of dataflow expressiveness.
- ✓ Solid theoretical foundation.
- ✓ Deterministic.

# Stateful Functions

```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}
```

# Stateful Functions

```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}  
  
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}
```



# Stateful Functions

```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}  
  
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}  
  
char* load(char* resource) {  
    static GHashTable* cache = g_hash_table_new();  
    char* contents = NULL;  
    // load file data from disk or cache (omitted)  
    return contents;  
}
```



# Stateful Functions



```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}  
  
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}  
  
char* load(char* resource) {  
    static GHashTable* cache = g_hash_table_new();  
    char* contents = NULL;  
    // load file data from disk or cache (omitted)  
    return contents;  
}
```



# Stateful Functions



```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```

State encapsulation



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}  
  
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}  
  
char* load(char* resource) {  
    static GHashTable* cache = g_hash_table_new();  
    char* contents = NULL;  
    // load file data from disk or cache (omitted)  
    return contents;  
}
```



# Local State

```
fn server(port) {  
    let (cnn, req)      = read(accept(port));  
    let resource        = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose(length), content)  
}
```



```
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}
```

State is local to each call-site!

Programming Model/  
Language



Compiler



Runtime System  
**Dataflow**

SFP:

- Variables
- Stateful Functions
- Composition

Programming Model/  
Language



Compiler



Dataflow  
Runtime System

SFP:

- Variables
- Stateful Functions
- Composition

**SFP: Stateful Functional Programming**

Programming Model/  
Language



Compiler



Dataflow  
Runtime System

SFP:

- Variables
- Stateful Functions
- Composition
- Code reuse!

SFP: Stateful Functional Programming

# Algorithms

```
fn server(port) {
    let (cnn, req)      = read(accept(port));
    let resource        = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

## Top-level Algorithms

### Abstraction

```
fn respond(cnn, content, length){
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let (cnn, req)      = read(accept(port));
    let resource        = parse(req);
    let (content, length) = load(resource);
    respond(cnn, content, length)
}
```

## Lambda Algorithms

```
fn server(port) {
    let respond = |cnn, content, length|
        reply(cnn, compose_header(length), content);
    let (cnn, req)      = read(accept(port));
    let resource        = parse(req);
    let (content, length) = load(resource);
    respond(cnn, content, length)
}
```

### Application

Programming Model/  
Language



Compiler



Runtime System  
**Dataflow**

SFP:

- Variables
- Stateful Functions
- Composition
- Algorithms

Programming Model/  
Language



Compiler

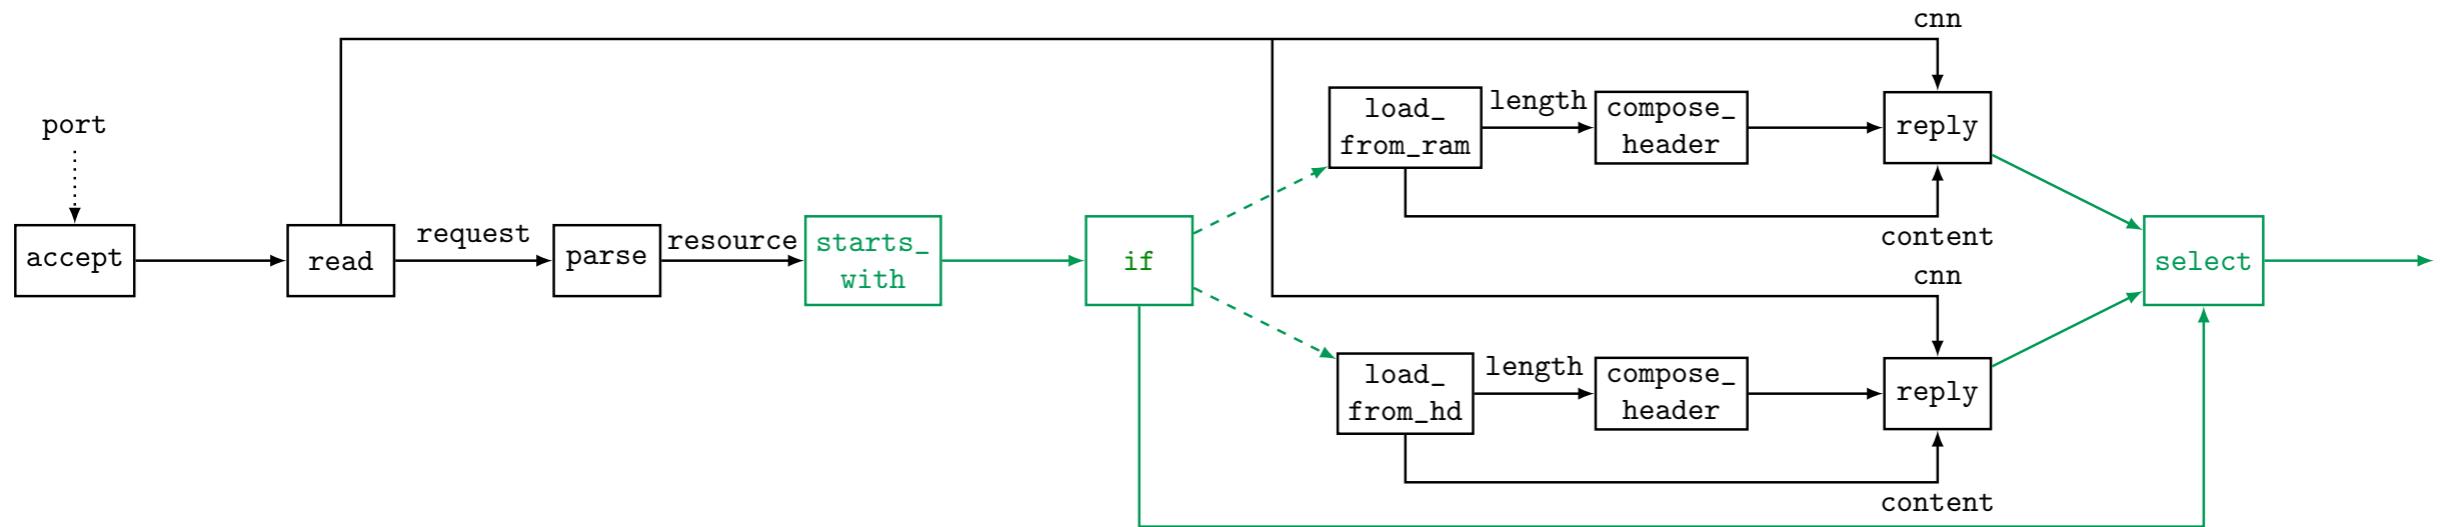


Runtime System  
**Dataflow**

SFP:

- Variables
- Stateful Functions
- Composition
- Algorithms
- Control Flow

# Control Flow



```

fn respond(cnn, content, length){
  reply(cnn, compose_header(length), content)
}
  
```

```

fn server(port) {
  let (cnn, req)      = read(accept(port));
  let resource        = parse(req);
  let (content, length) = load(resource);
  respond(cnn, content, length)
}
  
```

```

fn respond(cnn, content, length){
  reply(cnn, compose_header(length), content)
}
  
```

```

fn server(port) {
  let (cnn, req) = read(accept(port));
  let resource   = parse(req);
  if startsWith(resource, "news/") {
    let (content, length) = load_from_ram(resource);
    respond(cnn, content, length)
  } else {
    let (content, length) = load_from_hd(resource);
    reply(cnn, content, length)
  }
}
  
```

Programming Model/  
Language



Compiler



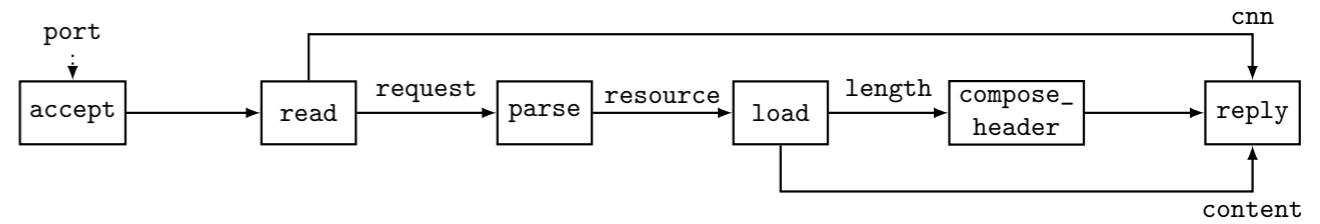
Runtime System  
**Dataflow**

SFP:

- Variables
- Stateful Functions
- Composition
- Algorithms
- Control Flow

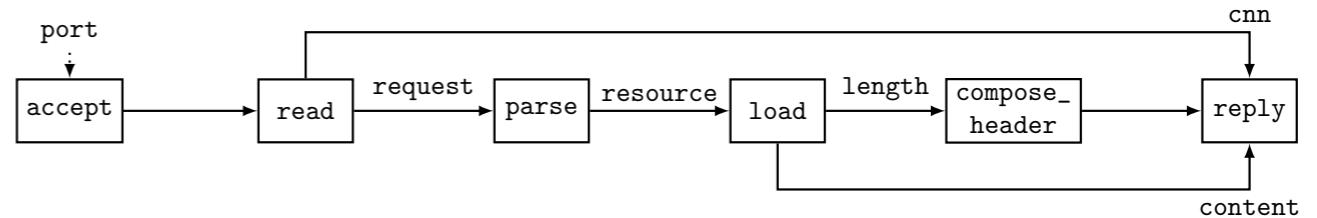
# Semantic Gap

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```



# Semantic Gap

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```



## Typical PL semantics:

The **server** function executes **once** for each invocation.

## Dataflow semantics:

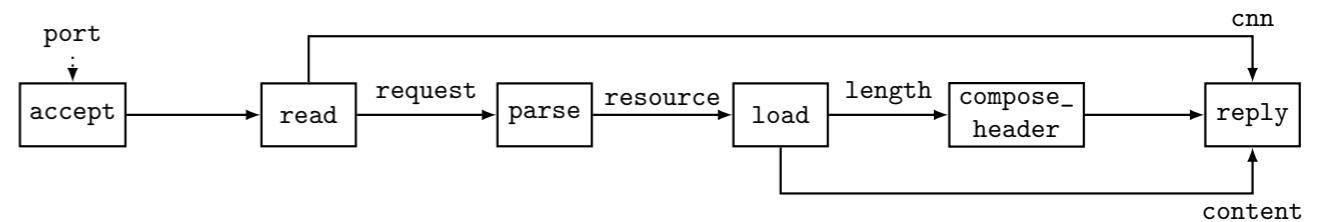
The **server** algorithm executes **forever** for each invocation.



The **accept** node is an I/O source operator, i.e., it finishes when all requests have been exhausted!

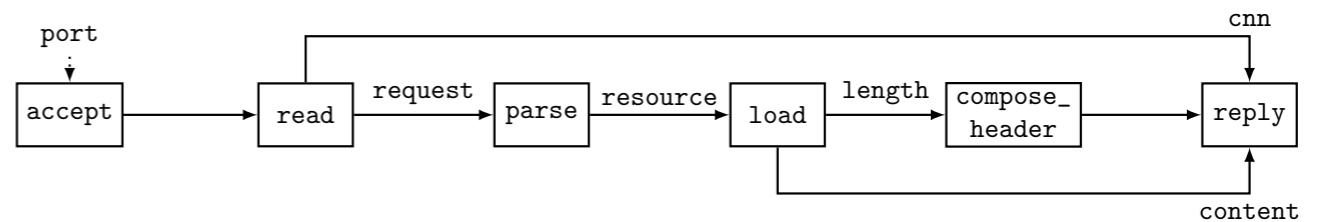
# SMap

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```



# SMap

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```

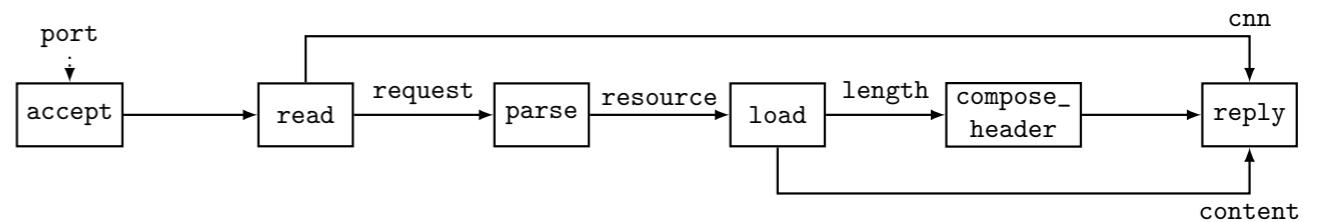


Abstract over processing:

```
fn request_handling(cnn) {  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}  
  
fn server(port) {  
    let cnn = accept(port);  
    request_handling(cnn)  
}
```

# SMap

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```

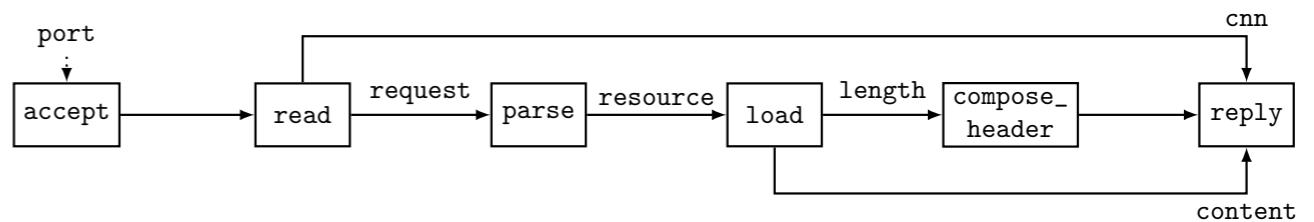


Abstract over processing:

```
fn request_handling(cnn) {  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}  
  
fn server(port) {  
    let cnn = accept(port);  
    request_handling(cnn) f  
}
```

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

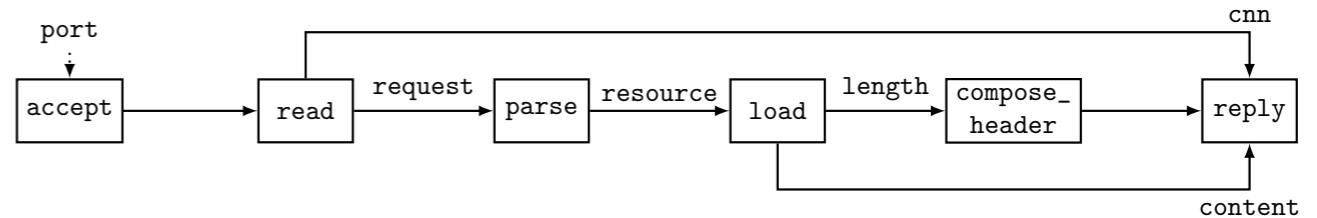
```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

$$I' = [b_1, b_2] = \text{map}(f, I)$$

$$\begin{aligned} b_1 &= f(a_1) \\ b_2 &= f(a_2) \end{aligned}$$

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

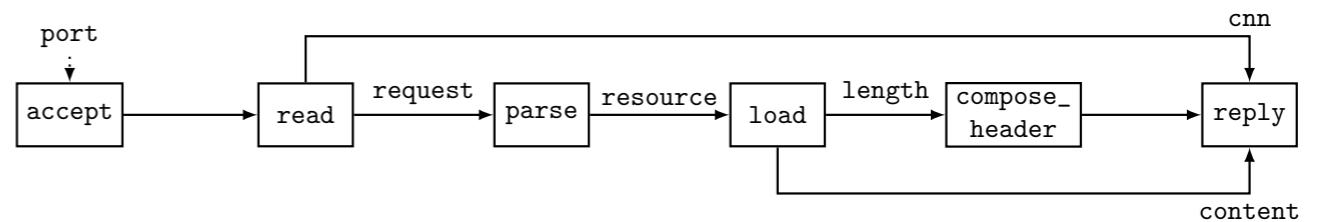
```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

$$\begin{array}{lll}
 I' = [b_1, b_2] = \text{map}(f, I) & = & I' = [b_1, b_2] = \text{smap}(f, I) \\
 b_1 = f(a_1) & & (s'_f, b_1) = f(s_f, a_1) \\
 b_2 = f(a_2) & & (s''_f, b_2) = f(s'_f, a_2)
 \end{array}$$

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

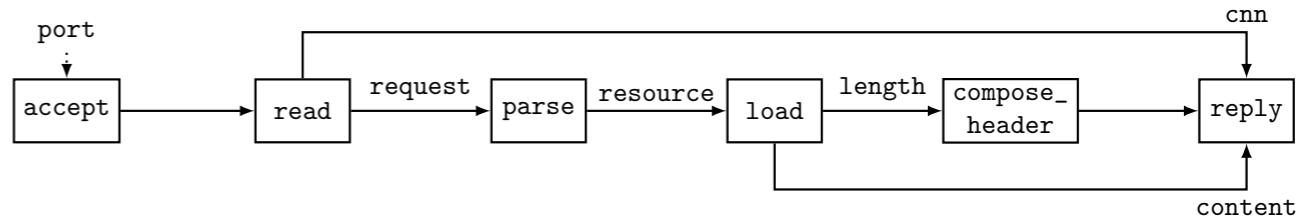
$$\begin{array}{lll}
 I' = [b_1, b_2] = \text{map}(f, I) & = & I' = [b_1, b_2] = \text{smap}(f, I) \\
 b_1 = f(a_1) & & (s'_f, b_1) = f(s_f, a_1) \\
 b_2 = f(a_2) & & (s''_f, b_2) = f(s'_f, a_2)
 \end{array}$$

data parallel

strictly sequential

# SMap

```
fn server(port) {
  let cnn = accept(port);
  let (cnn, req) = read(cnn);
  let resource = parse(req);
  let (content, length) = load(resource);
  reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
  let (cnn, req) = read(cnn);
  let resource = parse(req);
  let (content, length) = load(resource);
  reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
  let cnn = accept(port);      I = [a1, a2]
  request_handling(cnn)       f
}
```

$$\begin{array}{ccc} I' = [b_1, b_2] = \text{map}(f, I) & = & I' = [b_1, b_2] = \text{smap}(f, I) \\ b_1 = f(a_1) & & (s'_f, b_1) = f(s_f, a_1) \\ b_2 = f(a_2) & & (s''_f, b_2) = f(s'_f, a_2) \end{array}$$

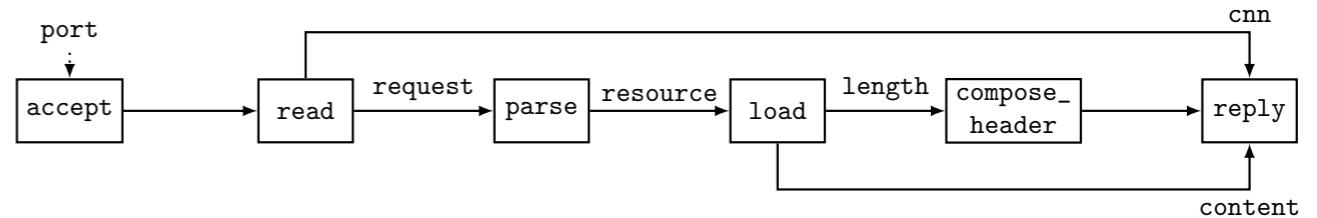
data parallel

```
fn f(a){  
  let x = g(a);  
  h(x)  
}  
→  
f = h ∘ g  
sf = (sg, sh)
```

strictly sequential

# SMap

```
fn server(port) {
  let cnn = accept(port);
  let (cnn, req) = read(cnn);
  let resource = parse(req);
  let (content, length) = load(resource);
  reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
  let (cnn, req) = read(cnn);
  let resource = parse(req);
  let (content, length) = load(resource);
  reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
  let cnn = accept(port);      I = [a1, a2]
  request_handling(cnn)       f
}
```

$$\begin{array}{lll} I' = [b_1, b_2] = \text{map}(f, I) & = & I' = [b_1, b_2] = \text{smap}(f, I) \\ b_1 = f(a_1) & & (s'_f, b_1) = f(s_g, a_1) \\ b_2 = f(a_2) & & (s''_f, b_2) = f(s'_g, a_2) \end{array}$$

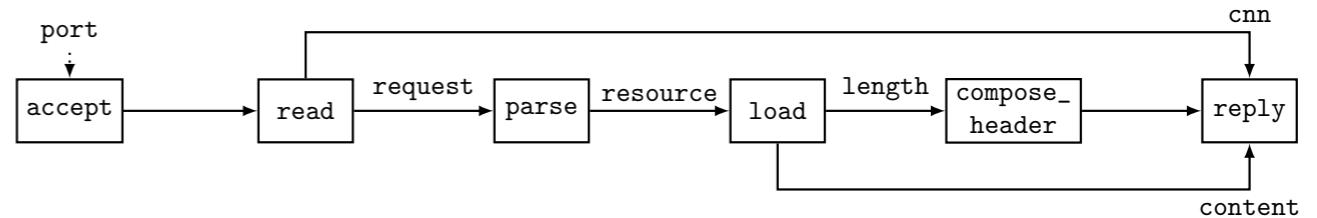
data parallel

$$\begin{array}{ll} \text{fn } f(a)\{\text{let } x = g(a); h(x)\} & (s'_g, x_1) = g(s_g, a_1) \\ \} & (s''_g, x_2) = g(s'_g, a_2) \\ \Rightarrow f = h \circ g & (s'_h, b_1) = h(s_g, x_1) \\ s_f = (s_g, s_h) & (s''_h, b_2) = h(s'_g, x_2) \end{array}$$

strictly sequential

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

$$\begin{array}{lll} I' = [b_1, b_2] = \text{map}(f, I) & = & I' = [b_1, b_2] = \text{smap}(f, I) \\ b_1 = f(a_1) & & (s'_f, b_1) = f(s_g, a_1) \\ b_2 = f(a_2) & & (s''_f, b_2) = f(s'_g, a_2) \end{array}$$

data parallel

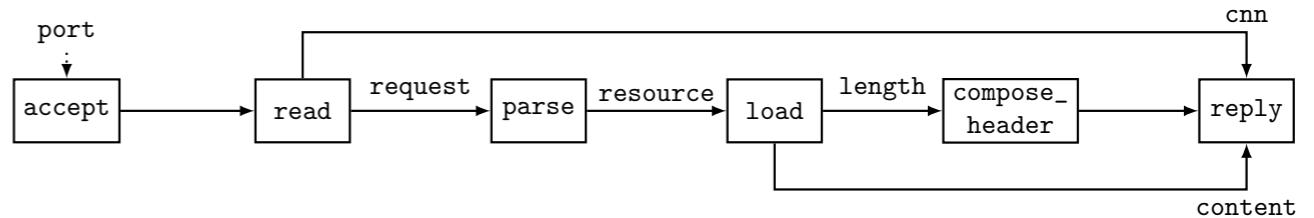
strictly sequential

$\begin{array}{l} \text{fn } f(a)\{\text{let } x = g(a); \\ \quad h(x) \}\\ \} \Rightarrow \\ f = h \circ g \\ s_f = (s_g, s_h) \end{array}$

$$\begin{array}{l} (s'_g, x_1) = g(s_g, a_1) \\ (s''_g, x_2) = g(s'_g, a_2) \\ (s'_h, b_1) = h(s_g, x_1) \\ (s''_h, b_2) = h(s'_h, x_2) \end{array}$$

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

$$I' = [b_1, b_2] = \text{map}(f, I) \quad \equiv \quad I' = [b_1, b_2] = \text{smap}(f, I)$$

$$\begin{aligned} b_1 &= f(a_1) & (s'_f, b_1) &= f(s_f, a_1) \\ b_2 &= f(a_2) & (s''_f, b_2) &= f(s'_f, a_2) \end{aligned}$$

data parallel

strictly sequential

$\text{fn } f(a)\{\$   
 $\text{let } x = g(a);$   
 $\text{h}(x)$   
 $\}$

➡

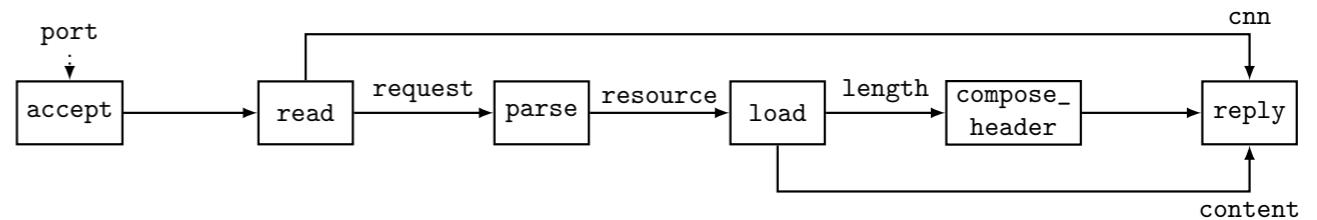
$f = h \circ g$   
 $s_f = (s_g, s_h)$

$(s'_g, x_1) = g(s_g, a_1)$	$(s'_g, x_1) = g(s_g, a_1)$
$(s''_g, x_2) = g(s'_g, a_2)$	$(s'_h, b_1) = h(s_h, x_1)$
$(s'_h, b_1) = h(s_h, x_1)$	$(s''_g, x_2) = g(s'_g, a_2)$
$(s''_h, b_2) = h(s'_h, x_2)$	$(s'_h, b_2) = h(s'_h, x_2)$

pipeline parallel

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



Abstract over processing:

```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnn = accept(port);      I = [a1, a2]
    request_handling(cnn)       f
}
```

$$I' = [b_1, b_2] = \text{map}(f, I) \quad = \quad I' = [b_1, b_2] = \text{smap}(f, I)$$

$$\begin{aligned} b_1 &= f(a_1) & (s'_f, b_1) &= f(s_f, a_1) \\ b_2 &= f(a_2) & (s''_f, b_2) &= f(s'_f, a_2) \end{aligned}$$

data parallel

Abstract over data:

Lists

- Finite
- Infinite (stream)

Generator

- Location independent

```
fn server(port) {
    let cnns:List<Cnn> = infinite_generator() || accept(port));
    let results:List<()> = smap(request_handling, cnns);
    ()
}
```

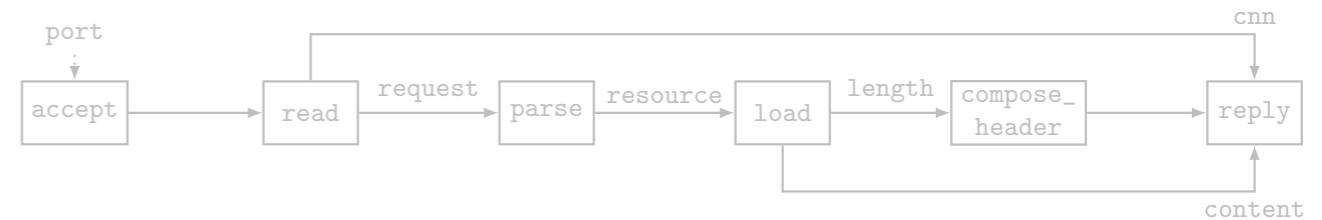
$$\begin{array}{lcl} \text{fn } f(a)\{ & (s'_g, x_1) = g(s_g, a_1) & (s'_g, x_1) = g(s_g, a_1) \\ \quad \text{let } x = g(a); & (s''_g, x_2) = g(s'_g, a_2) & (s'_h, b_1) = h(s_h, x_1) \\ \quad h(x) & (s'_h, b_1) = h(s_h, x_1) & \cancel{(s'_h, b_1) = h(s_h, x_1)} \\ \}\Rightarrow & (s''_h, b_2) = h(s'_h, x_2) & (s''_g, x_2) = g(s'_g, a_2) \\ f = h \circ g & & (s''_h, b_2) = h(s'_h, x_2) \\ s_f = (s_g, s_h) & & (s'_h, b_2) = h(s'_h, x_2) \end{array}$$

strictly sequential

pipeline parallel

# SMap

```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

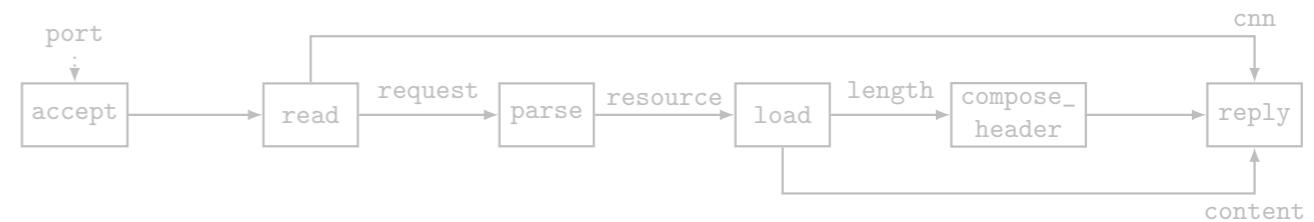


```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnns:List<Cnn> = infinite_generator() || accept(port));
    let results:List<()> = smap(request_handling, cnns);
    ()
}
```

# SMap as a Loop

```
fn server(port) {  
    let cnn = accept(port);  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```



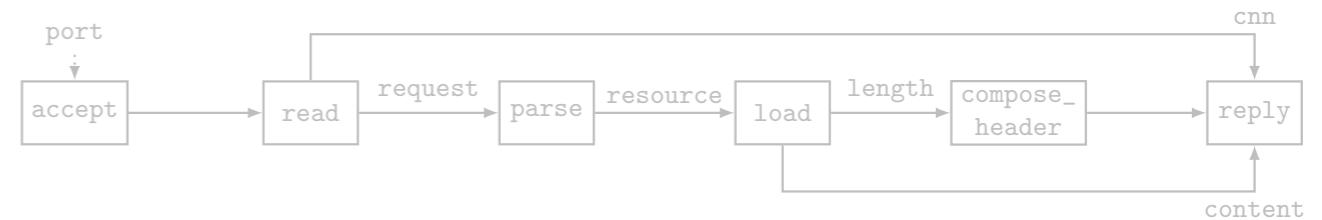
```
fn request_handling(cnn) {  
    let (cnn, req) = read(cnn);  
    let resource = parse(req);  
    let (content, length) = load(resource);  
    reply(cnn, compose_header(length), content)  
}
```

```
fn server(port) {  
    let cnns:List<Cnn> = infinite_generator(|| accept(port));  
    let results:List<()> = for cnn in cnns {  
        request_handling(cnn)  
    };  
    ()  
}
```

```
fn server(port) {  
    let cnns:List<Cnn> = infinite_generator(|| accept(port));  
    let results:List<()> = smap(request_handling, cnns);  
    ()  
}
```

# From SMap to Dataflow

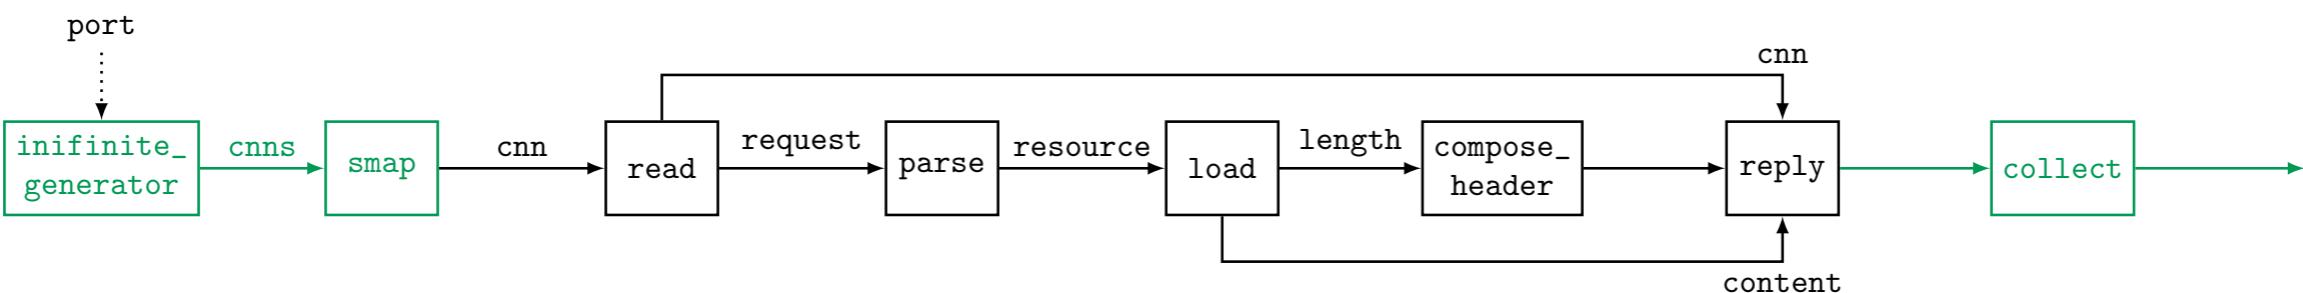
```
fn server(port) {
    let cnn = accept(port);
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```



```
fn request_handling(cnn) {
    let (cnn, req) = read(cnn);
    let resource = parse(req);
    let (content, length) = load(resource);
    reply(cnn, compose_header(length), content)
}
```

```
fn server(port) {
    let cnns:List<Cnn> = infinite_generator() || accept(port));
    let results:List<()> = for cnn in cnns {
        request_handling(cnn)
    };
    ()
}
```

```
fn server(port) {
    let cnns:List<Cnn> = infinite_generator() || accept(port));
    let results:List<()> = smap(request_handling, cnns);
    ()
}
```



Programming Model/  
Language



Compiler



Dataflow  
Runtime System

SFP:

- Variables
- Stateful Functions
- Composition
- Algorithms
- Control Flow
- SMap

Programming Model/  
Language



Compiler

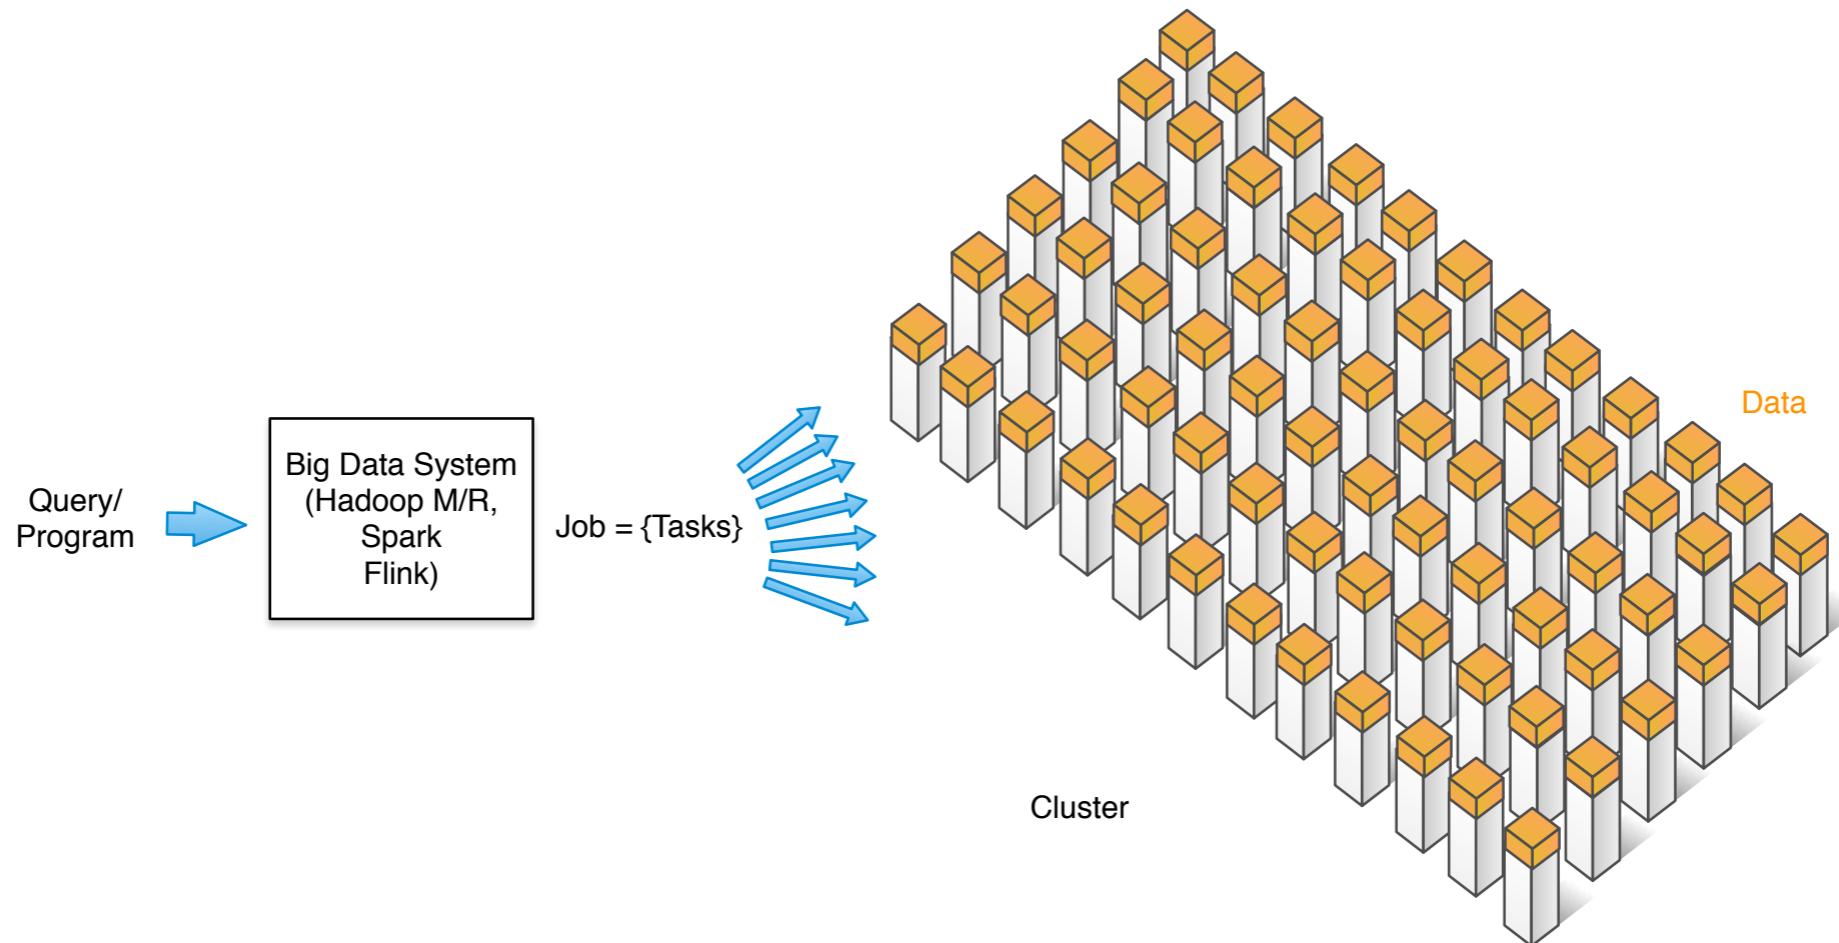


Dataflow  
Runtime System

- No concurrency abstractions.
- Practicality:
  - Integration into existing programming models (e.g., OOP)/languages.
  - Gradual switch/reuse of existing code base.

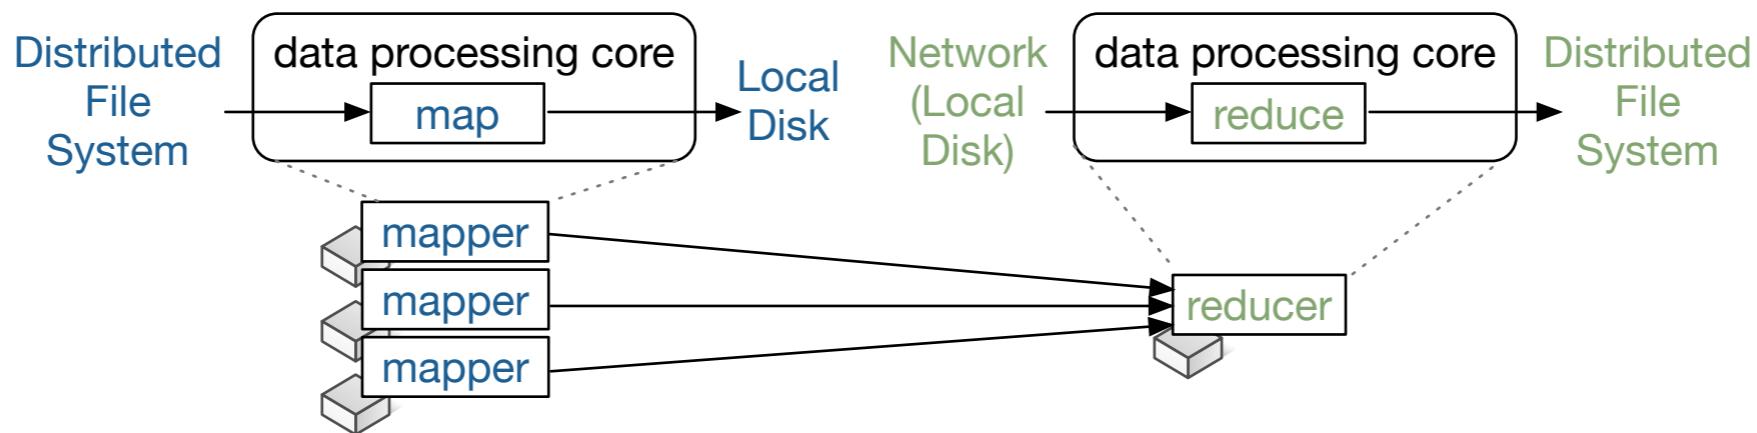
# Big Data Systems

1. Big Data Systems (BDSs) scale with the number of cores in the cluster
2. General wisdom: The main bottleneck is I/O (disk and network)



# Big Data Systems

1. Big Data Systems (BDSs) scale with the number of cores in the cluster **for independent tasks**.
2. General wisdom: The main bottleneck is I/O (disk and network)



---

Dean J, Ghemawat S. MapReduce: simplified data processing on large clusters. Communications of the ACM. 2008.

# Big Data Systems

1. Big Data Systems (BDSs) scale with the number of cores in the cluster **for independent tasks**.
2. General wisdom: The main bottleneck is I/O (disk and network) **for simple data**.

Jobs are CPU-bound!

WordCount, Sort	Analytics Queries (in Hive)
Simple Data Formats	vs.
Uncompressed	Complex Data Formats (Parquet, Tables, JSON)
	Compressed

---

Kay Ousterhout, Ryan Rasti, Sylvia Ratnasamy, Scott Shenker, and Byung-Gon Chun. 2015. Making sense of performance in data analytics frameworks.(NSDI'15).

# Big Data Systems

1. Big Data Systems (BDSs) scale with the number of cores in the cluster **for independent tasks**.
2. General wisdom: The main bottleneck is I/O (disk and network) **for simple data**.

Jobs are CPU-bound!

WordCount, Sort	Analytics Queries (in Hive)
Simple Data Formats	vs.
Uncompressed	Complex Data Formats (Parquet, Tables, JSON)
	Compressed

---

Kay Ousterhout, Ryan Rasti, Sylvia Ratnasamy, Scott Shenker, and Byung-Gon Chun. 2015. Making sense of performance in data analytics frameworks.(NSDI'15).

BDS cores/data processing pipelines do not scale well on multicores!

- Local optimizations do not solve this problem.
  - BDSs do not benefit from new network HW.
- Rewrite data processing cores!

---

Animesh Trivedi, Patrick Stuedi, Jonas Pfefferle, Radu Stoica, Bernard Metzler, Ioannis Koltsidas, and Nikolas Ioannou. 2016. On the [ir]relevance of network performance for data processing. (HotCloud'16).

# Data Processing Cores

The diagram illustrates the processing flow for three data processing cores: Hadoop, Spark, and Flink. It shows three main stages: **context**, **computation**, and **context**. In the first stage, **read I/O** leads to **deserialize**. In the **computation** stage, data undergoes **map, reduce, query graph**. In the final stage, **compress** leads to **write I/O**.

Stage	Hadoop	Spark	Flink
<b>context</b>	Context ctxt	TaskContext ctxt	TaskContext ctxt
<b>computation</b>	map, reduce, query graph	map, reduce, query graph	map, reduce, query graph
<b>context</b>	Context ctxt	TaskContext ctxt	TaskContext ctxt

```

(a) Hadoop
public class Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {
    /* The default implementation
     * is the identity function. */
    protected void map(KEYIN key, VALUEIN value,
                      Context ctxt)
        throws IOException {
        ctxt.write((KEYOUT) key,
                  (VALUEOUT) value);
    }
    public void run(Context ctxt) {
        while (ctxt.nextKeyValue())
            map(ctxt.getCurrentKey(),
                 ctxt.getCurrentValue(),
                 ctxt);
    }
}

(b) Spark
private[spark] class
ShuffleMapTask(partitionId: Int,
                partition: Partition)
  extends Task[MapStatus] {
  override def runTask(ctxt: TaskContext)
    throws IOException {
    val rdd = ...
    val writer = new BlockWriter[Product2[Any, Any]](
      partitionId, ctxt)
    writer.write(
      rdd.iterator(partition, ctxt)
        .asInstanceOf[Iterator[_ <:> Product2[Any, Any]]])
    writer.stop(success = true).get}
}

(c) Flink
public class DataSourceTask<OT>
  extends AbstractInvokable {
  private InputFormat<OT, InputSplit> format;
  private Collector<OT> output;
  ...
  public void execute() throws Exception {
    ...
    OT returned;
    if ((returned =
        format.nextRecord(reuse)) != null)
      output.collect(returned);
  }
}

```

# Rewriting Hadoop Map/Reduce

```
fn output_side(writer, key, value) {
    output(writer, key, value)
}

fn compute_and_output(mapper, writer, split) {
    let (line, content) = split;
    let kv_pairs = hmr_map(mapper, line, content);
    for kv_pair in kv_pairs {
        let (k, v) = kv_pair;
        output_side(writer, k, v)
    }
}

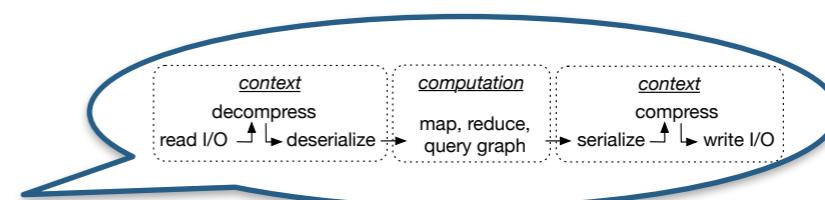
fn coarse(reader: org.apache.hadoop.mapreduce.Mapper$Context,
          mapper: org.apache.hadoop.mapreduce.Mapper,
          writer: org.apache.hadoop.mapreduce.Mapper$Context) {
    let splits = finite_generator(|| load_split(reader),
                                || has_more_data(reader));
    for split in splits {
        compute_and_output(mapper, writer, split)
    }
}
```

# Rewriting Hadoop Map/Reduce

```
fn output_side(writer, key, value) {
    output(writer, key, value)
}

fn compute_and_output(mapper, writer, split) {
    let (line, content) = split;
    let kv_pairs = hmr_map(mapper, line, content);
    for kv_pair in kv_pairs {
        let (k, v) = kv_pair;
        output_side(writer, k, v)
    }
}

fn coarse(reader: org.apache.hadoop.mapreduce.Mapper$Context,
          mapper: org.apache.hadoop.mapreduce.Mapper,
          writer: org.apache.hadoop.mapreduce.Mapper$Context) {
    let splits = finite_generator(|| load_split(reader),
                                || has_more_data(reader));
    for split in splits {
        compute_and_output(mapper, writer, split)
    }
}
```

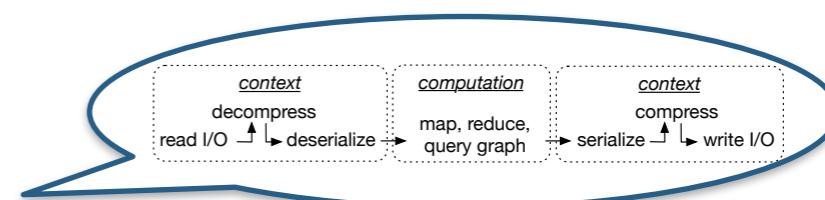


# Rewriting Hadoop Map/Reduce

```
fn output_side(writer, key, value) {
    output(writer, key, value)
}

fn compute_and_output(mapper, writer, split) {
    let (line, content) = split;
    let kv_pairs = hmr_map(mapper, line, content);
    for kv_pair in kv_pairs {
        let (k, v) = kv_pair;
        output_side(writer, k, v)
    }
}

fn coarse(reader: org.apache.hadoop.mapreduce.Mapper$Context,
          mapper: org.apache.hadoop.mapreduce.Mapper,
          writer: org.apache.hadoop.mapreduce.Mapper$Context) {
    let splits = finite_generator(|| load_split(reader),
                                || has_more_data(reader));
    for split in splits {
        compute_and_output(mapper, writer, split)
    }
}
```

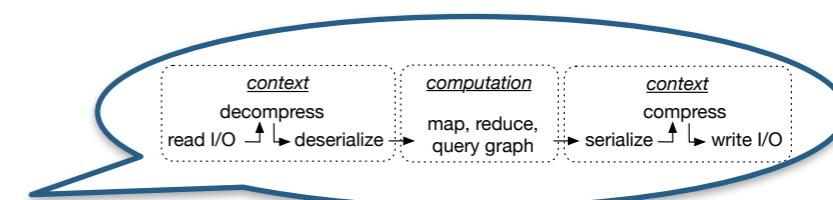


# Rewriting Hadoop Map/Reduce

```
fn output_side(writer, key, value) {
    output(writer, key, value)
}
    writer.write(key, value)

fn compute_and_output(mapper, writer, split) {
    let (line, content) = split;
    let kv_pairs = hmr_map(mapper, line, content);
    for kv_pair in kv_pairs {
        let (k, v) = kv_pair;
        output_side(writer, k, v)
    }
}

fn coarse(reader: org.apache.hadoop.mapreduce.Mapper$Context,
          mapper: org.apache.hadoop.mapreduce.Mapper,
          writer: org.apache.hadoop.mapreduce.Mapper$Context) {
    let splits = finite_generator(|| load_split(reader),
                                || has_more_data(reader));
    for split in splits {
        compute_and_output(mapper, writer, split)
    }
}
```



# Rewriting Hadoop Map/Reduce

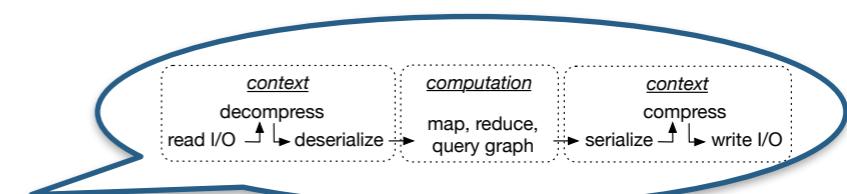
```
fn output_side(writer, key, value) {
    output(writer, key, value)
}
    writer.write(key, value)

fn compute_and_output(mapper, writer, split) {
    let (line, content) = split;
    let kv_pairs = hmr_map(mapper, line, content);
    for kv_pair in kv_pairs {
        let (k, v) = kv_pair;
        output_side(writer, k, v)
    }
}

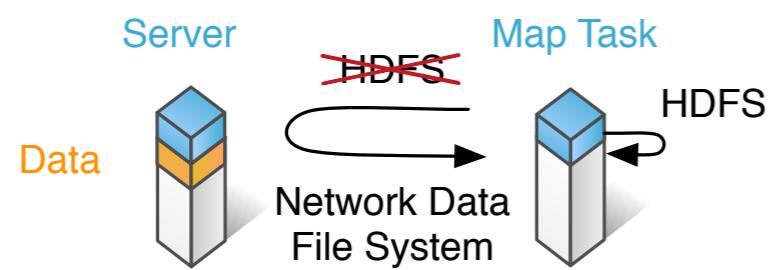
fn coarse(reader: org.apache.hadoop.mapreduce.Mapper$Context,
          mapper: org.apache.hadoop.mapreduce.Mapper,
          writer: org.apache.hadoop.mapreduce.Mapper$Context) {
    let splits = finite_generator(|| load_split(reader),
                                || has_more_data(reader));
    for split in splits {
        compute_and_output(mapper, writer, split)
    }
}
```

## 4 variants:

- Coarse (C)
- Coarse Input Fine Output (CIFO)
- Fine Input Coarse Output (FICO)
- Fine (F)



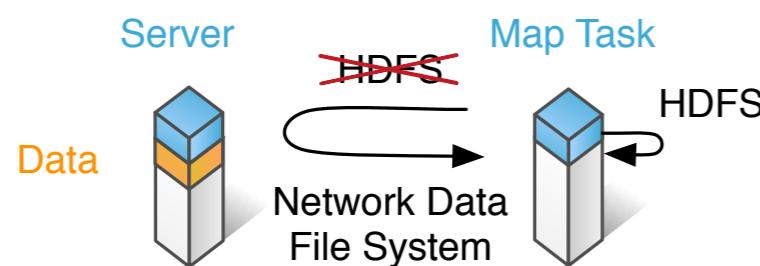
# Evaluation



- SequenceFile
- JSON
- Snappy, LZO
- TPC-H Parts table

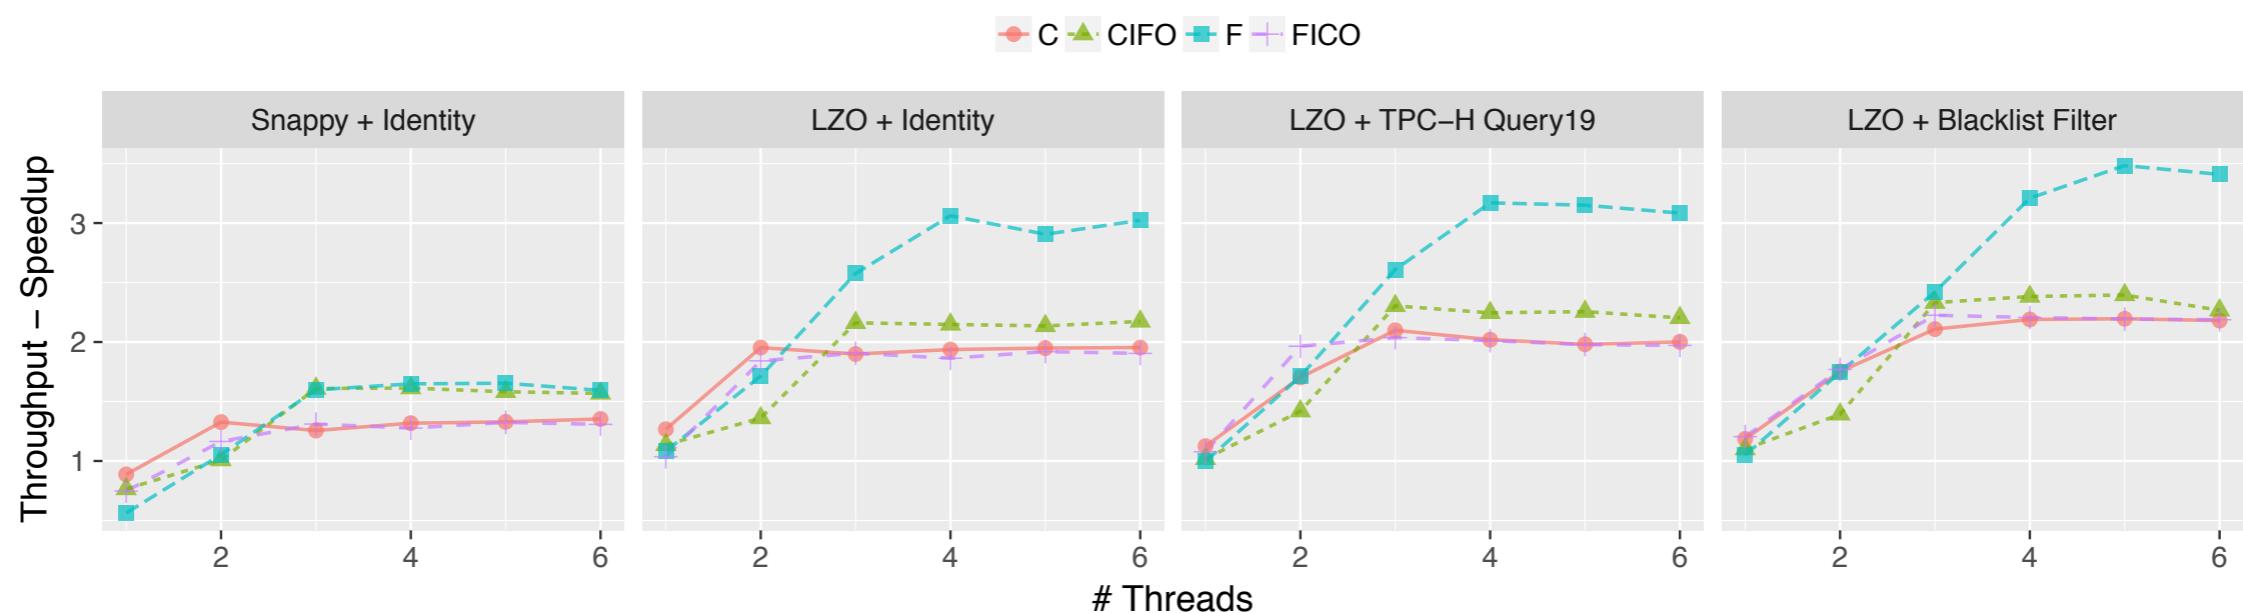
Almost the same setup as for  
a reduce task!

# Throughput Analysis



- SequenceFile
- JSON
- Snappy, LZO
- TPC-H Parts table

Almost the same setup as for  
a reduce task!



Speedup of up to 3.5x for compute-intensive configurations!

Programming Model/  
Language



Compiler



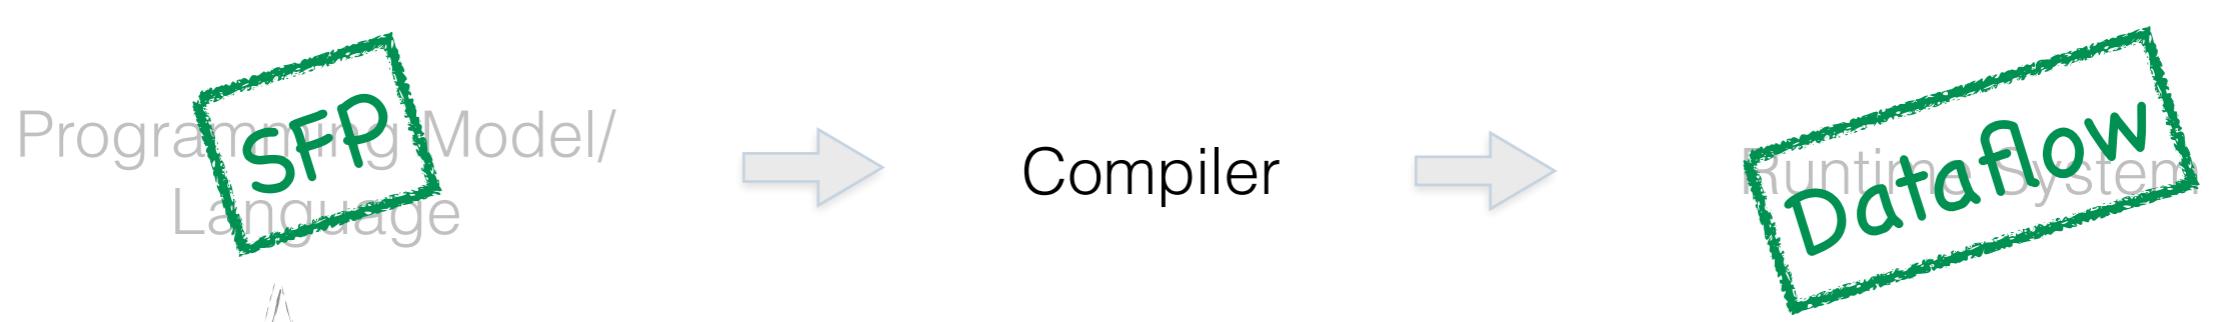
Runtime System  
**Dataflow**

SFP:

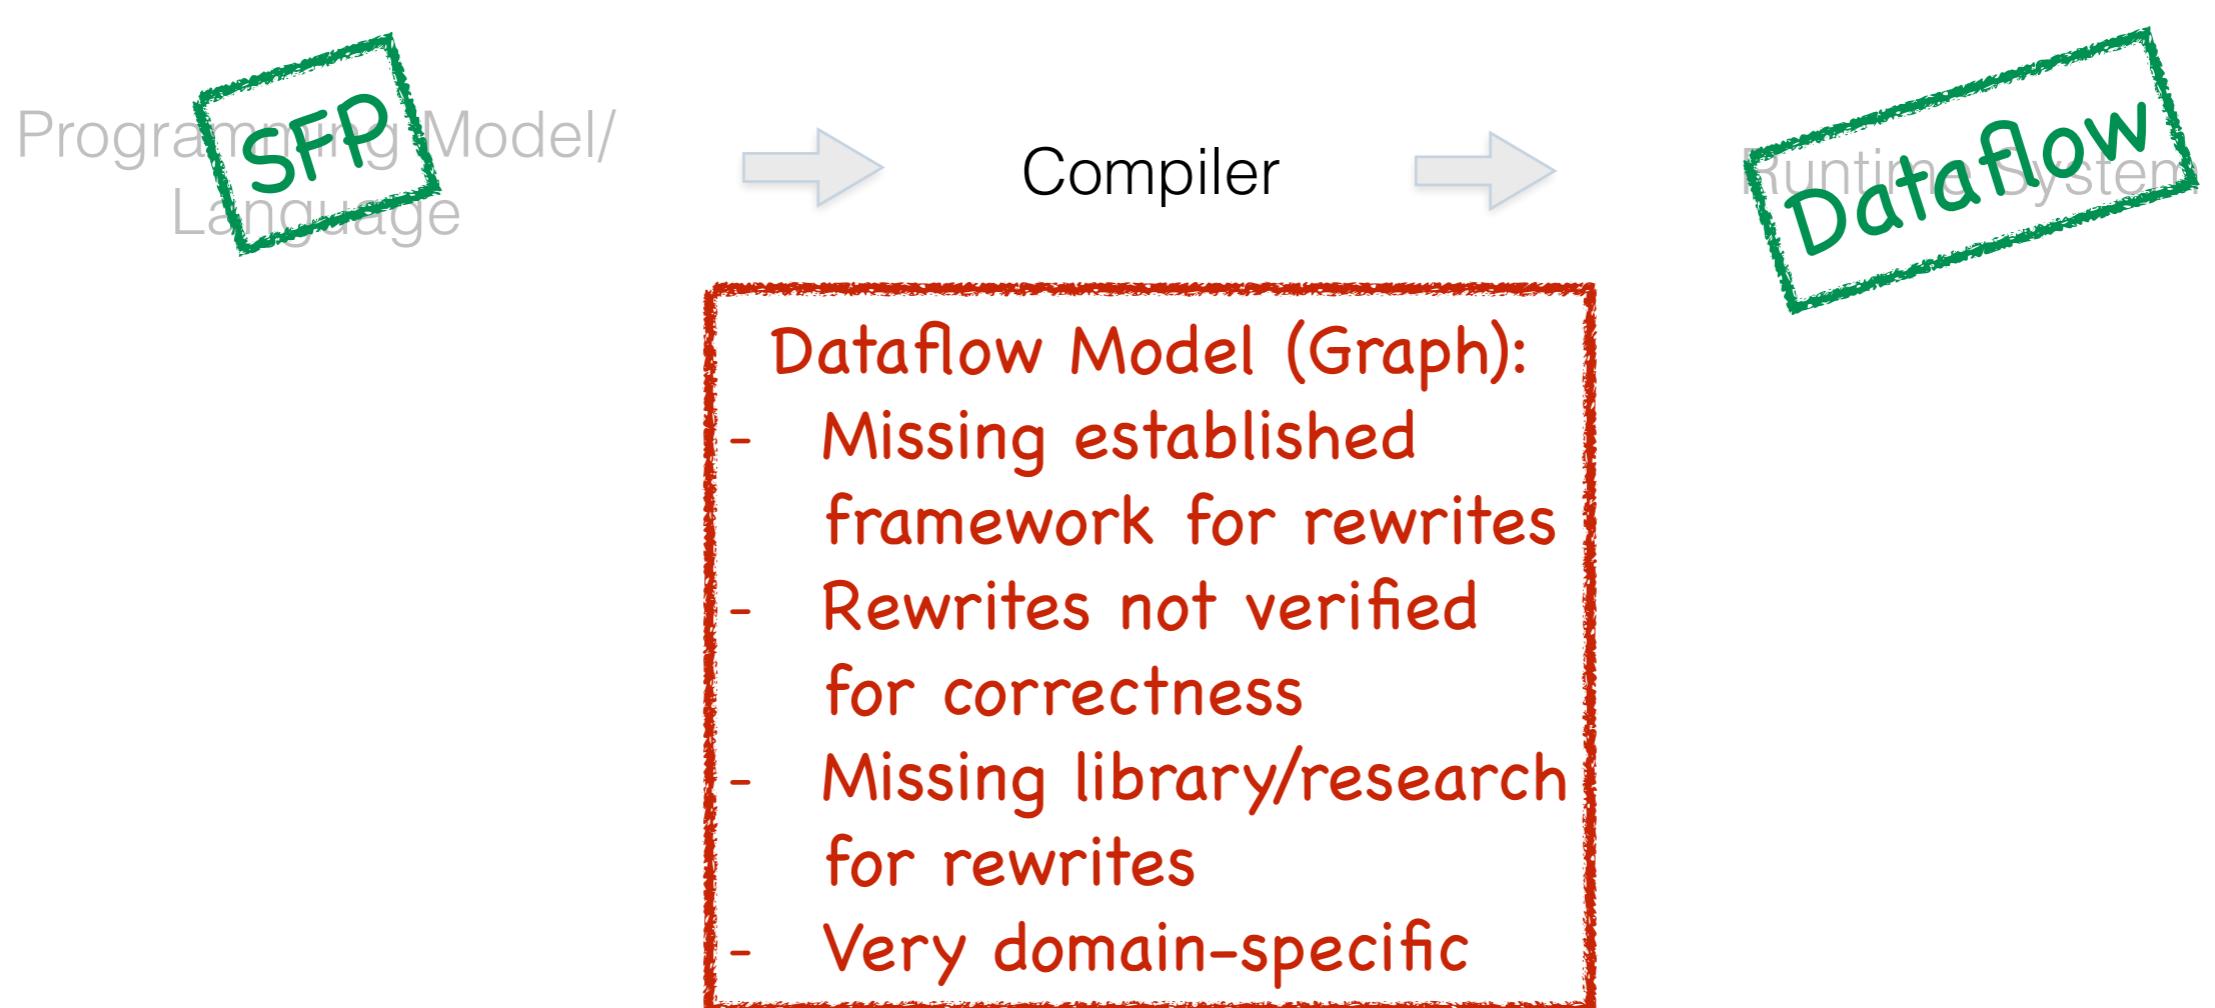
- Variables
- Stateful Functions
- Composition
- Algorithms
- Control Flow
- SMap

Sebastian Ertel, Christof Fetzer, and Pascal Felber. 2015. Ohua: Implicit Dataflow Programming for Concurrent Systems. In Proceedings of the Principles and Practices of Programming on The Java Platform (PPPJ '15). ACM.

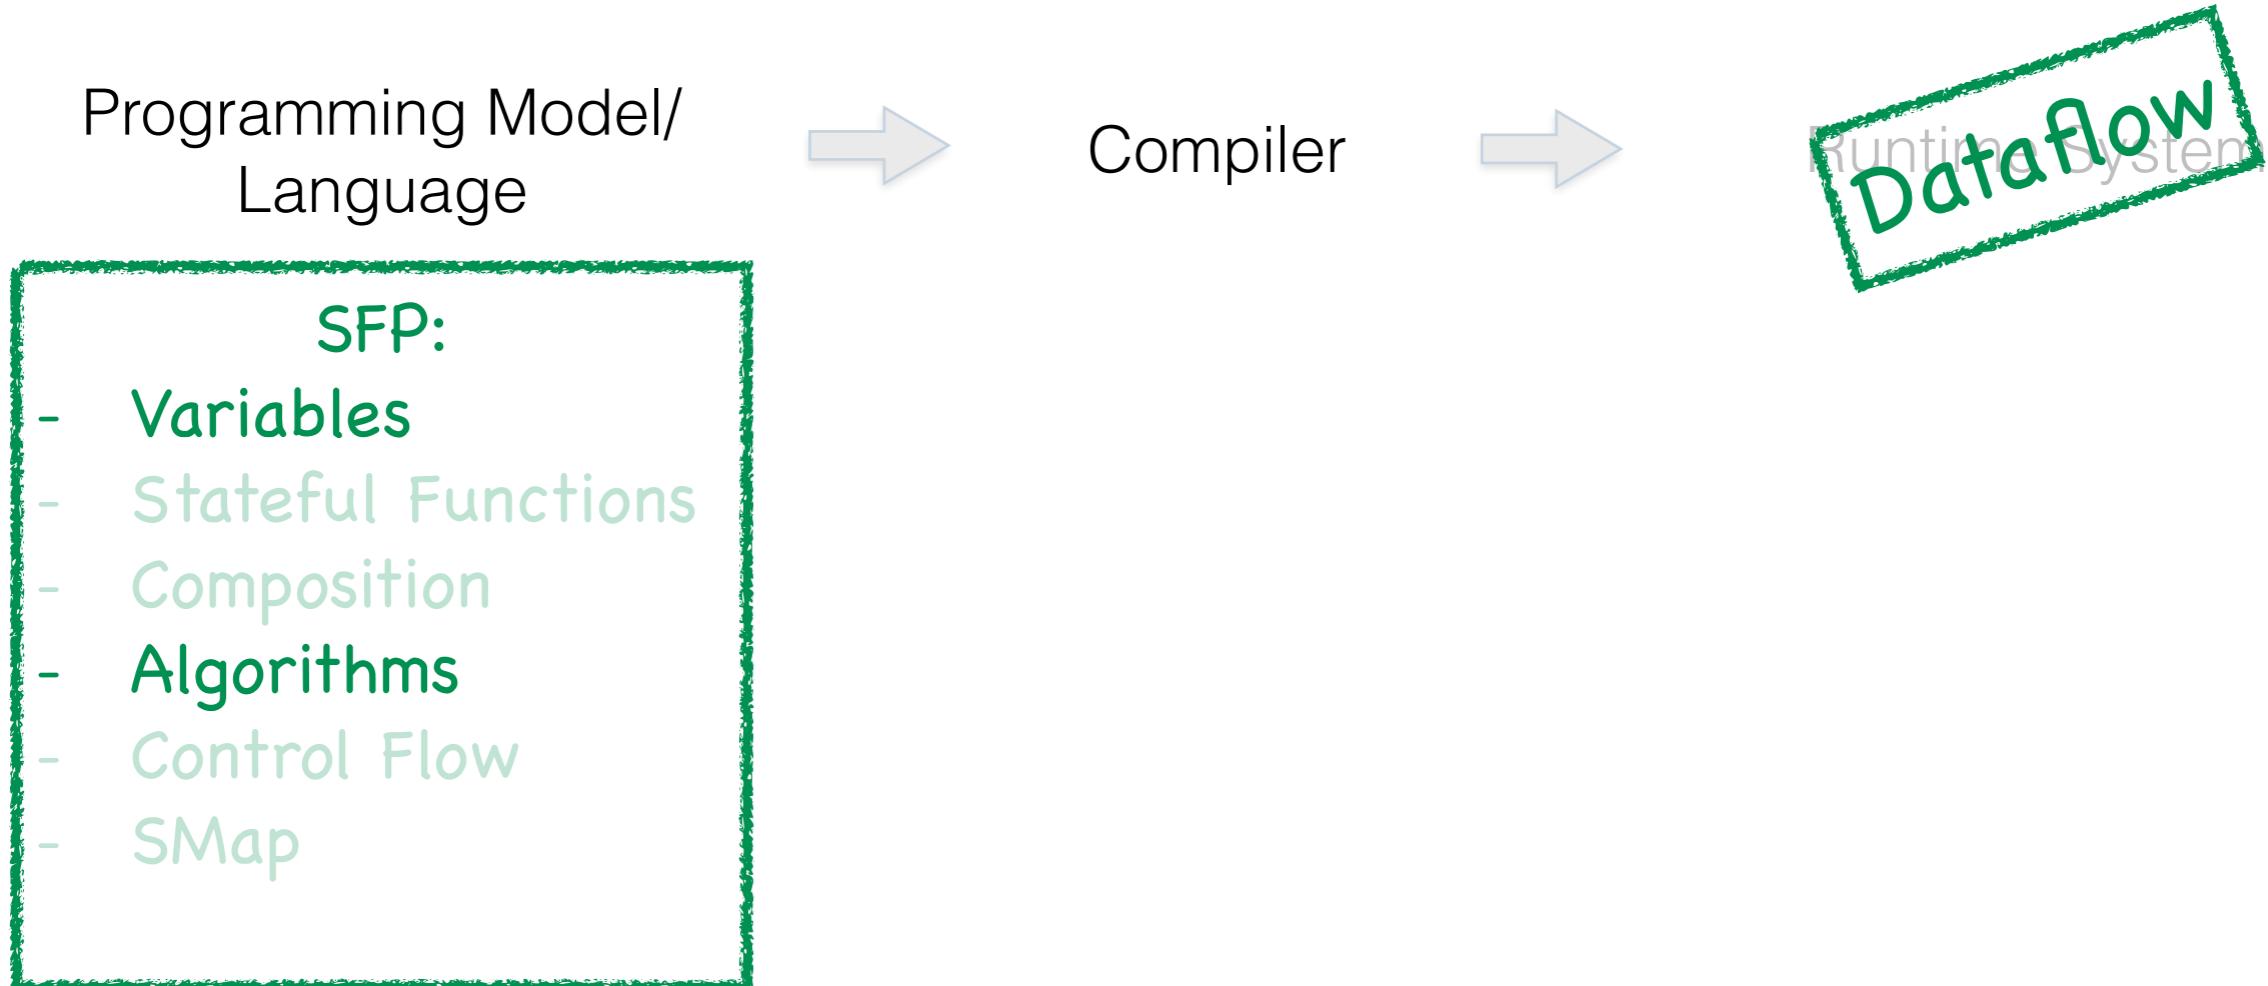
Sebastian Ertel, Justus Adam, and Jeronimo Castrillon. 2018. Supporting Fine-grained Dataflow Parallelism in Big Data Systems. In Proceedings of the 9th International Workshop on Programming Models and Applications for Multicores and Manycores (PMAM'18). ACM.



- No concurrency abstractions.
- Practicality:
  - Integration into existing programming models (e.g., OOP)/ languages.
  - Gradual switch/reuse of existing code base.

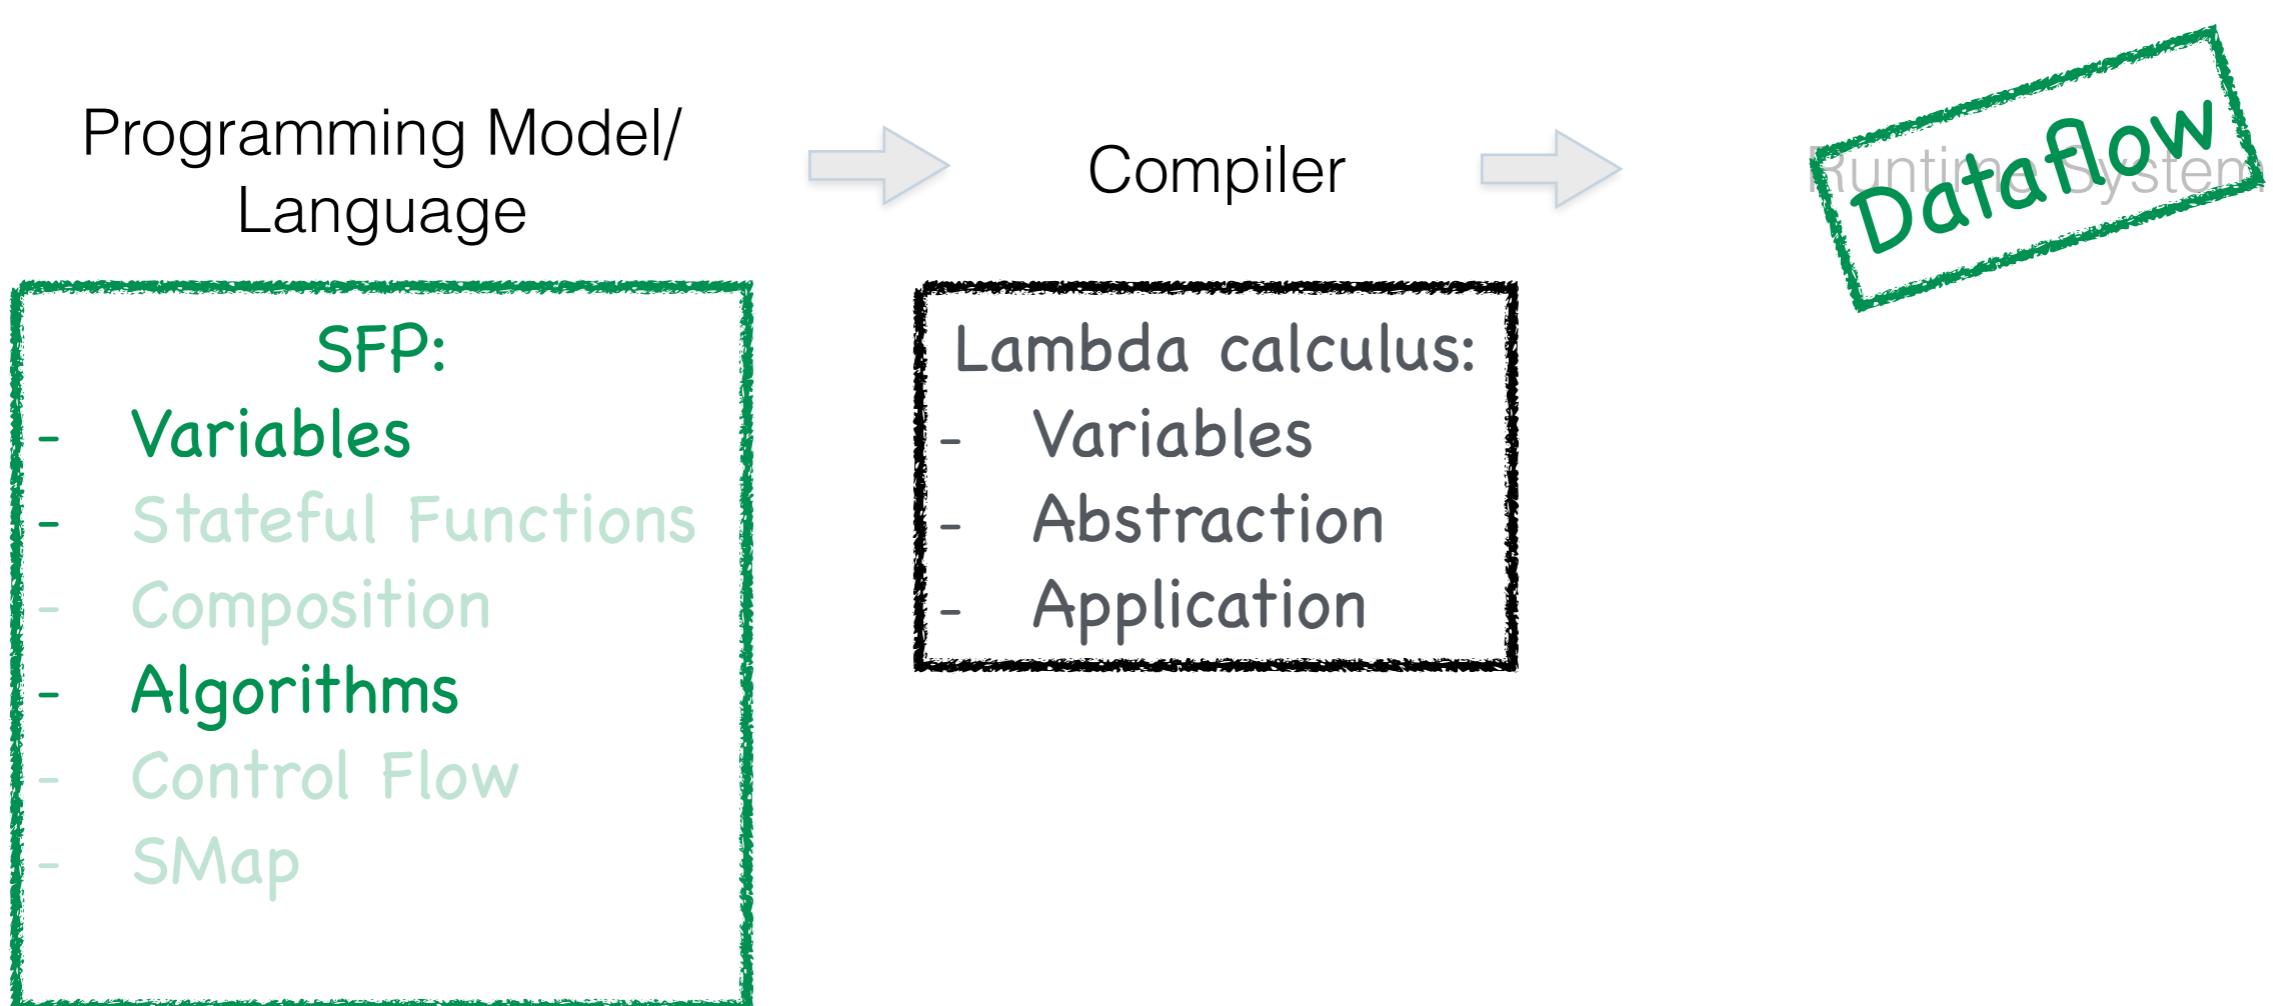


# PL/Compiler Co-Design



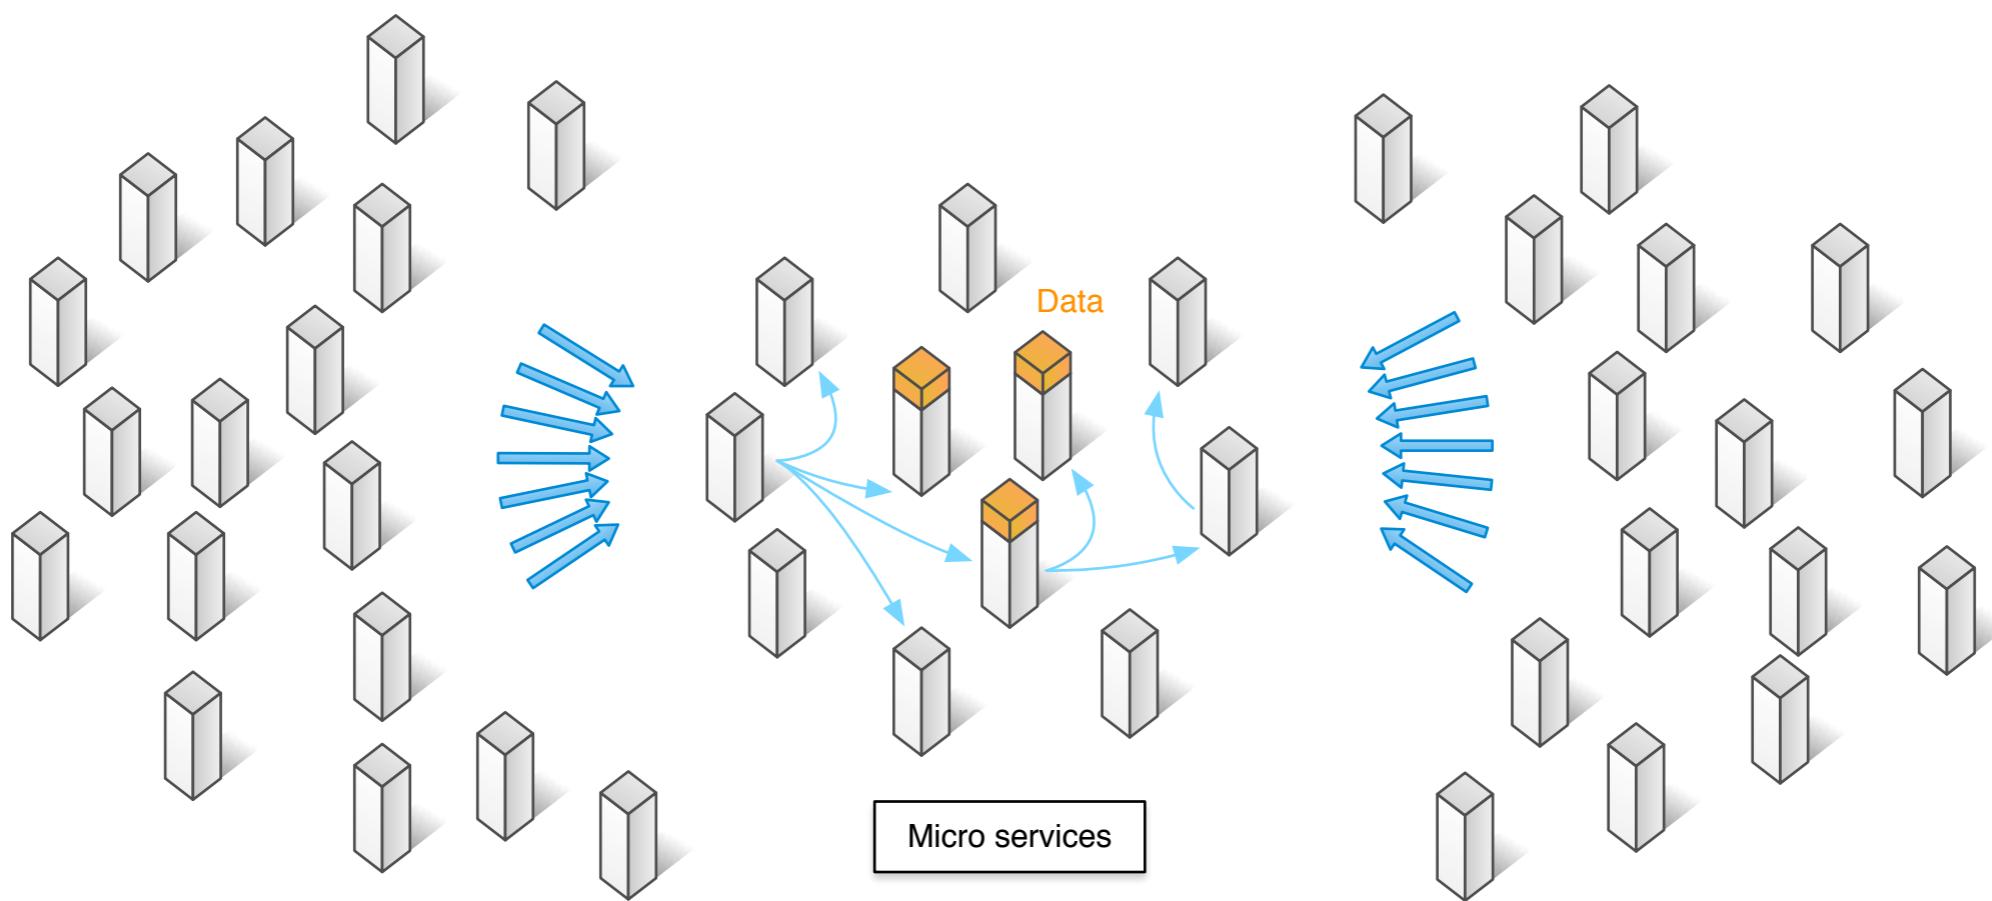
J. Backus. Can programming be liberated from the von neumann style?:A functional style and its algebra of programs. Commun.ACM,1978.  
Barendregt, Hendrik Pieter. "Functional programming and lambda calculus." *Formal models and semantics*. Elsevier, 1990.

# PL/Compiler Co-Design



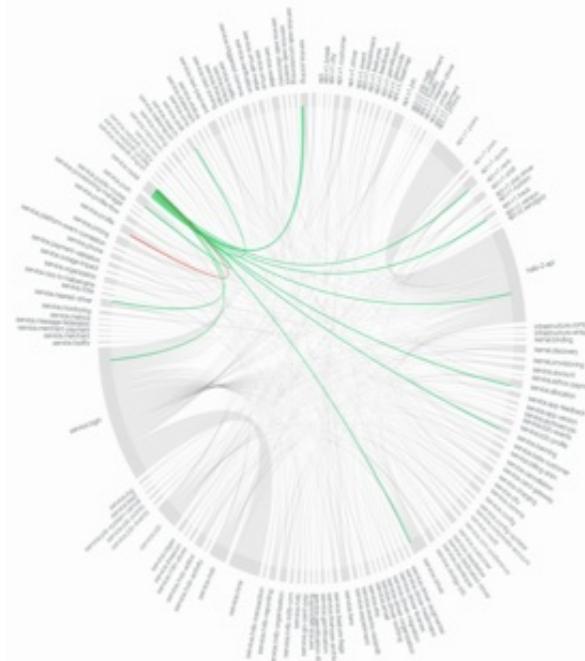
J. Backus. Can programming be liberated from the von neumann style?:A functional style and its algebra of programs. Commun.ACM,1978.  
Barendregt, Hendrik Pieter. "Functional programming and lambda calculus." *Formal models and semantics*. Elsevier, 1990.

# Microservice Architectures



# Redefining the Death Star

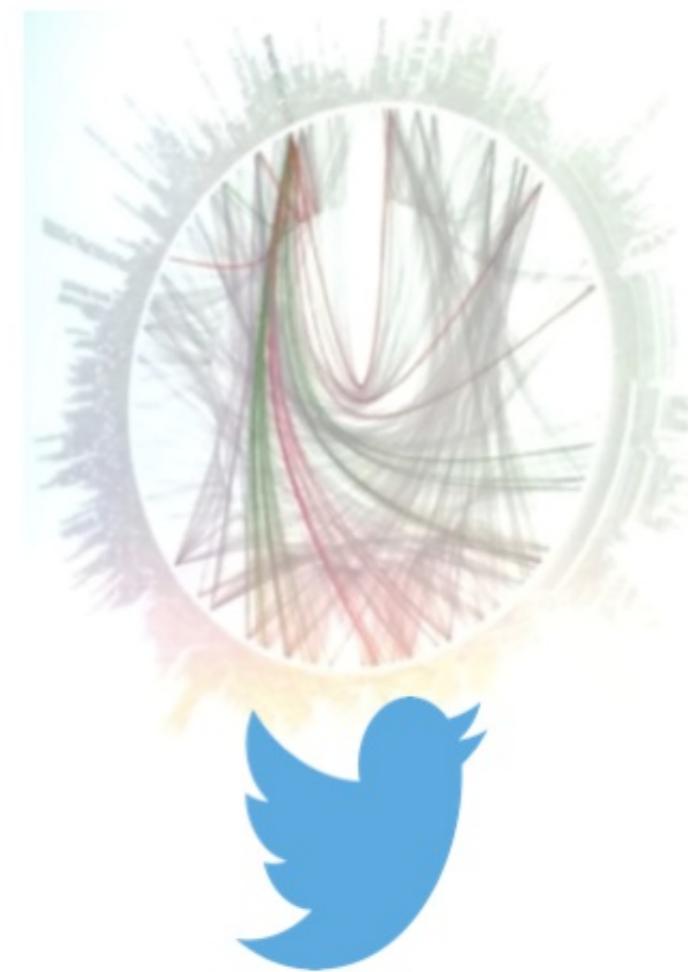
450 microservices



500+ microservices



500+ microservices



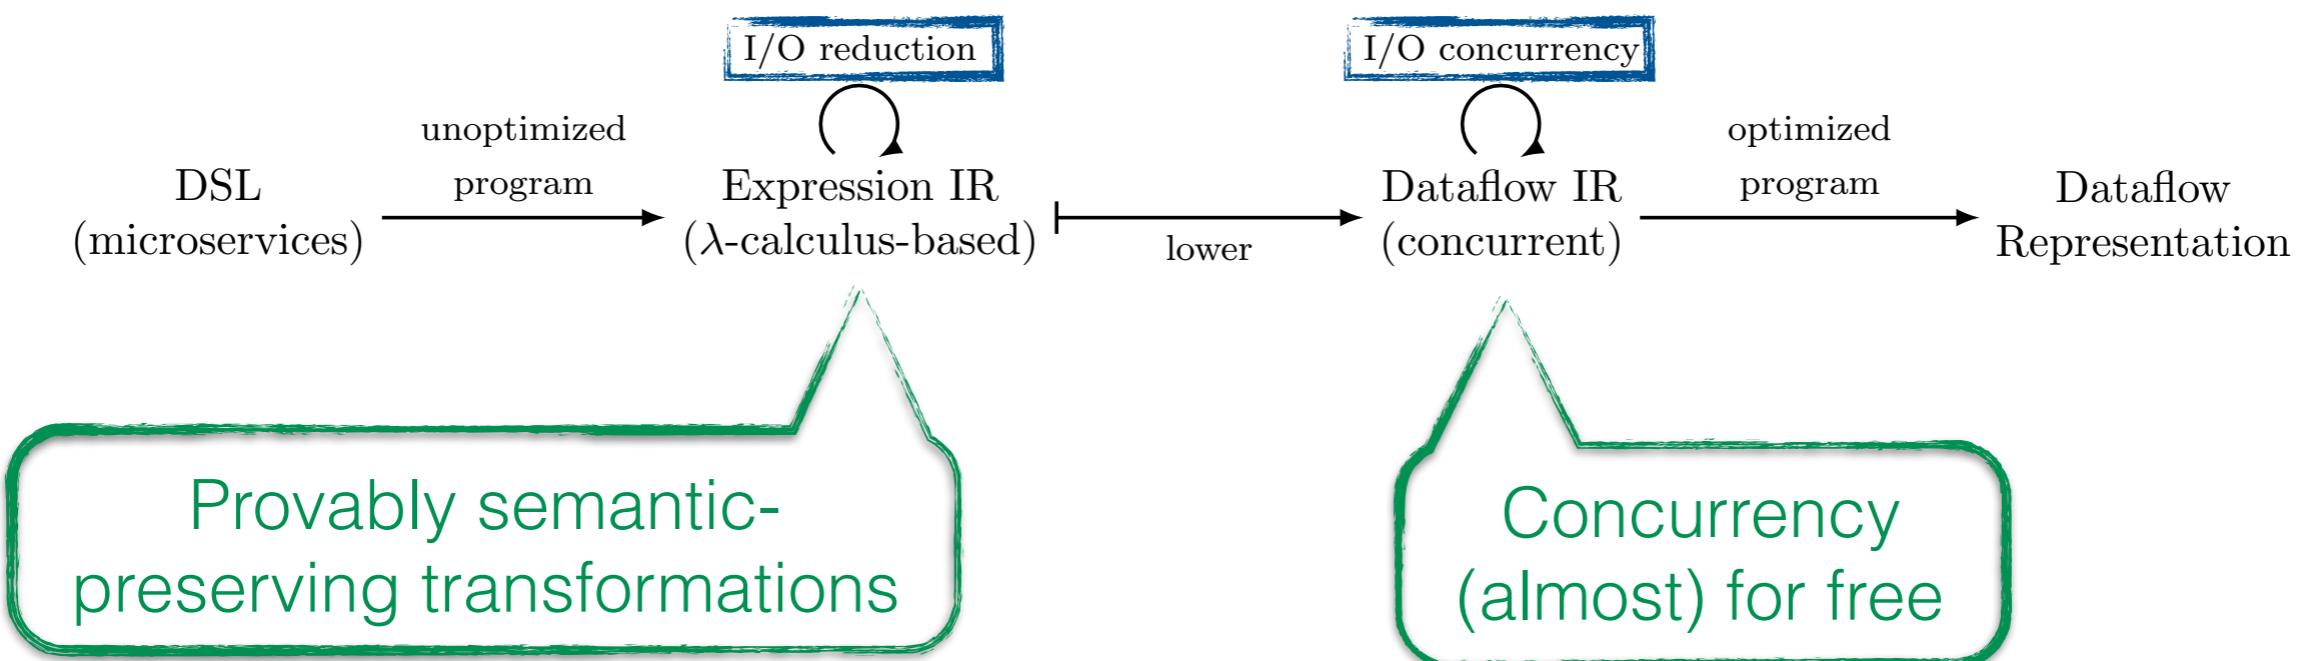
Source:

Netflix: <http://www.slideshare.net/BruceWong3/the-case-for-chaos>

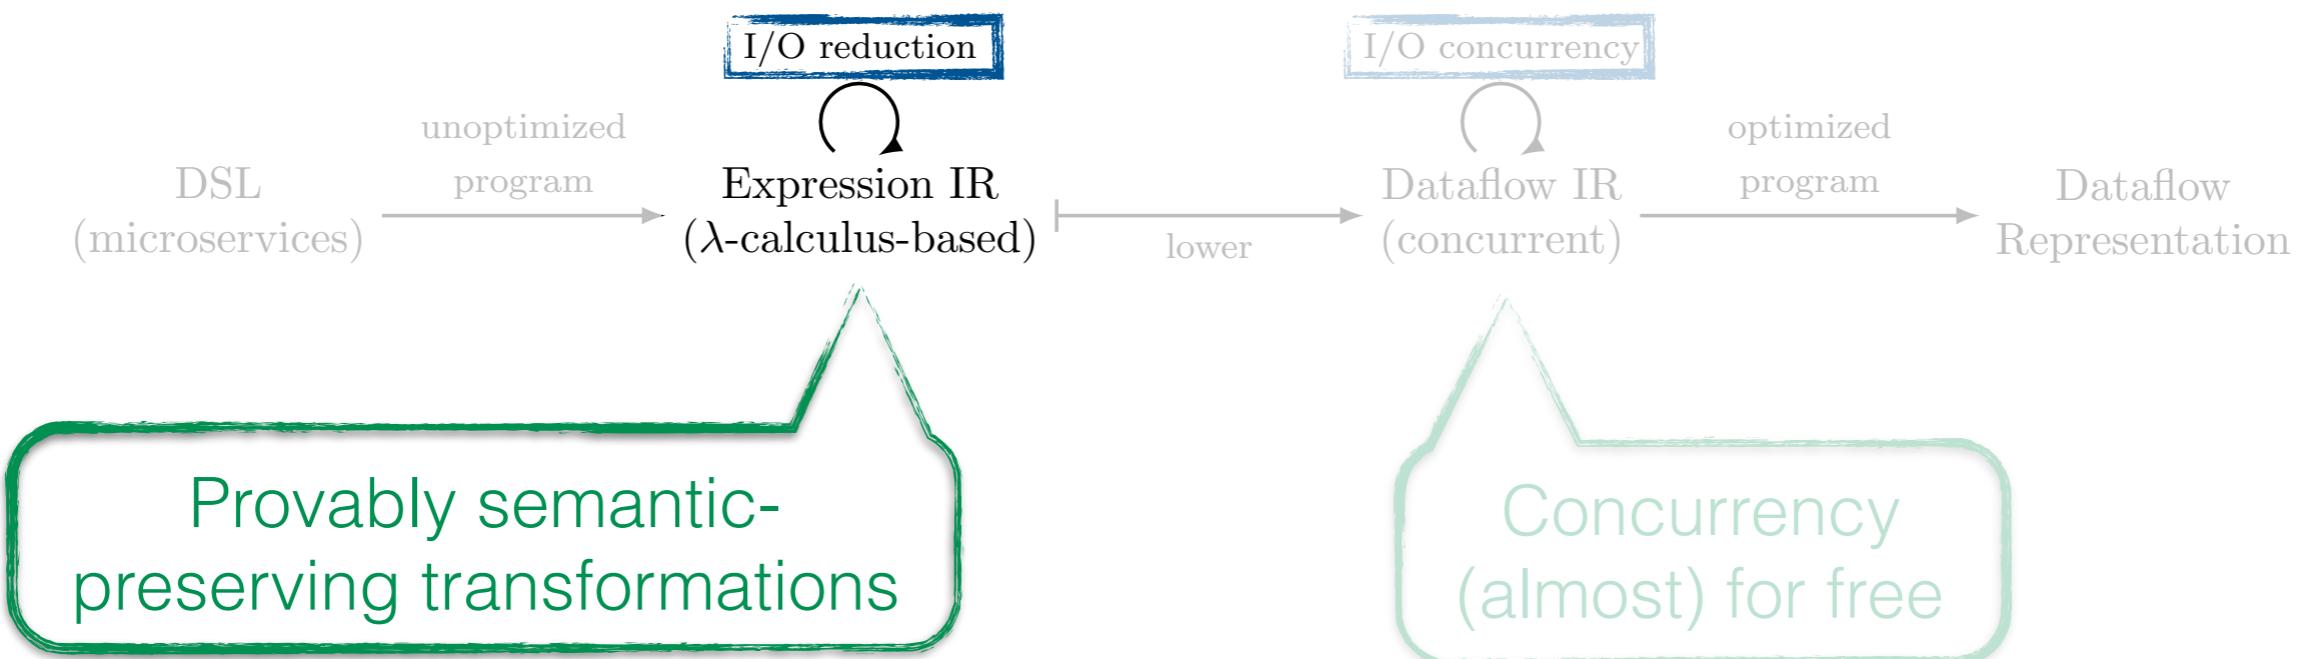
Twitter: <https://twitter.com/adrianco/status/441883572618948608>

Hail-o: <https://sudo.hailoapp.com/services/2015/03/09/journey-into-a-microservice-world-part-3/>

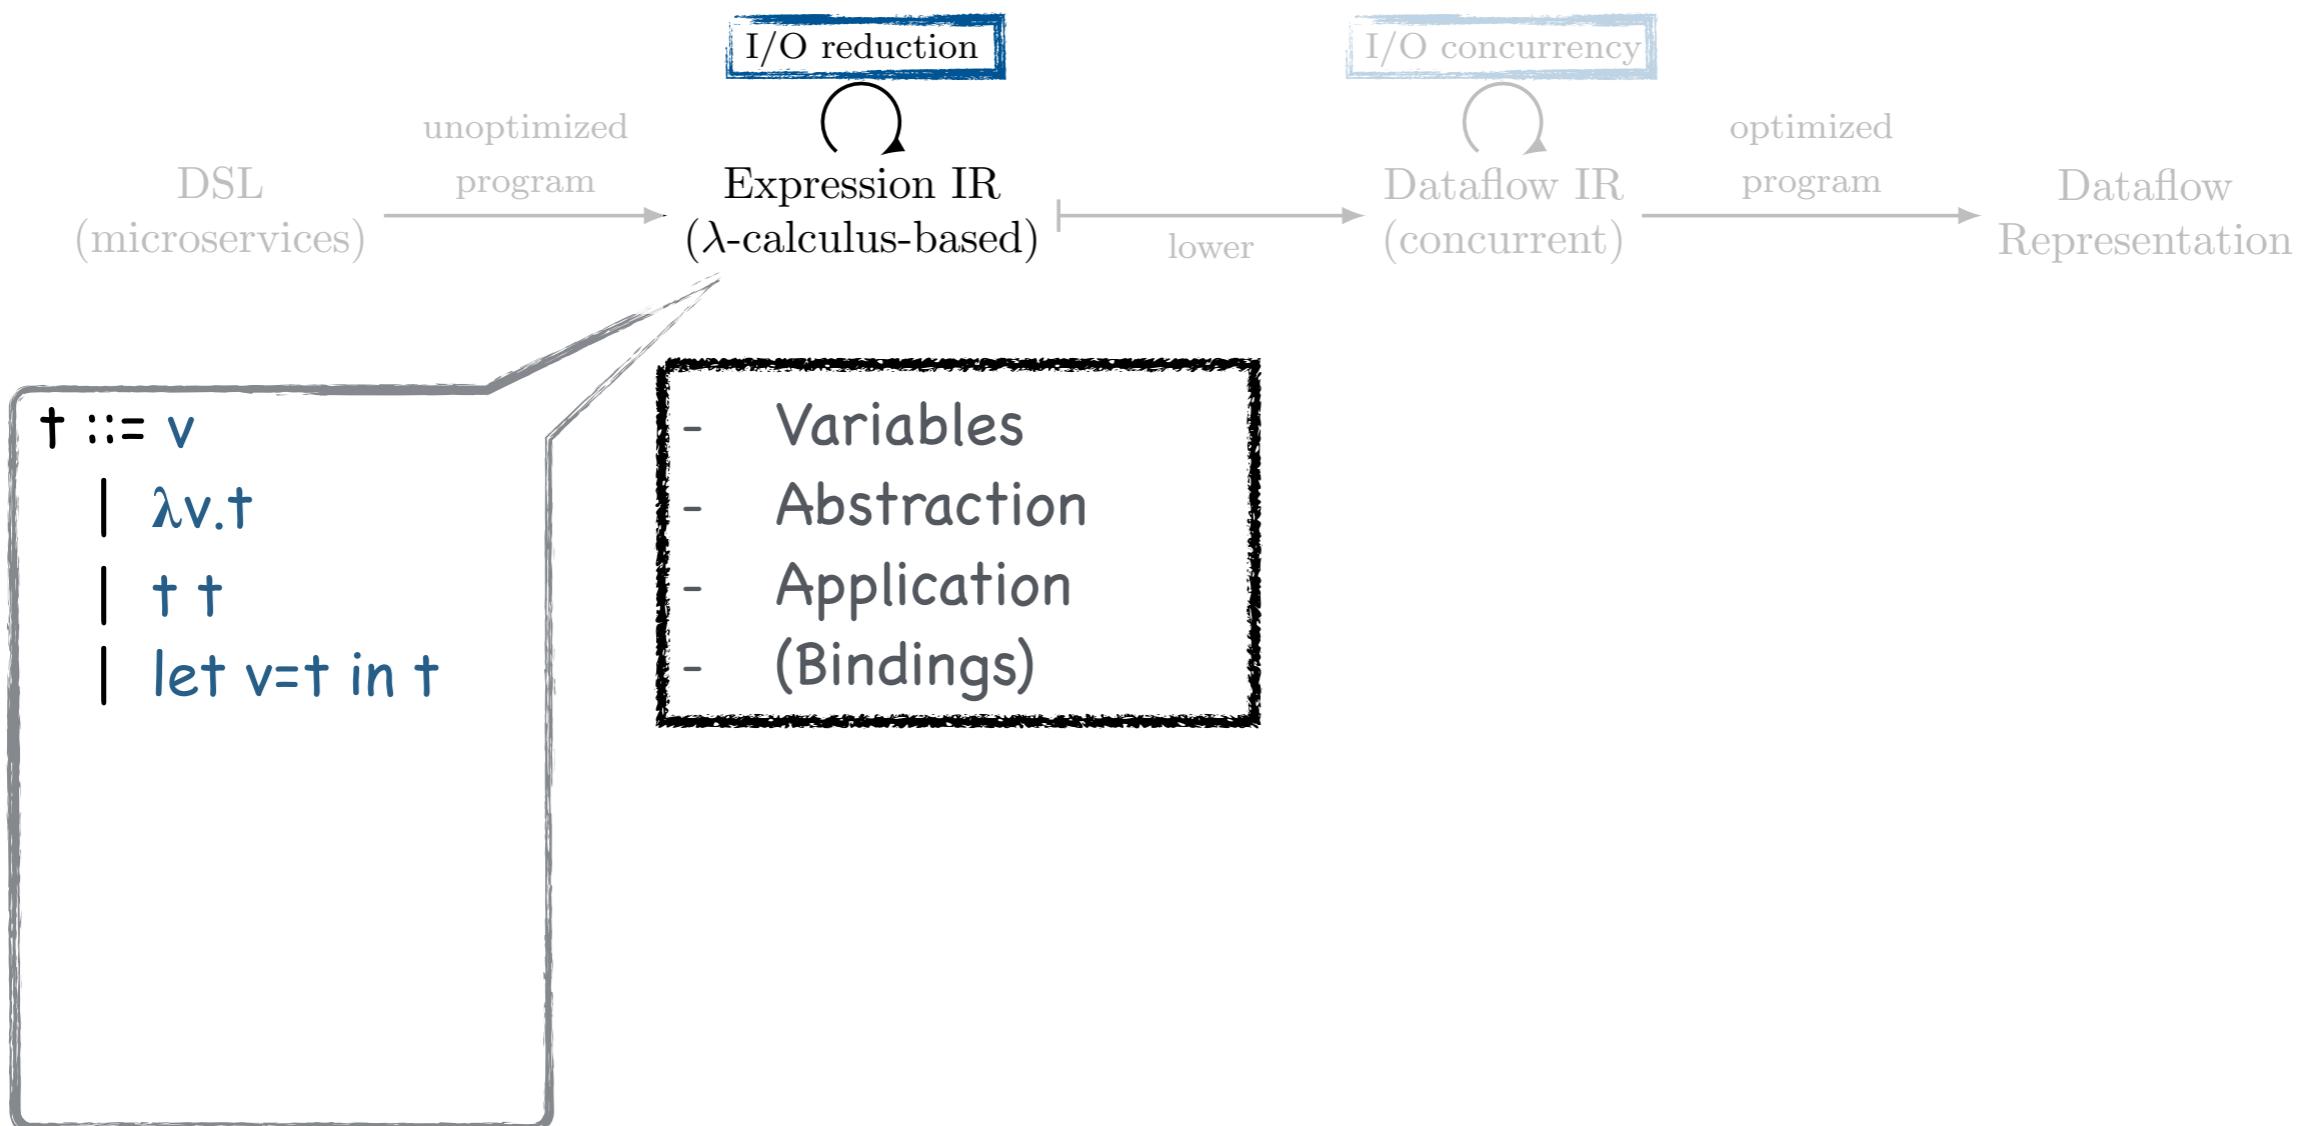
# Compiler



# Compiler



# Foundation: Call-by-need Lambda Calculus

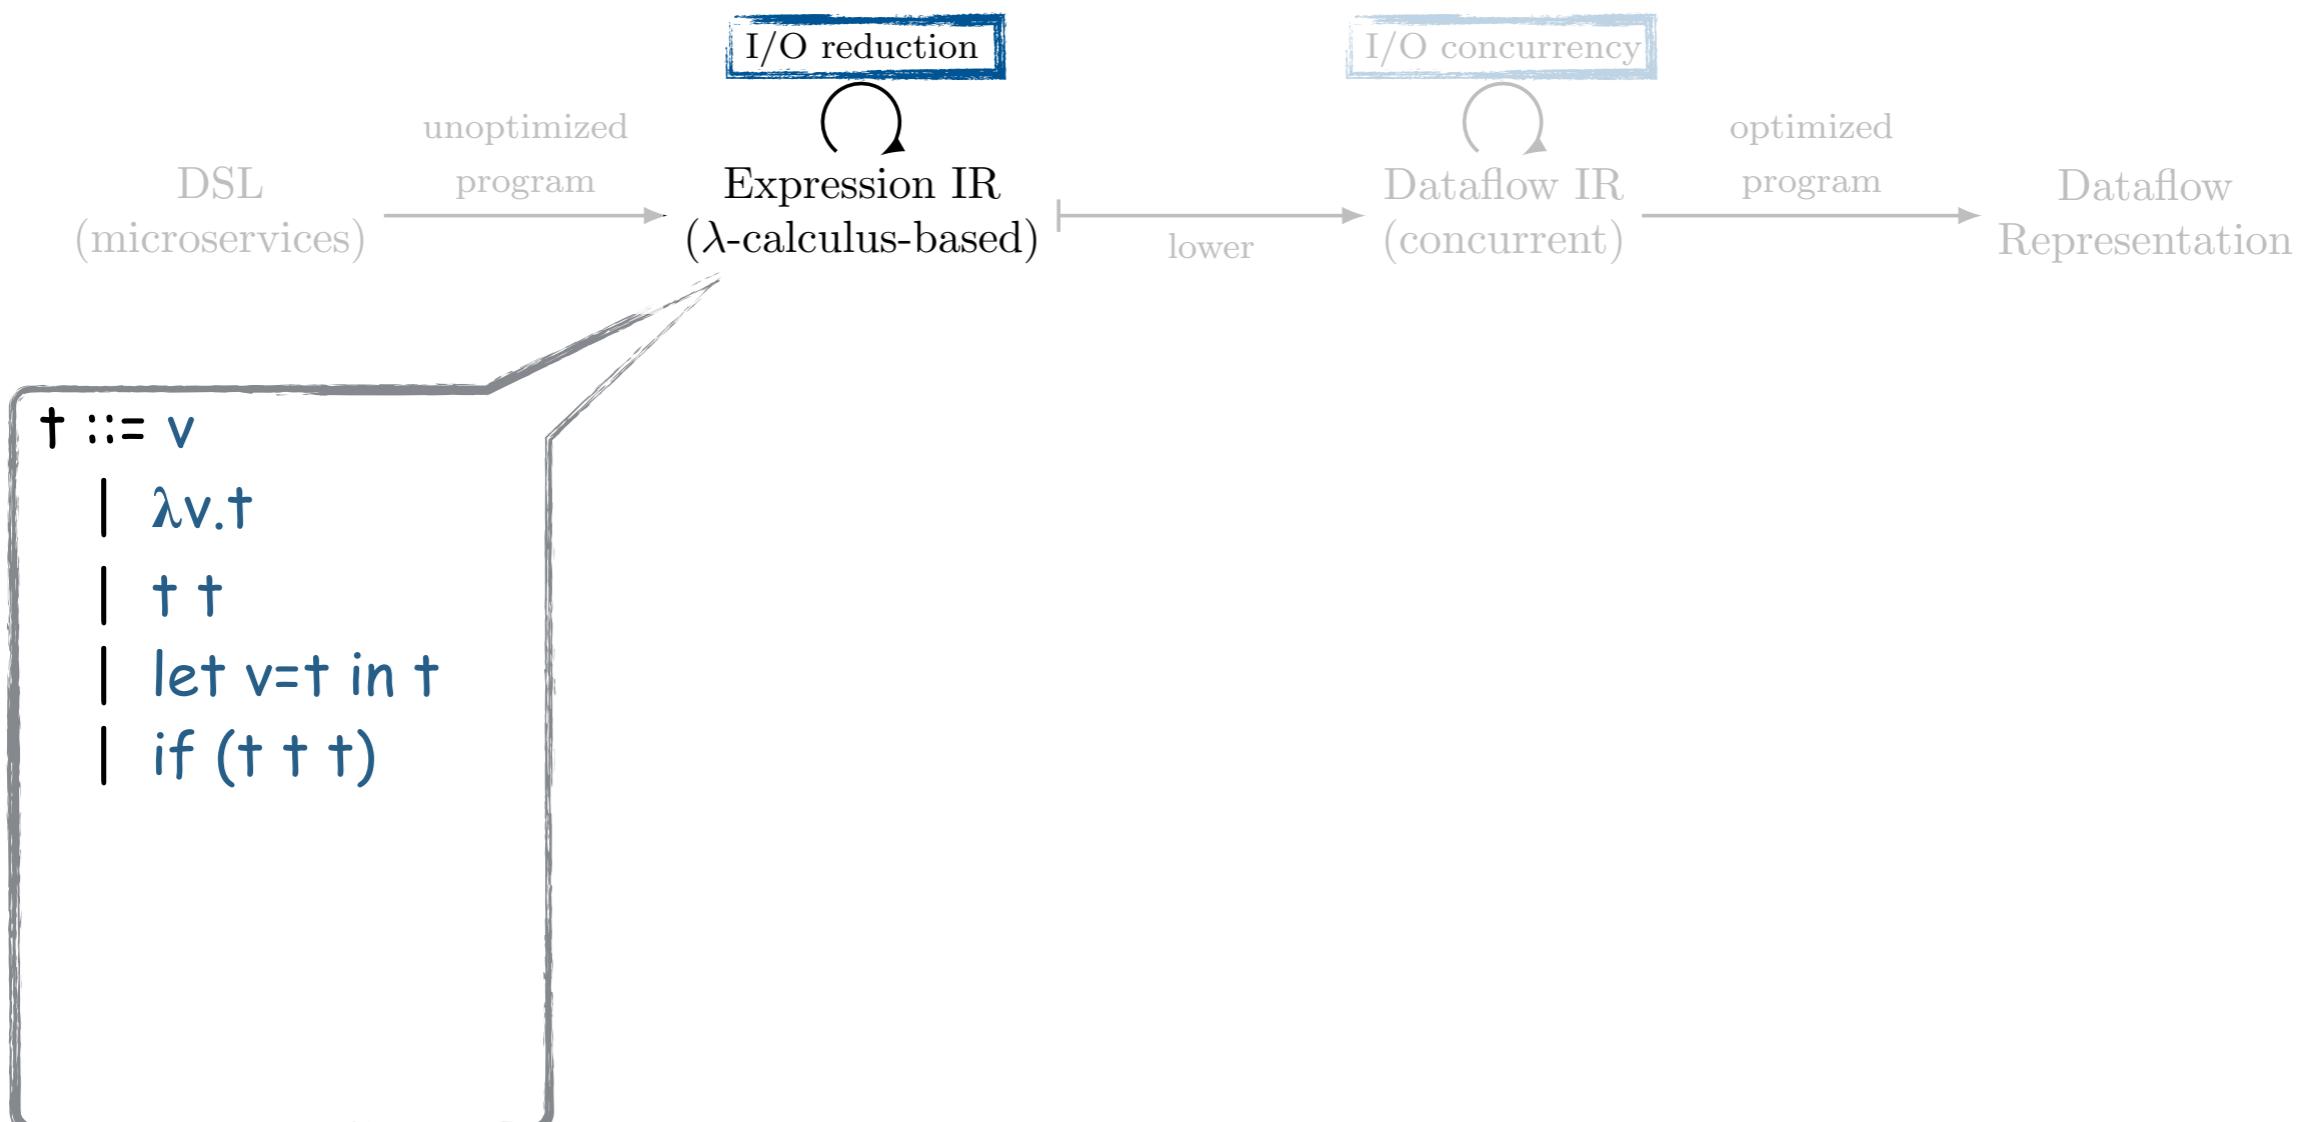


Ariola ZM, Felleisen M. The call-by-need lambda calculus. Journal of functional programming. 1997.

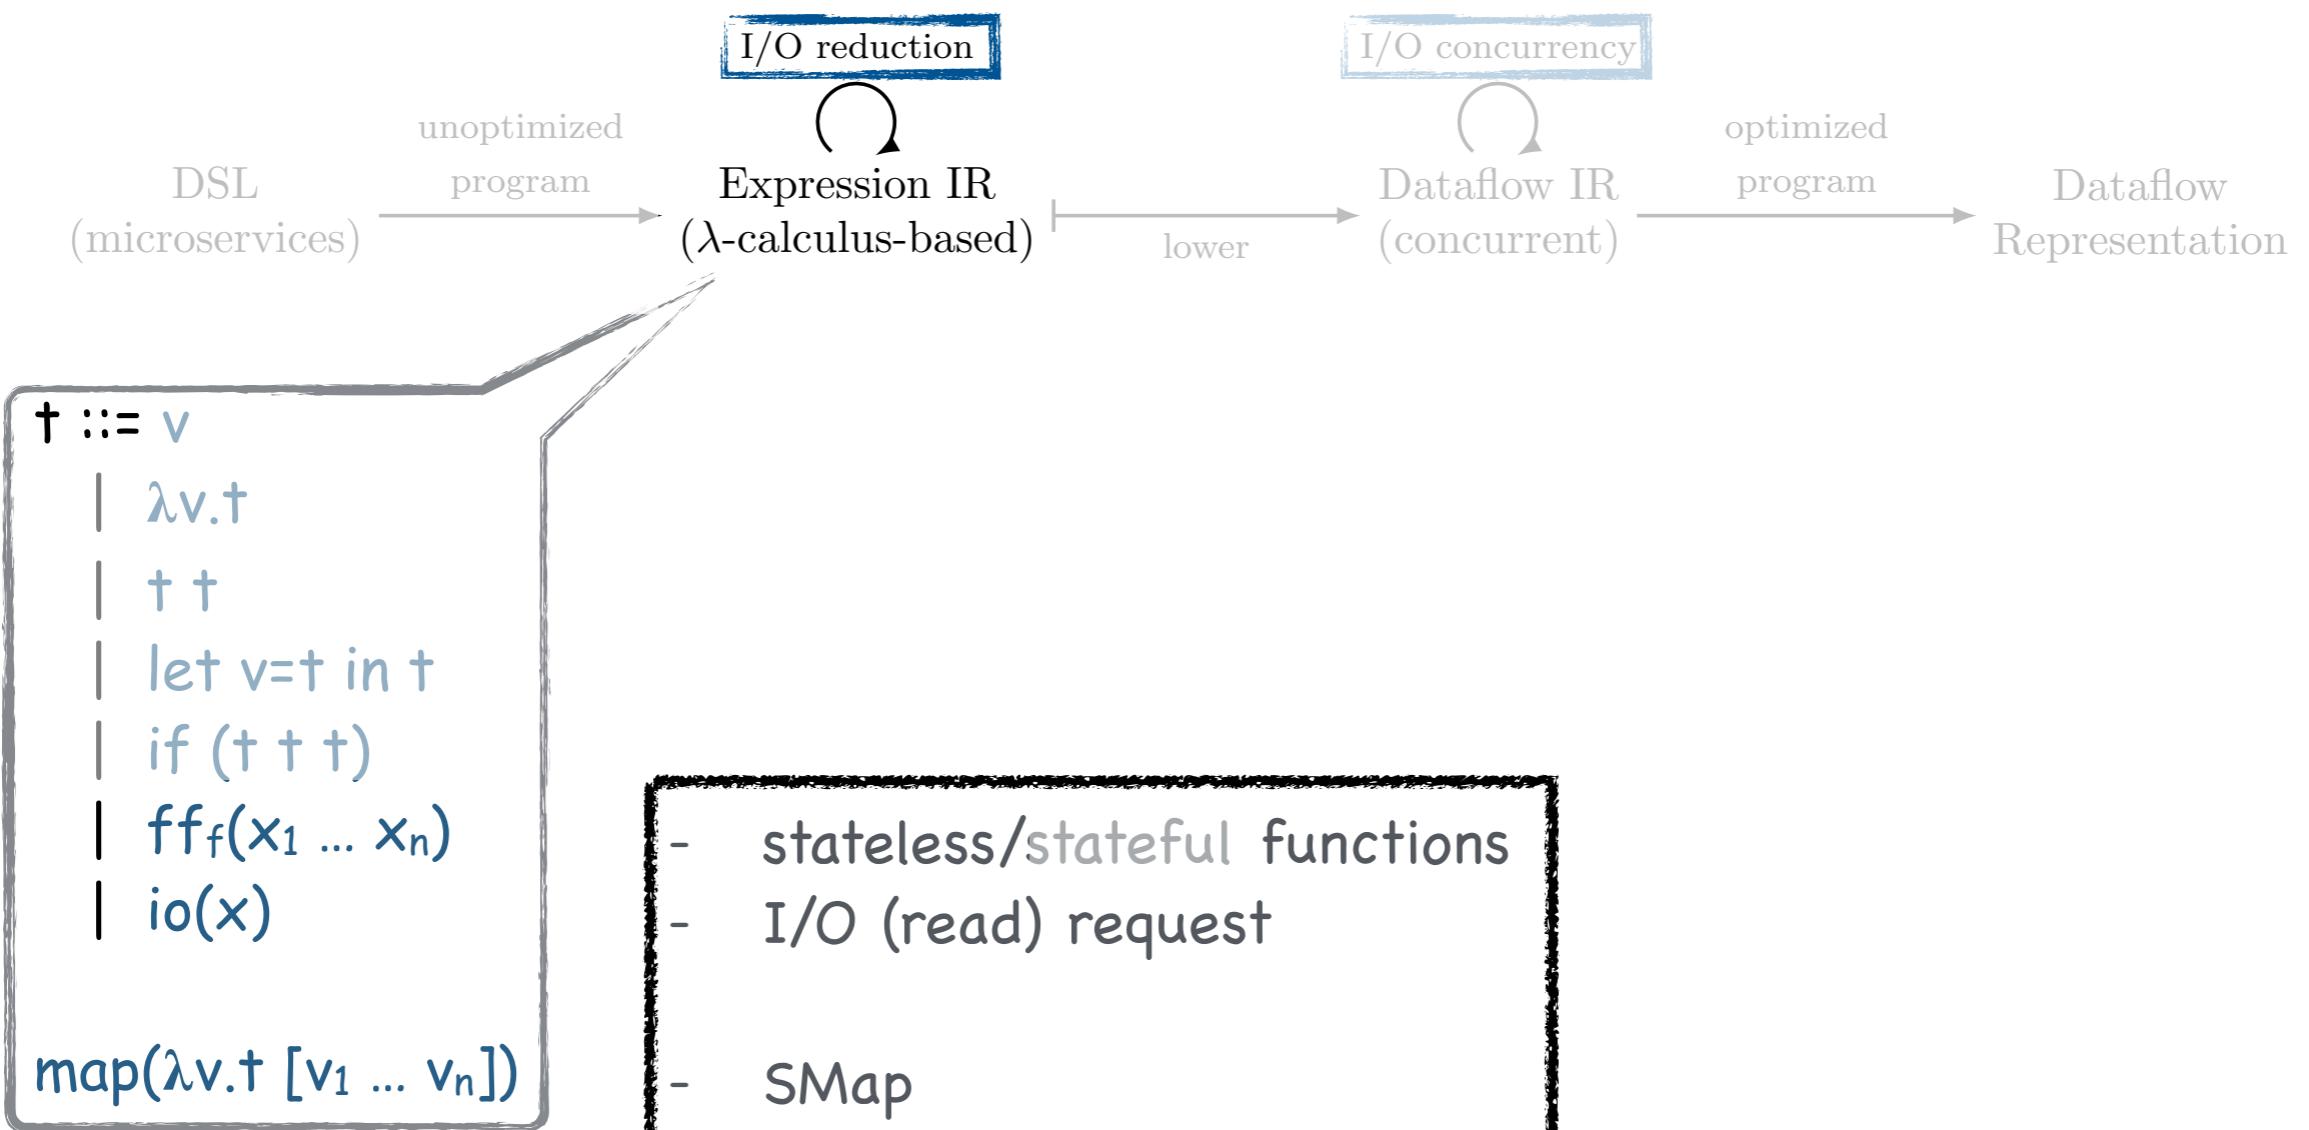
Maraist J, Odersky M, Wadler P. The call-by-need lambda calculus. Journal of functional programming. 1998.

Chang S, Felleisen M. The call-by-need lambda calculus, revisited. In European Symposium on Programming 2012. Springer.

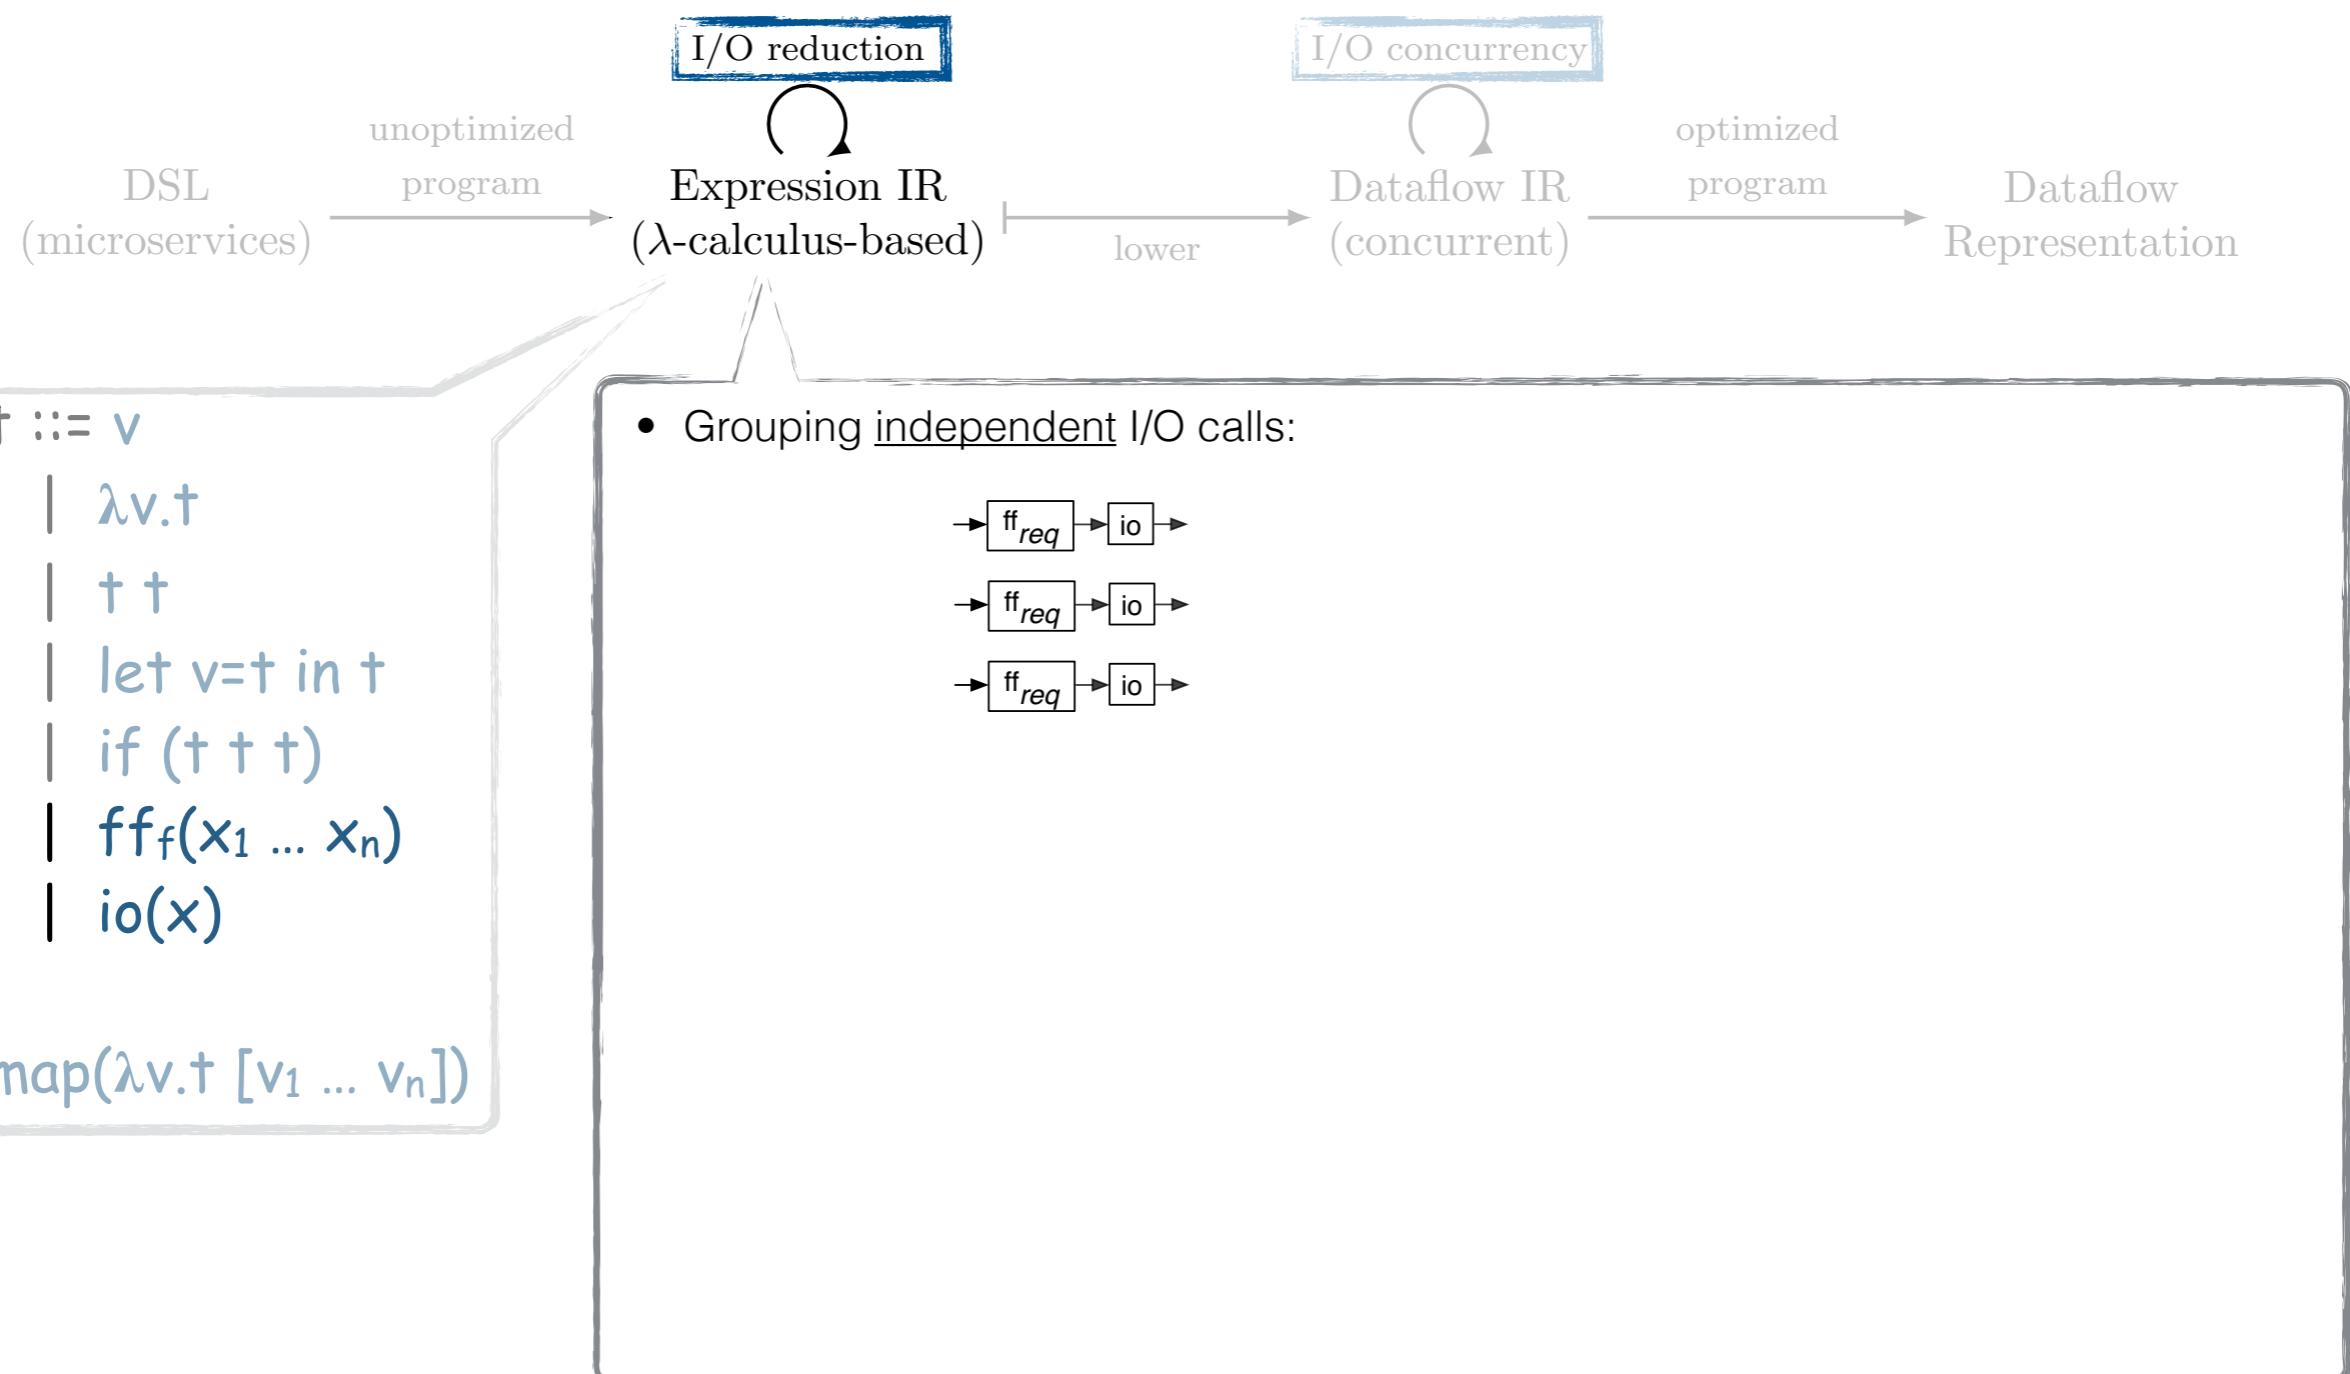
# Foundation: Call-by-need Lambda Calculus



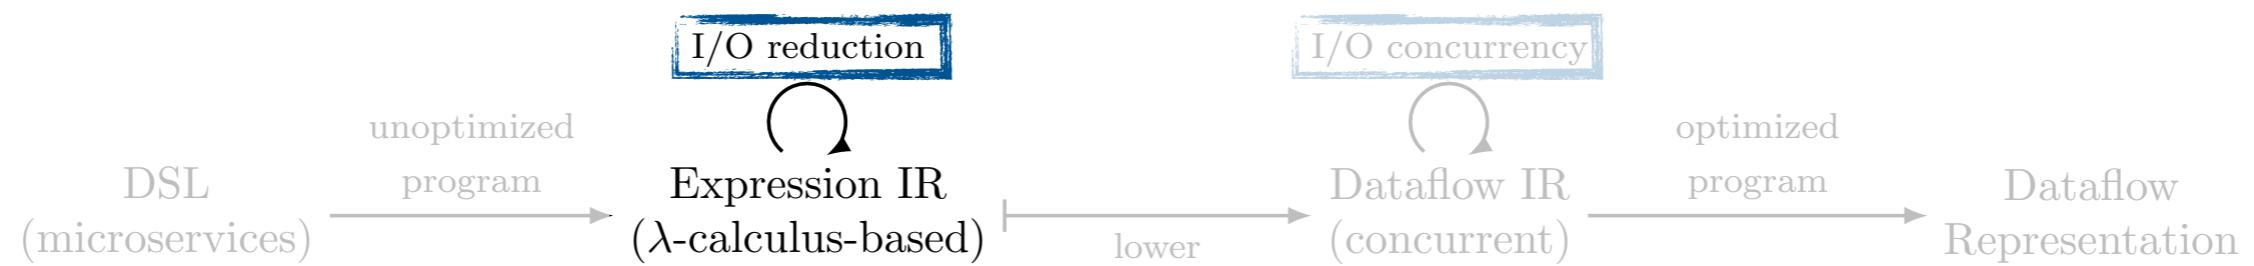
# Domain-specific Aspects



# I/O Transformations

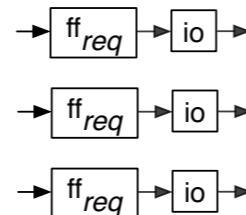


# I/O Transformations



$\dagger ::= v$   
 |  $\lambda v. \dagger$   
 |  $\dagger \dagger$   
 |  $\text{let } v = \dagger \text{ in } \dagger$   
 |  $\text{if } (\dagger \dagger \dagger)$   
 |  $\text{ff}_f(x_1 \dots x_n)$   
 |  $\text{io}(x)$   
  
 $\text{map}(\lambda v. \dagger [v_1 \dots v_n])$

- Grouping independent I/O calls:



Remember SMap

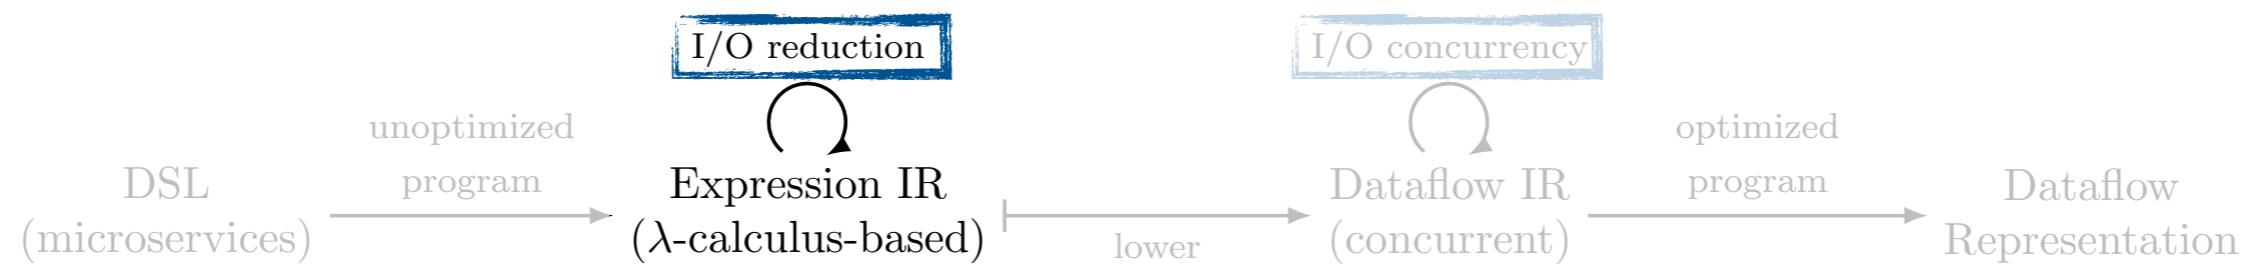
$$\begin{aligned} (s'_g, x_1) &= g(s_g, a_1) \\ (s''_g, x_2) &= g(s'_g, a_2) \\ (s'_h, b_1) &= h(s_h, x_1) \\ (s''_h, b_2) &= h(s'_h, x_2) \end{aligned}$$



$$\begin{aligned} (s'_g, x_1) &= g(s_g, a_1) \\ (s'_h, b_1) &= h(s_h, x_1) \\ (s''_g, x_2) &= g(s'_g, a_2) \\ (s''_h, b_2) &= h(s'_h, x_2) \end{aligned}$$

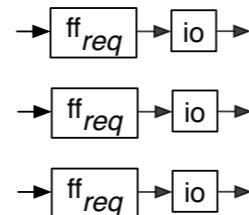
pipeline parallel

# I/O Transformations



$\dagger ::= v$   
 |  $\lambda v. \dagger$   
 |  $\dagger \dagger$   
 |  $\text{let } v = \dagger \text{ in } \dagger$   
 |  $\text{if } (\dagger \dagger \dagger)$   
 |  $\text{ff}_f(x_1 \dots x_n)$   
 |  $\text{io}(x)$   
  
 $\text{map}(\lambda v. \dagger [v_1 \dots v_n])$

- Grouping independent I/O calls:

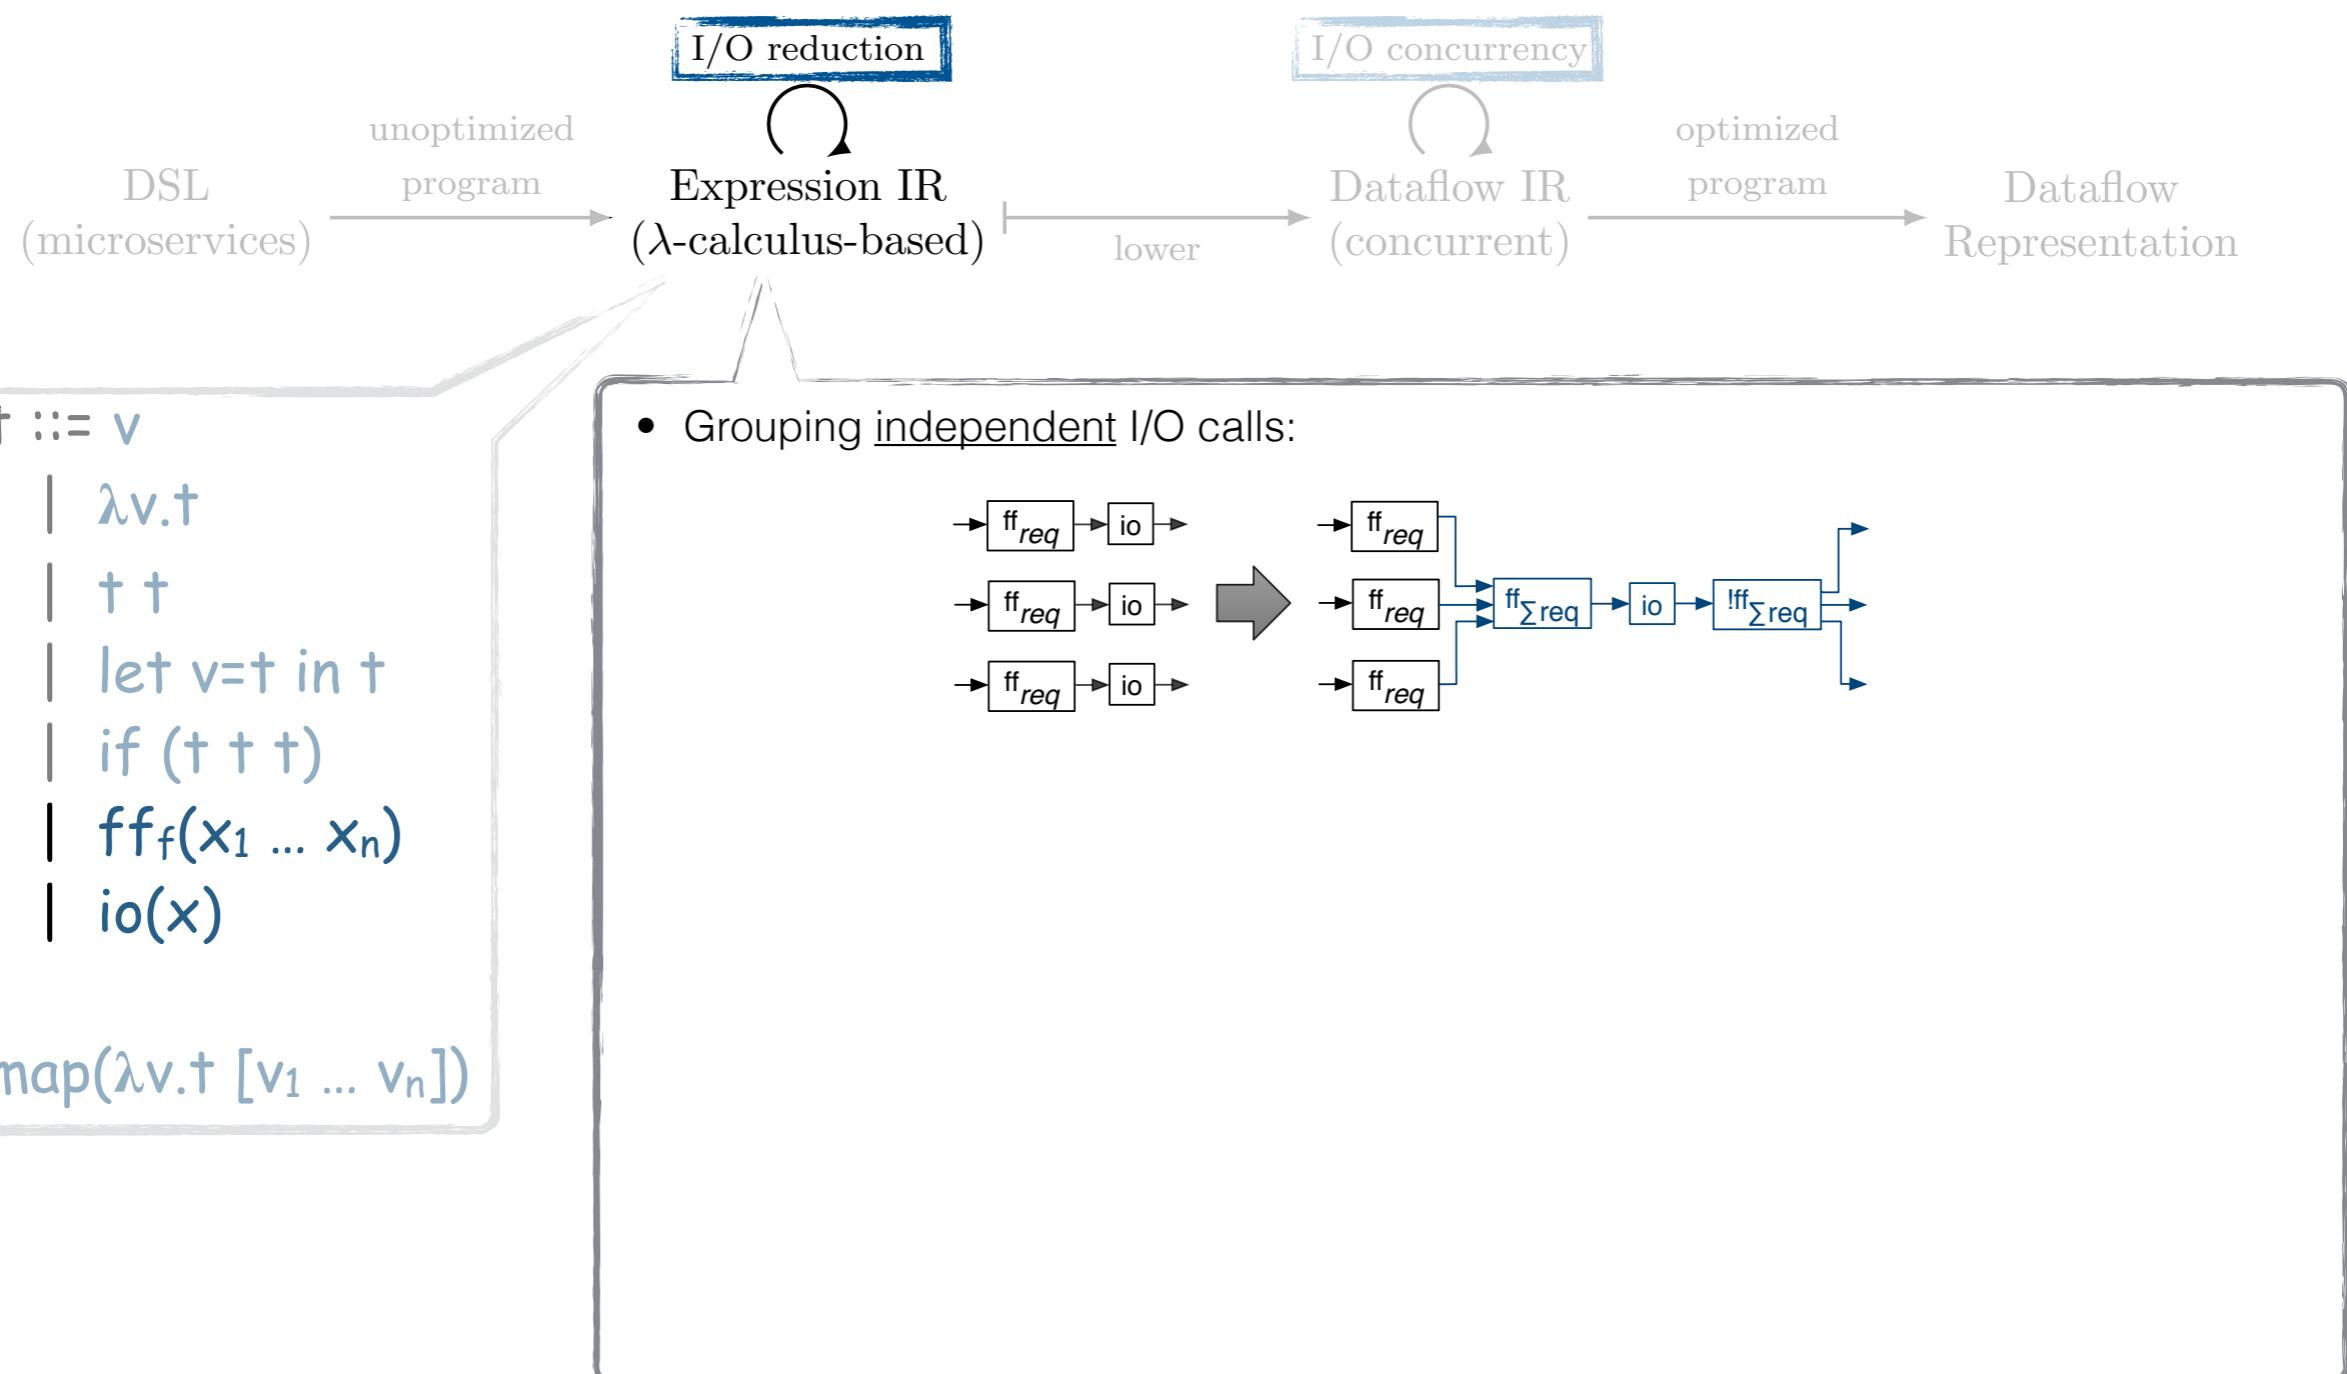


## Let-floating

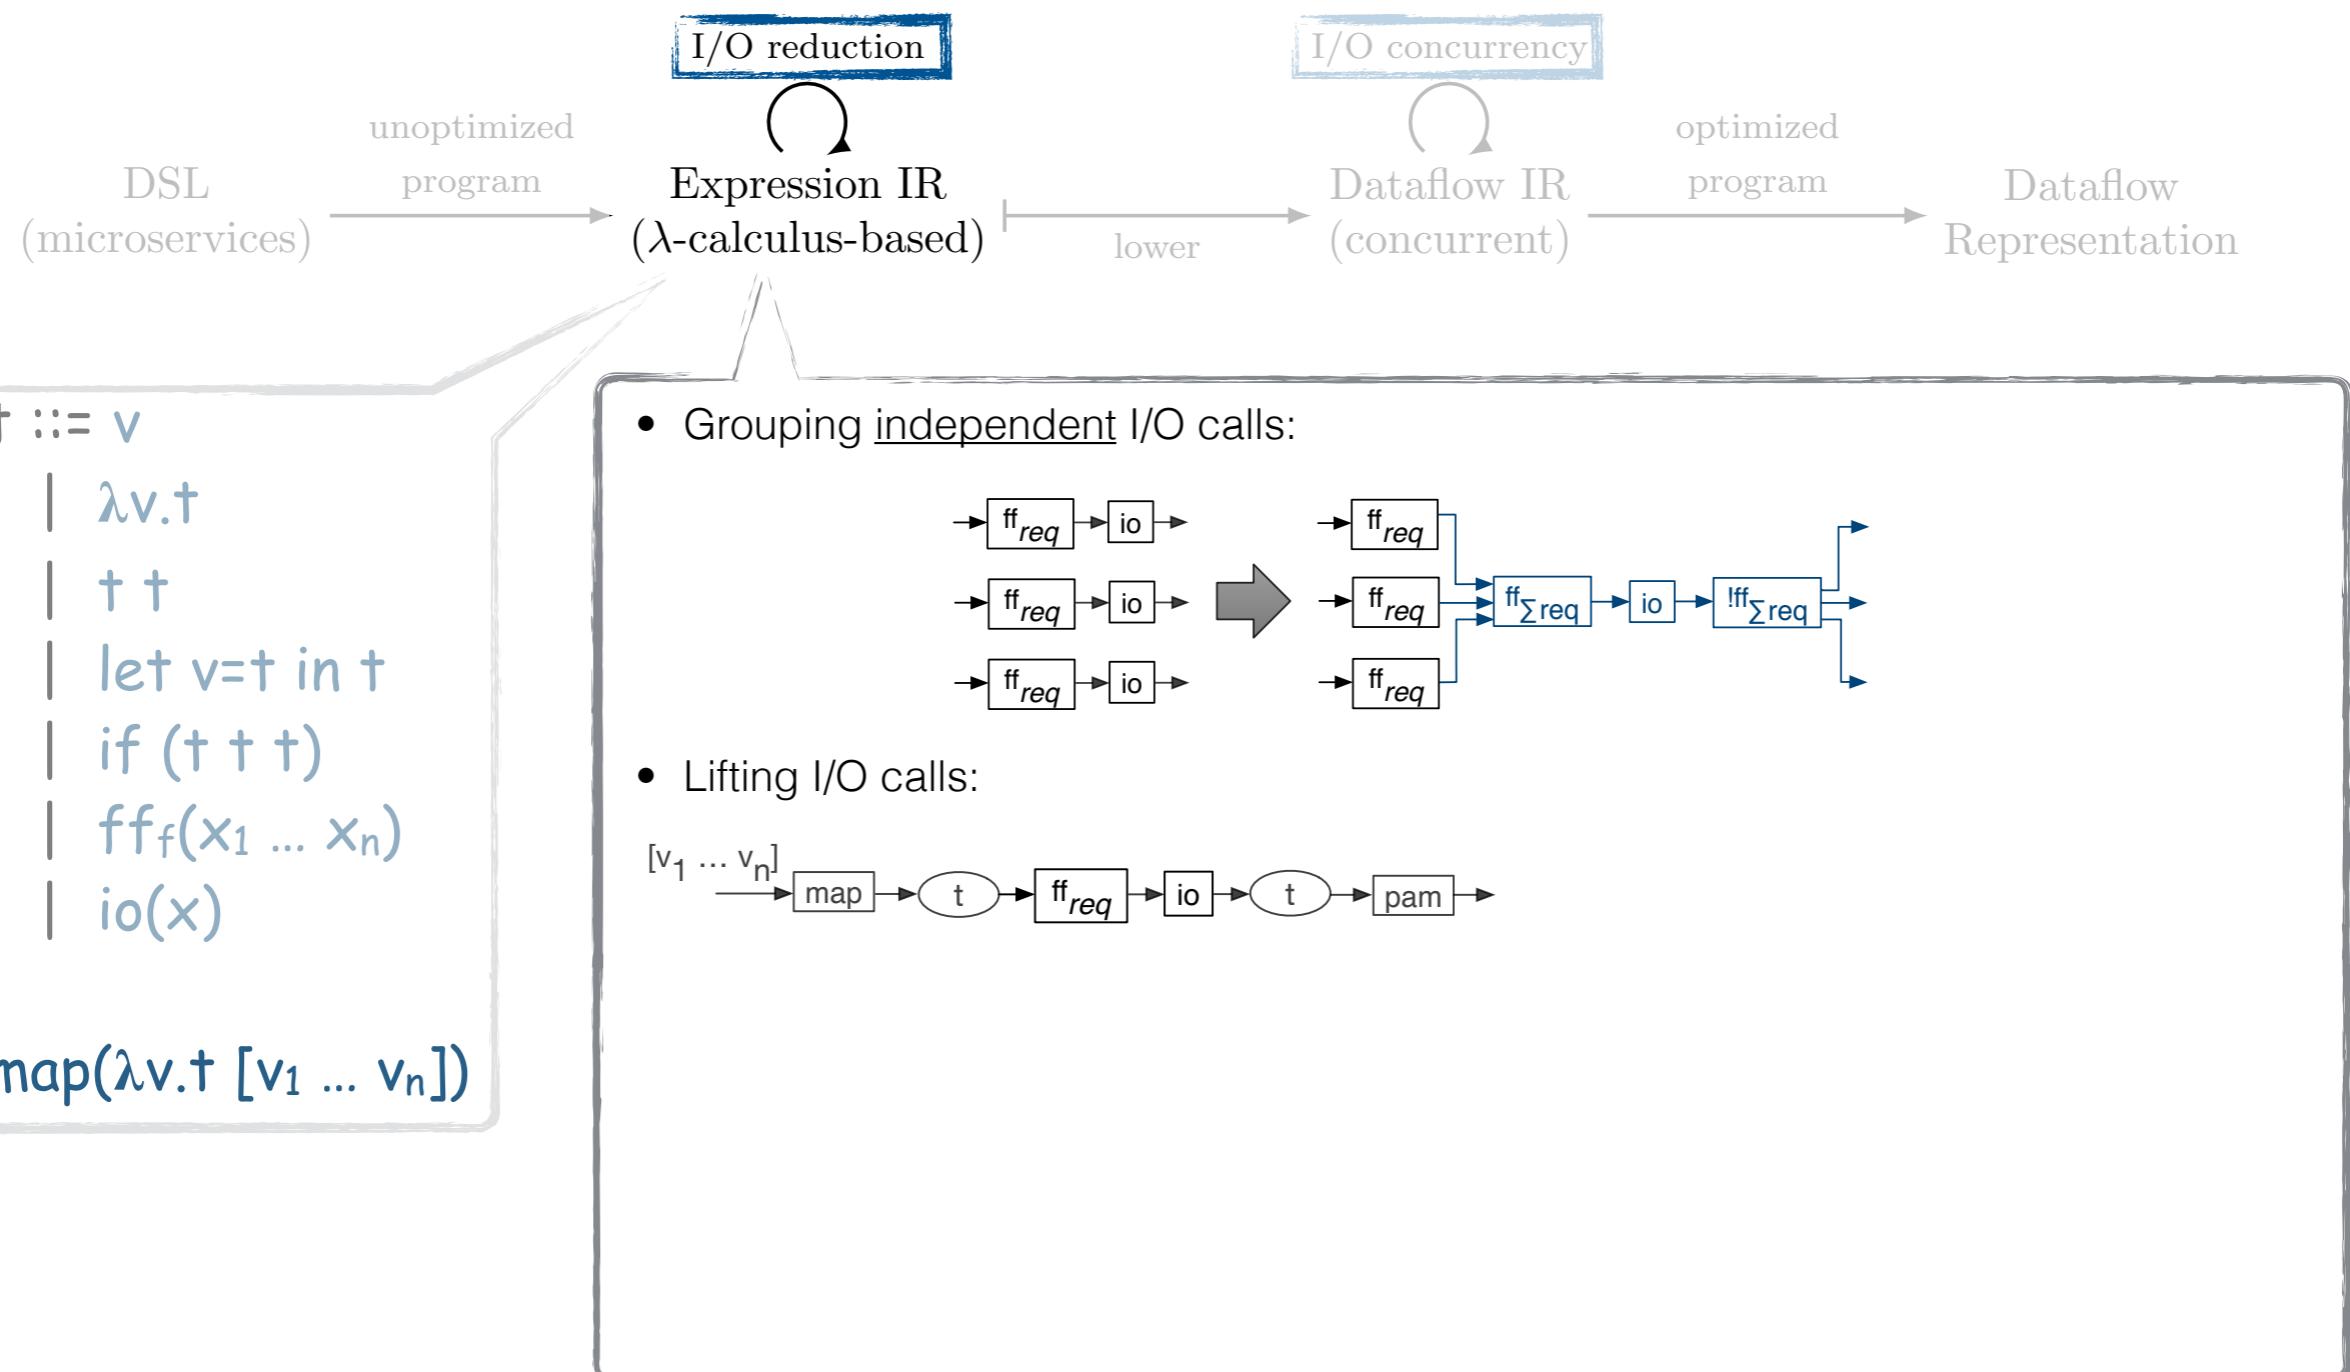
$$\begin{aligned}
 \text{smap}(f, [a_1, a_2]) = & \text{let } (s'_g, x_1) = g(s_g, a_1) \text{ in} \\
 & \text{let } (s''_g, x_2) = g(s'_g, a_2) \text{ in} \\
 & \text{let } (s'_h, b_1) = h(s_h, x_1) \text{ in} \\
 & \text{let } (s''_h, b_2) = h(s'_h, x_2) \text{ in} \\
 & [b_1, b_2]
 \end{aligned}
 \quad = \quad
 \begin{aligned}
 & \text{let } (s'_g, x_1) = g(s_g, a_1) \text{ in} \\
 & \cancel{\text{let } (s'_h, b_1) = h(s_h, x_1) \text{ in}} \\
 & \cancel{\text{let } (s''_g, x_2) = g(s'_g, a_2) \text{ in}} \\
 & \cancel{\text{let } (s''_h, b_2) = h(s'_h, x_2) \text{ in}} \\
 & [b_1, b_2]
 \end{aligned}$$

pipeline parallel

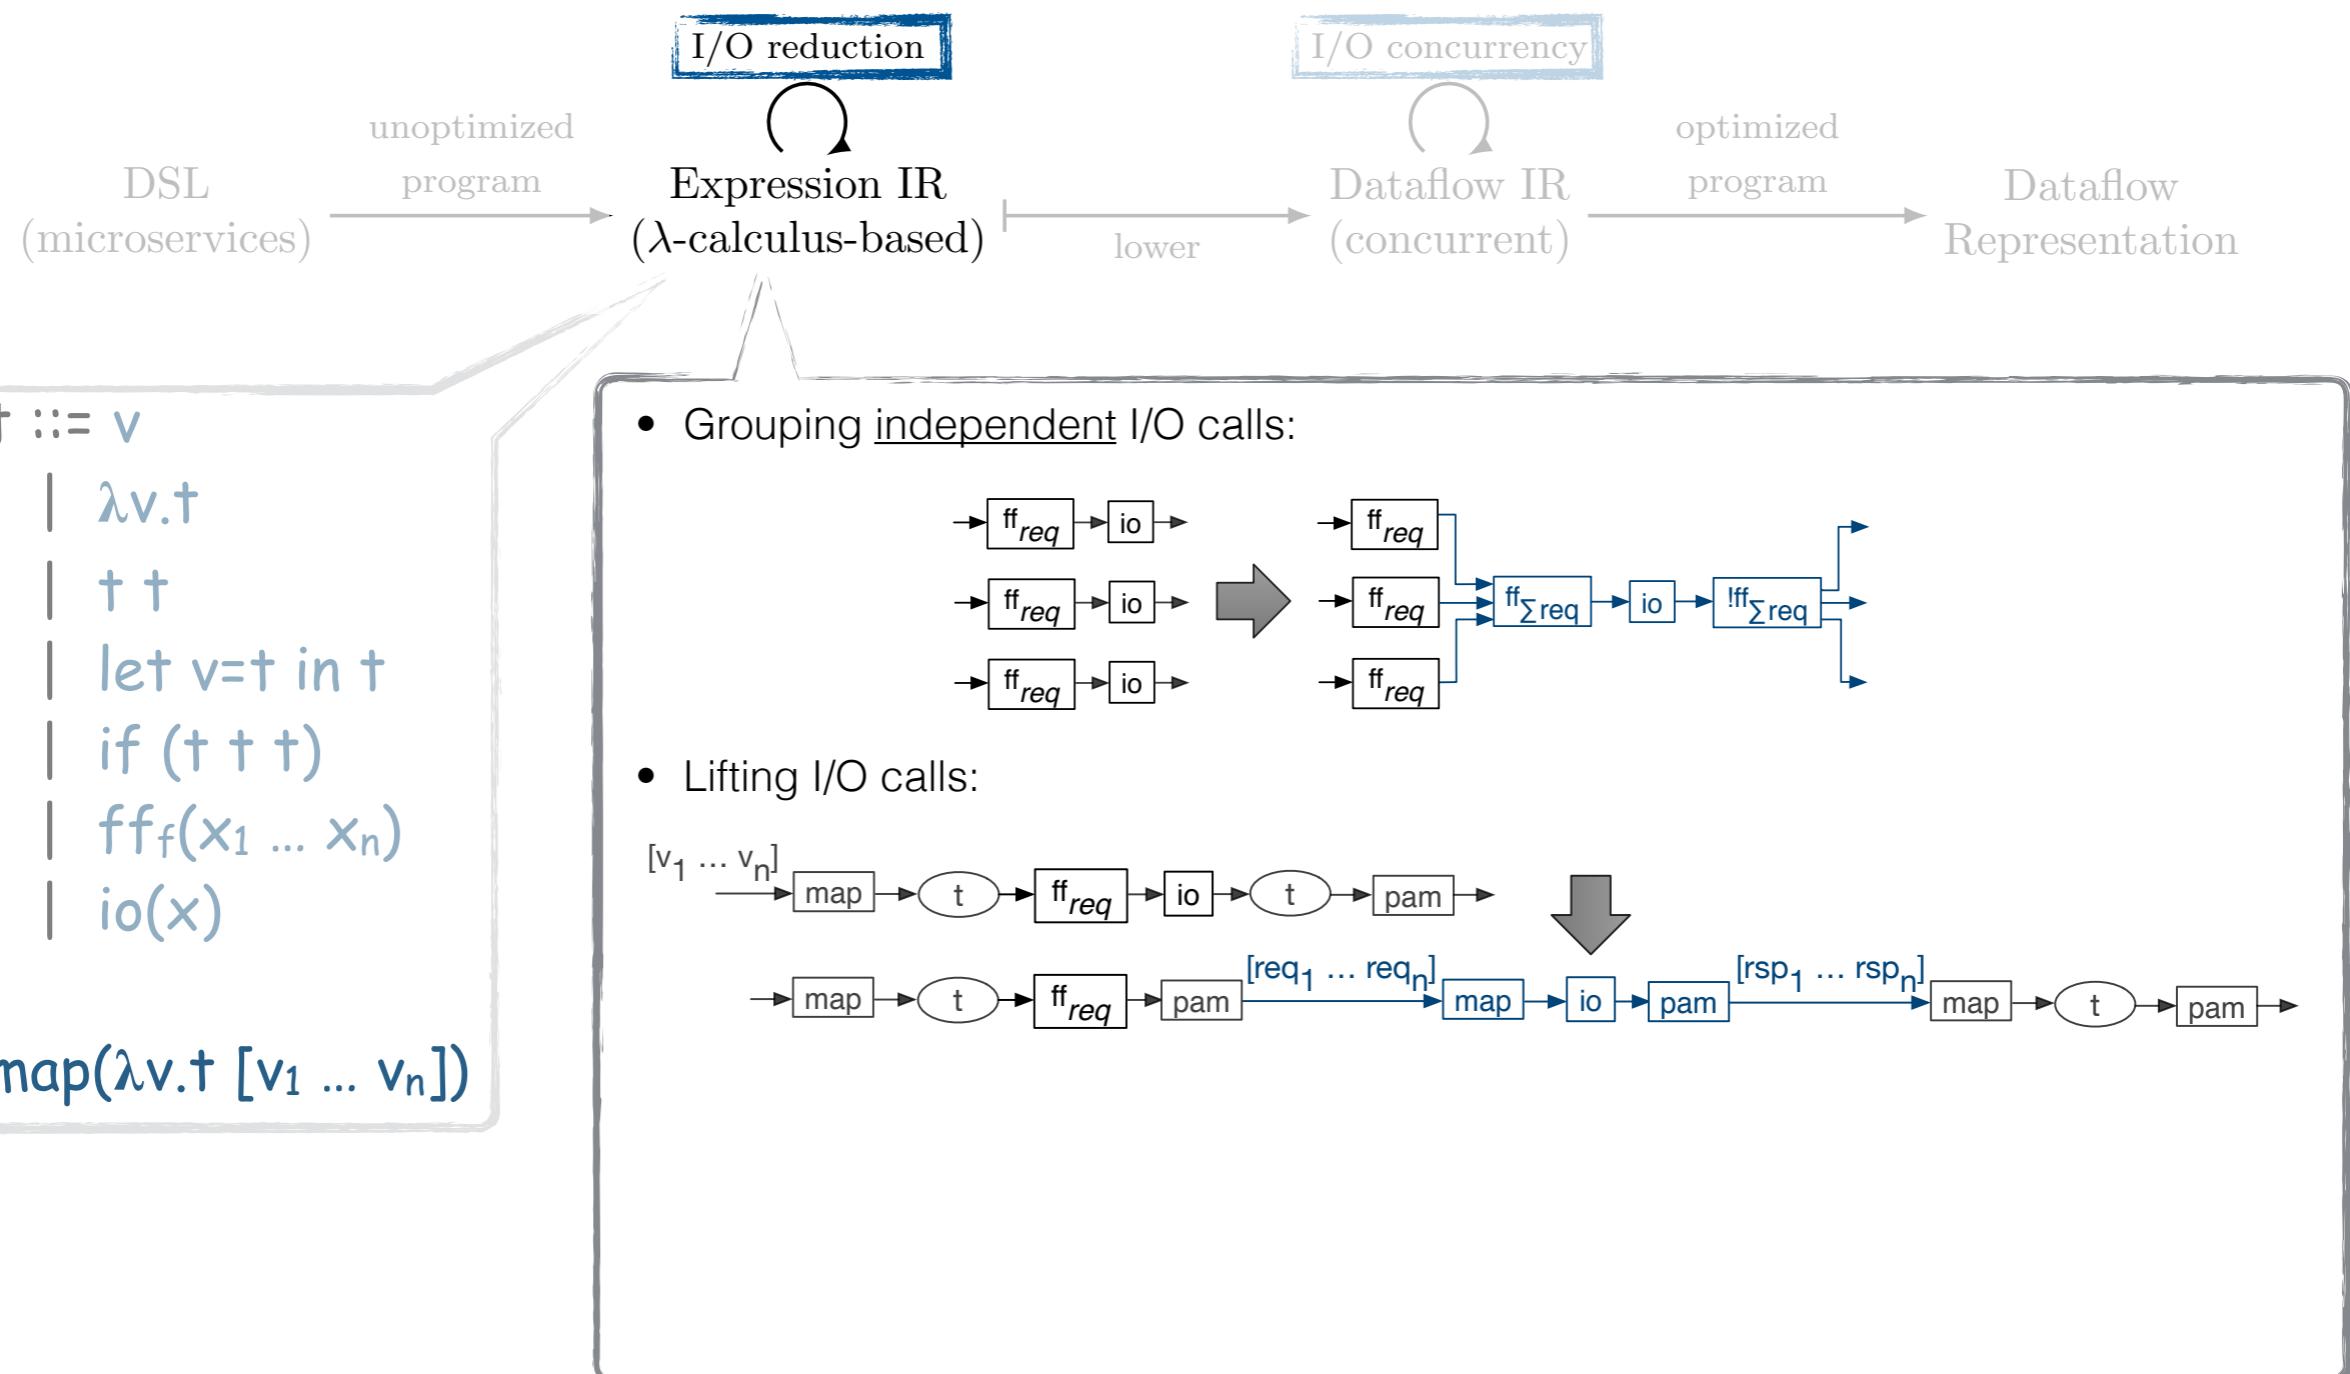
# I/O Transformations



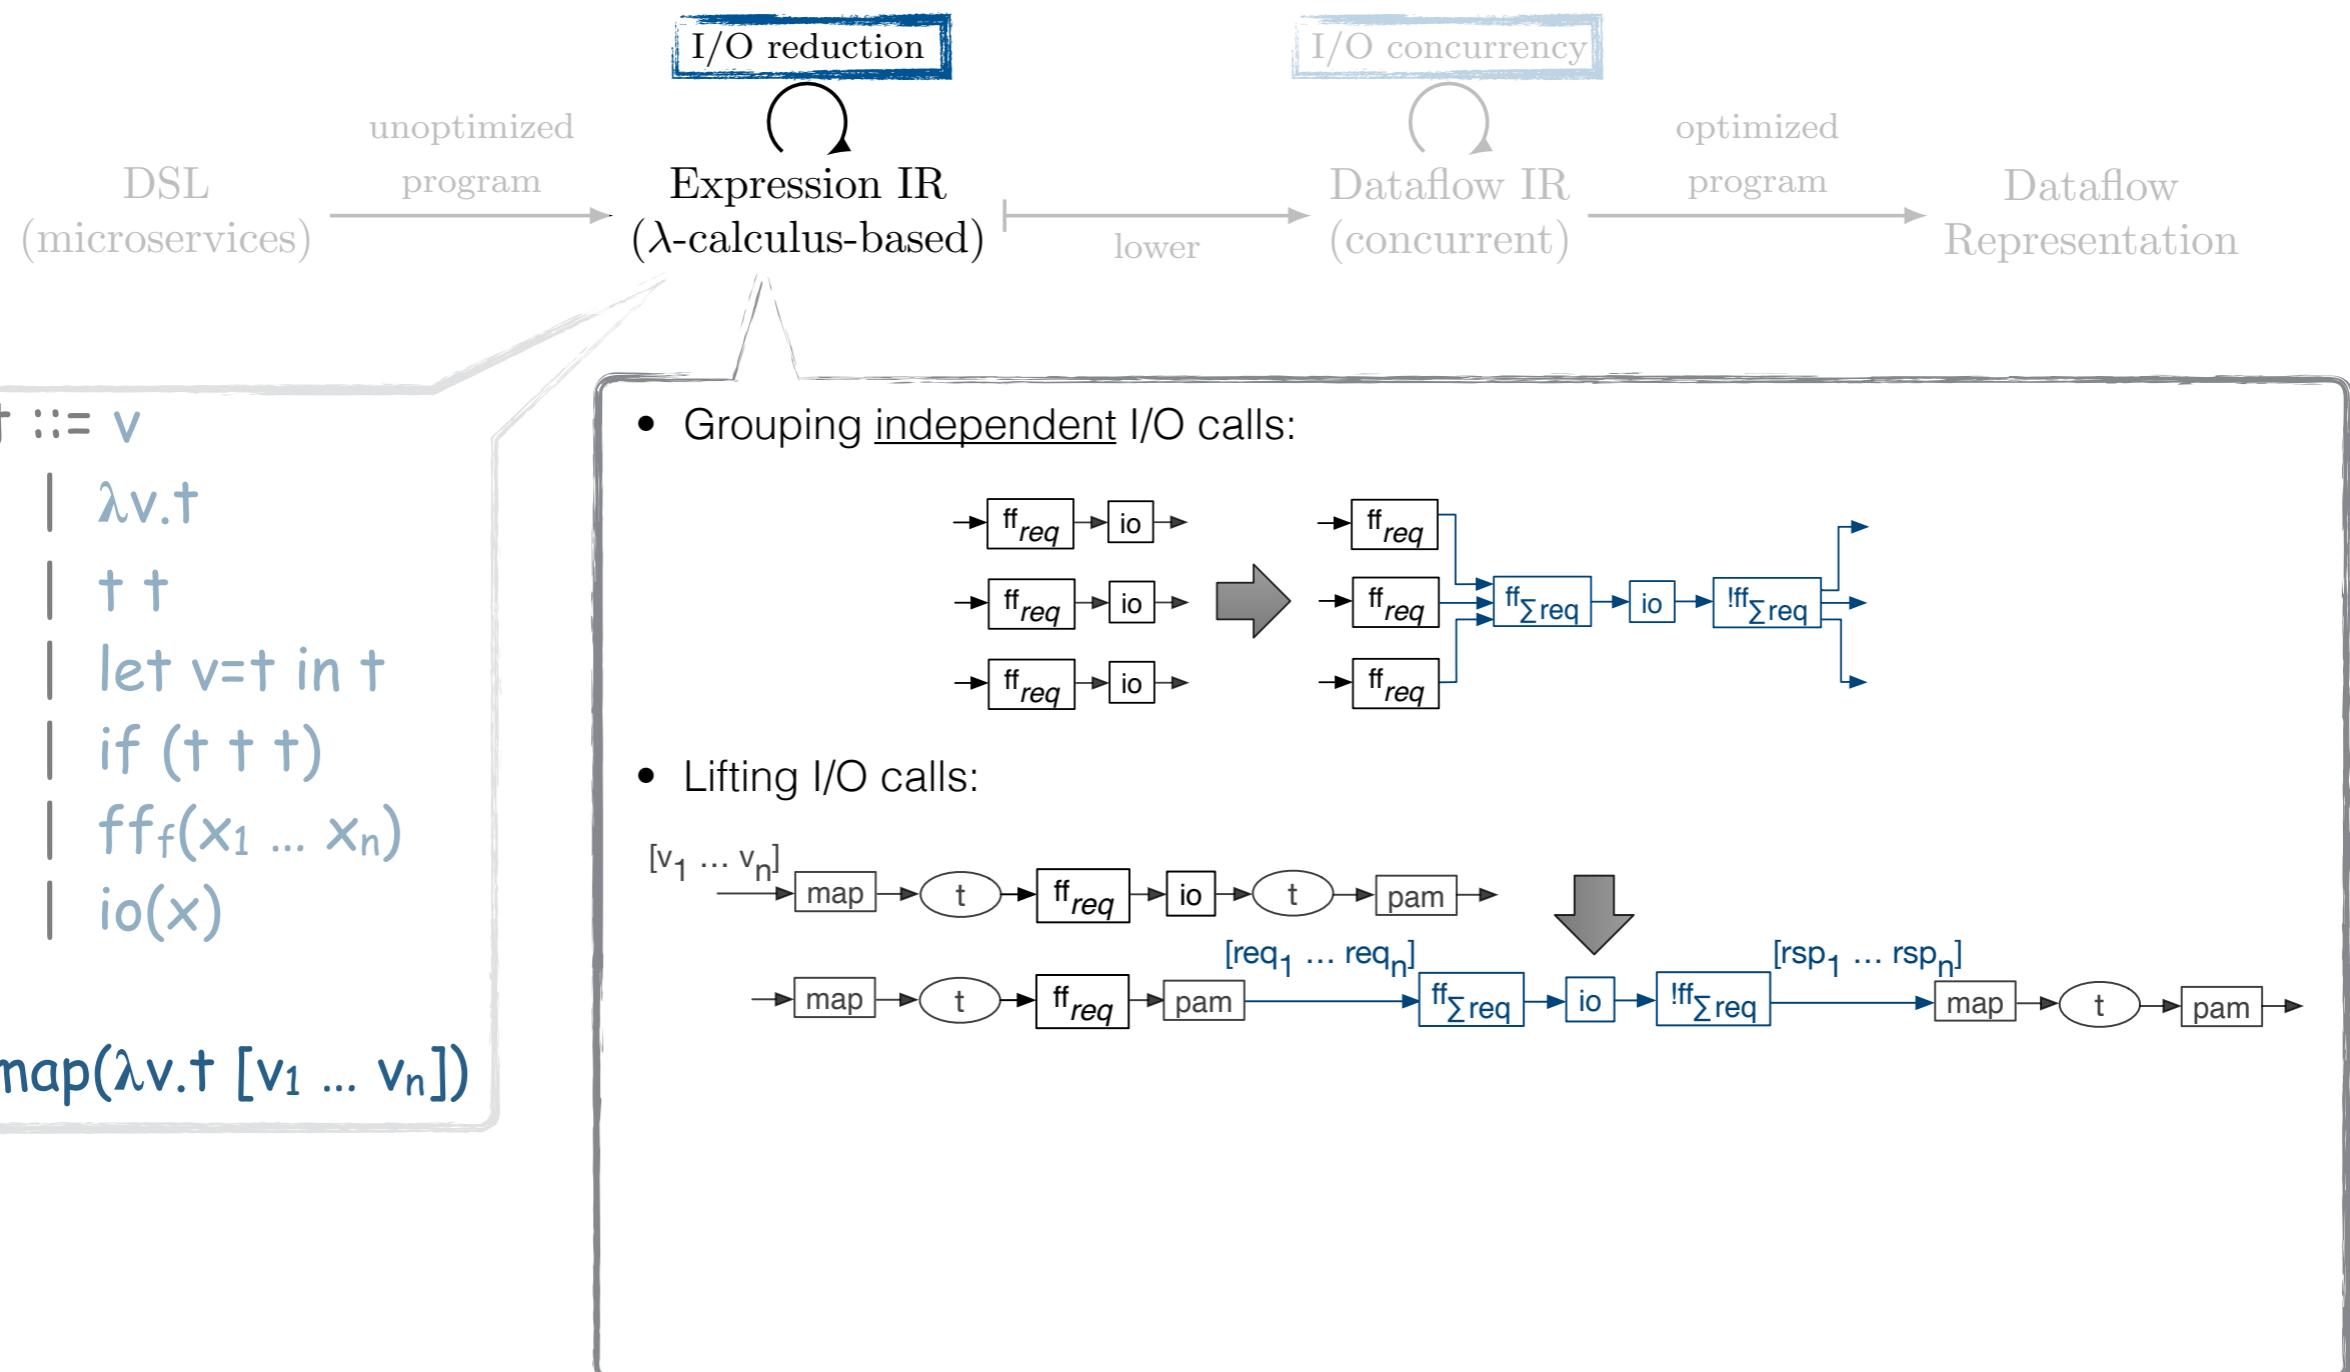
# I/O Transformations



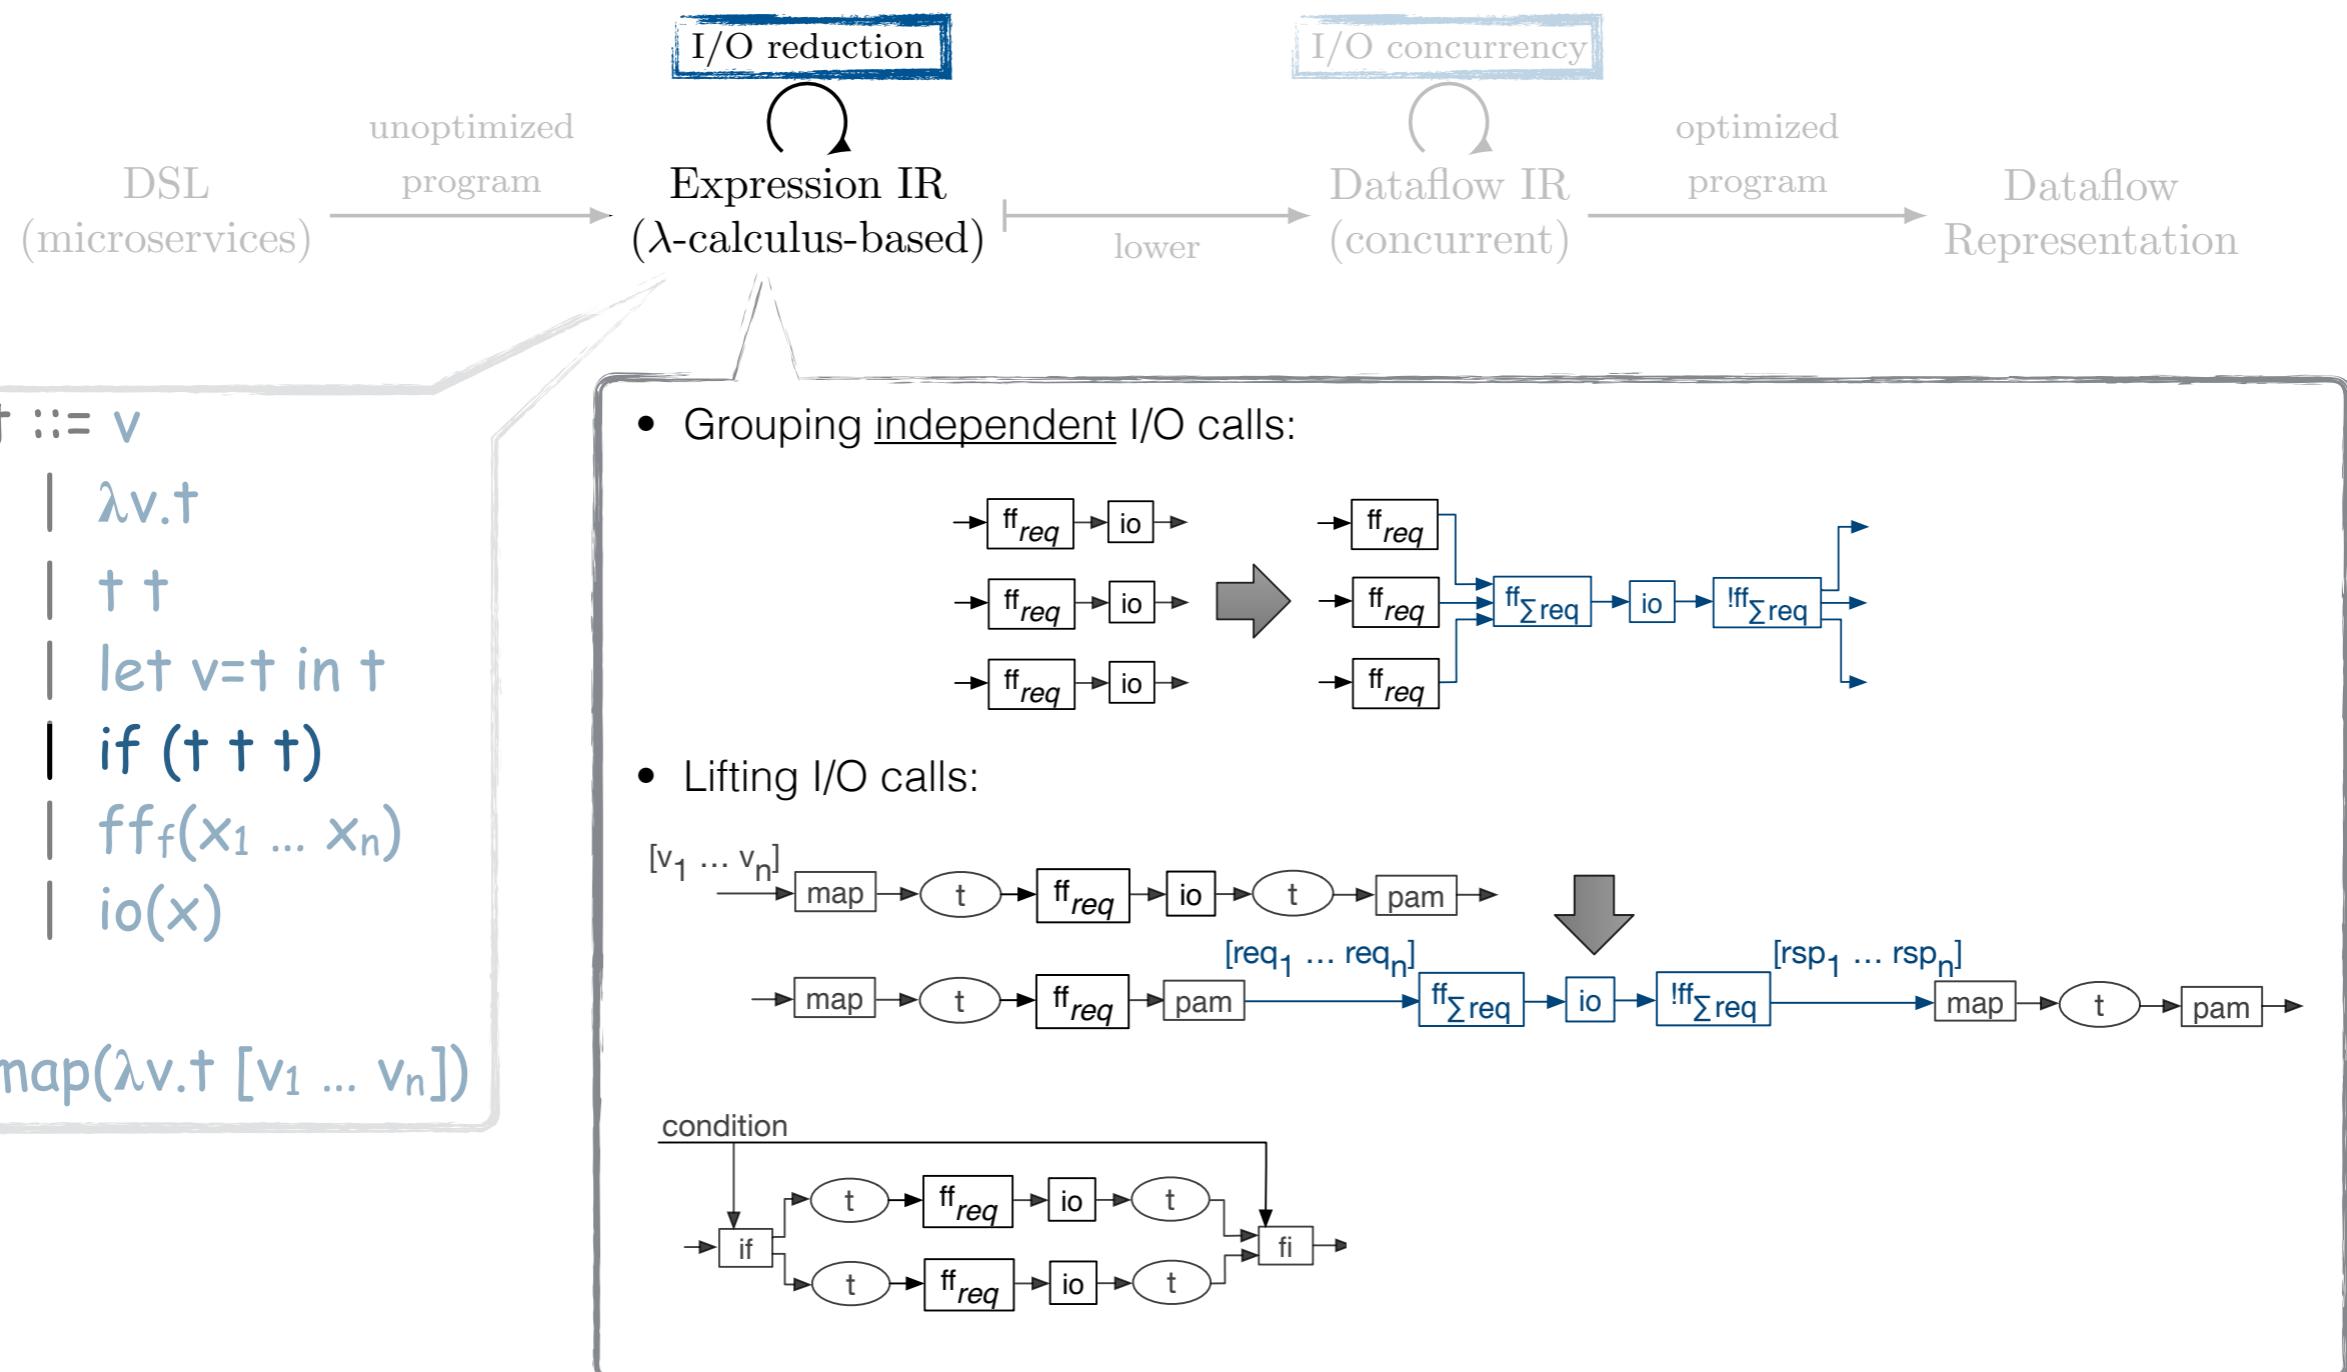
# I/O Transformations



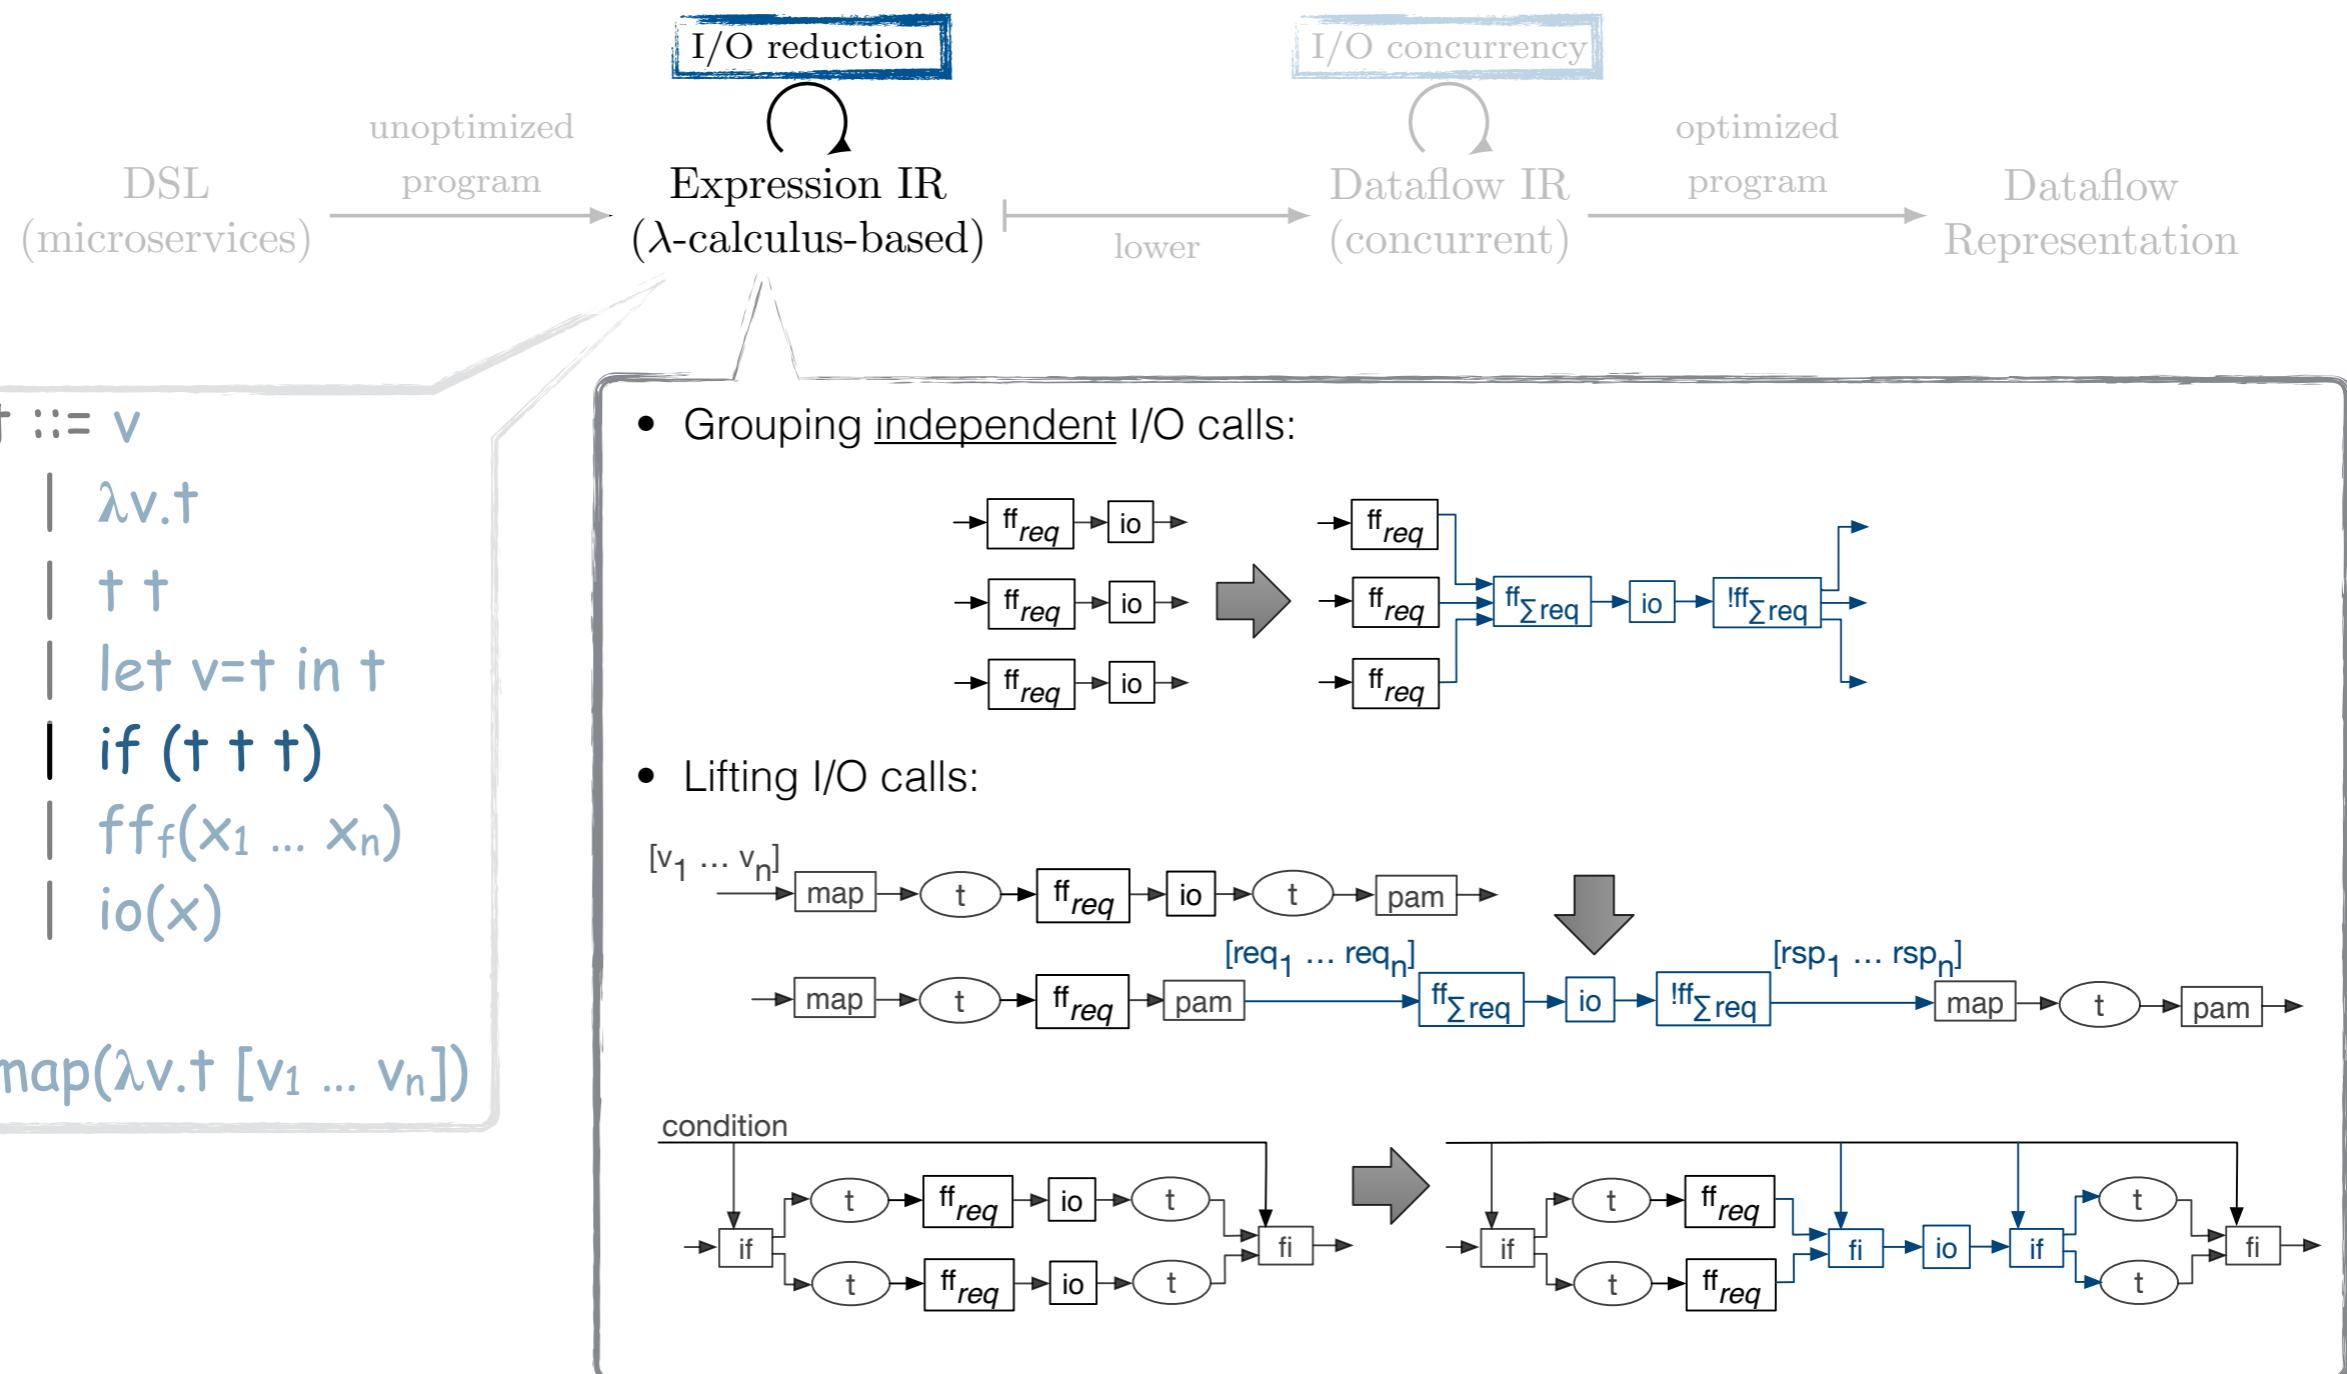
# I/O Transformations



# I/O Transformations

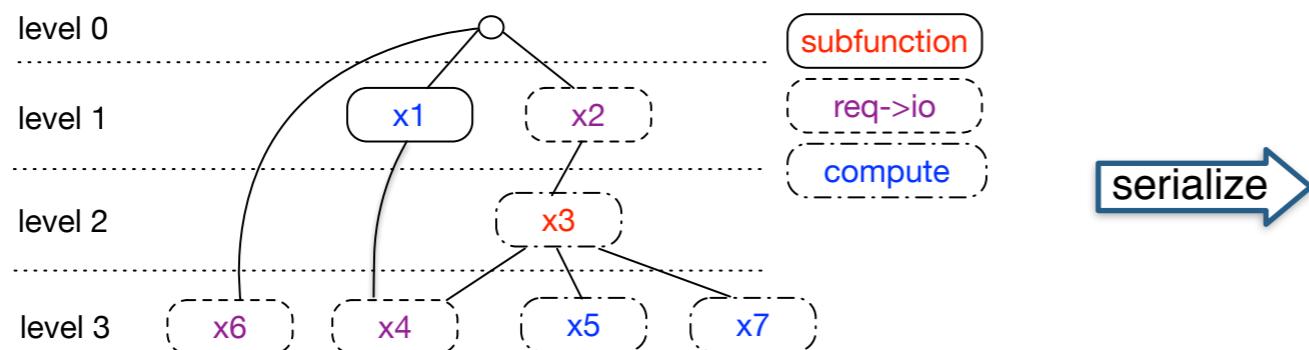


# I/O Transformations



# Evaluation

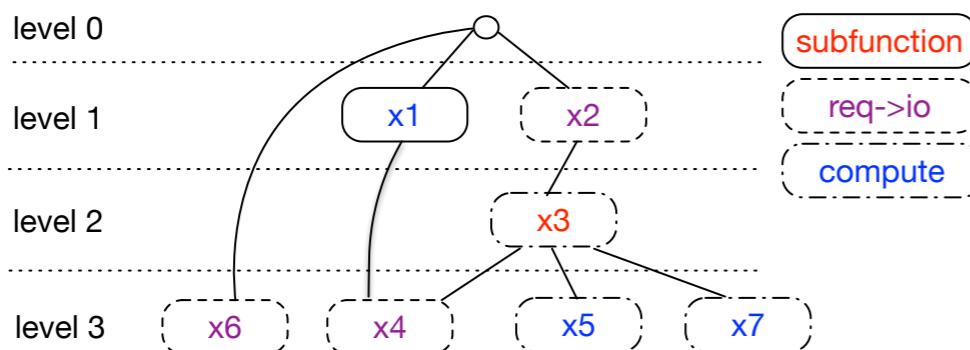
Level-Graphs:



Programs:

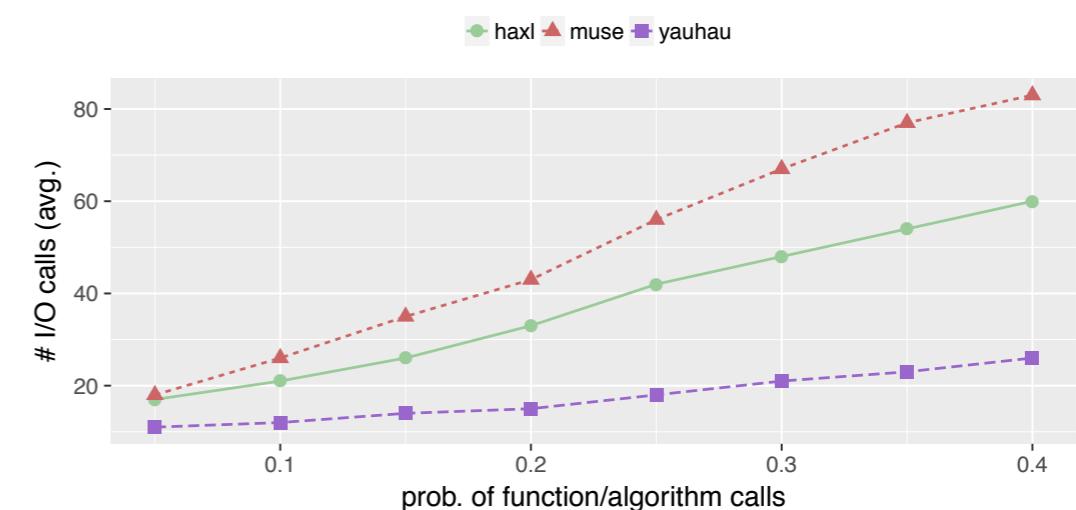
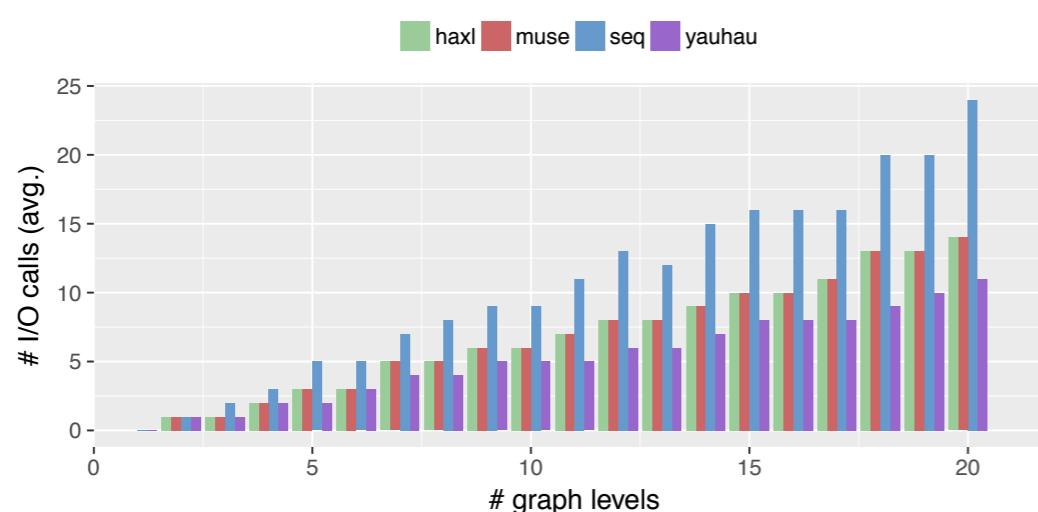
# Evaluation

Level-Graphs:



Programs:

- Haxl
- Muse
- Seq
- Yauhau



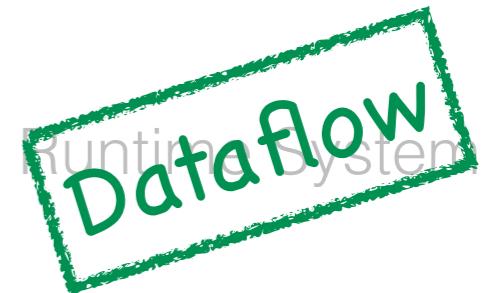
Simon Marlow, Louis Brandy, Jonathan Coens, and Jon Purdy. 2014. There is no fork: an abstraction for efficient, concurrent, and concise data access. ICFP '14.

Alexey Kachayev. 2015. Reinventing Haxl: Efficient, Concurrent and Concise Data Access. Presentation at EuroClojure 2015.

Programming Model/  
Language

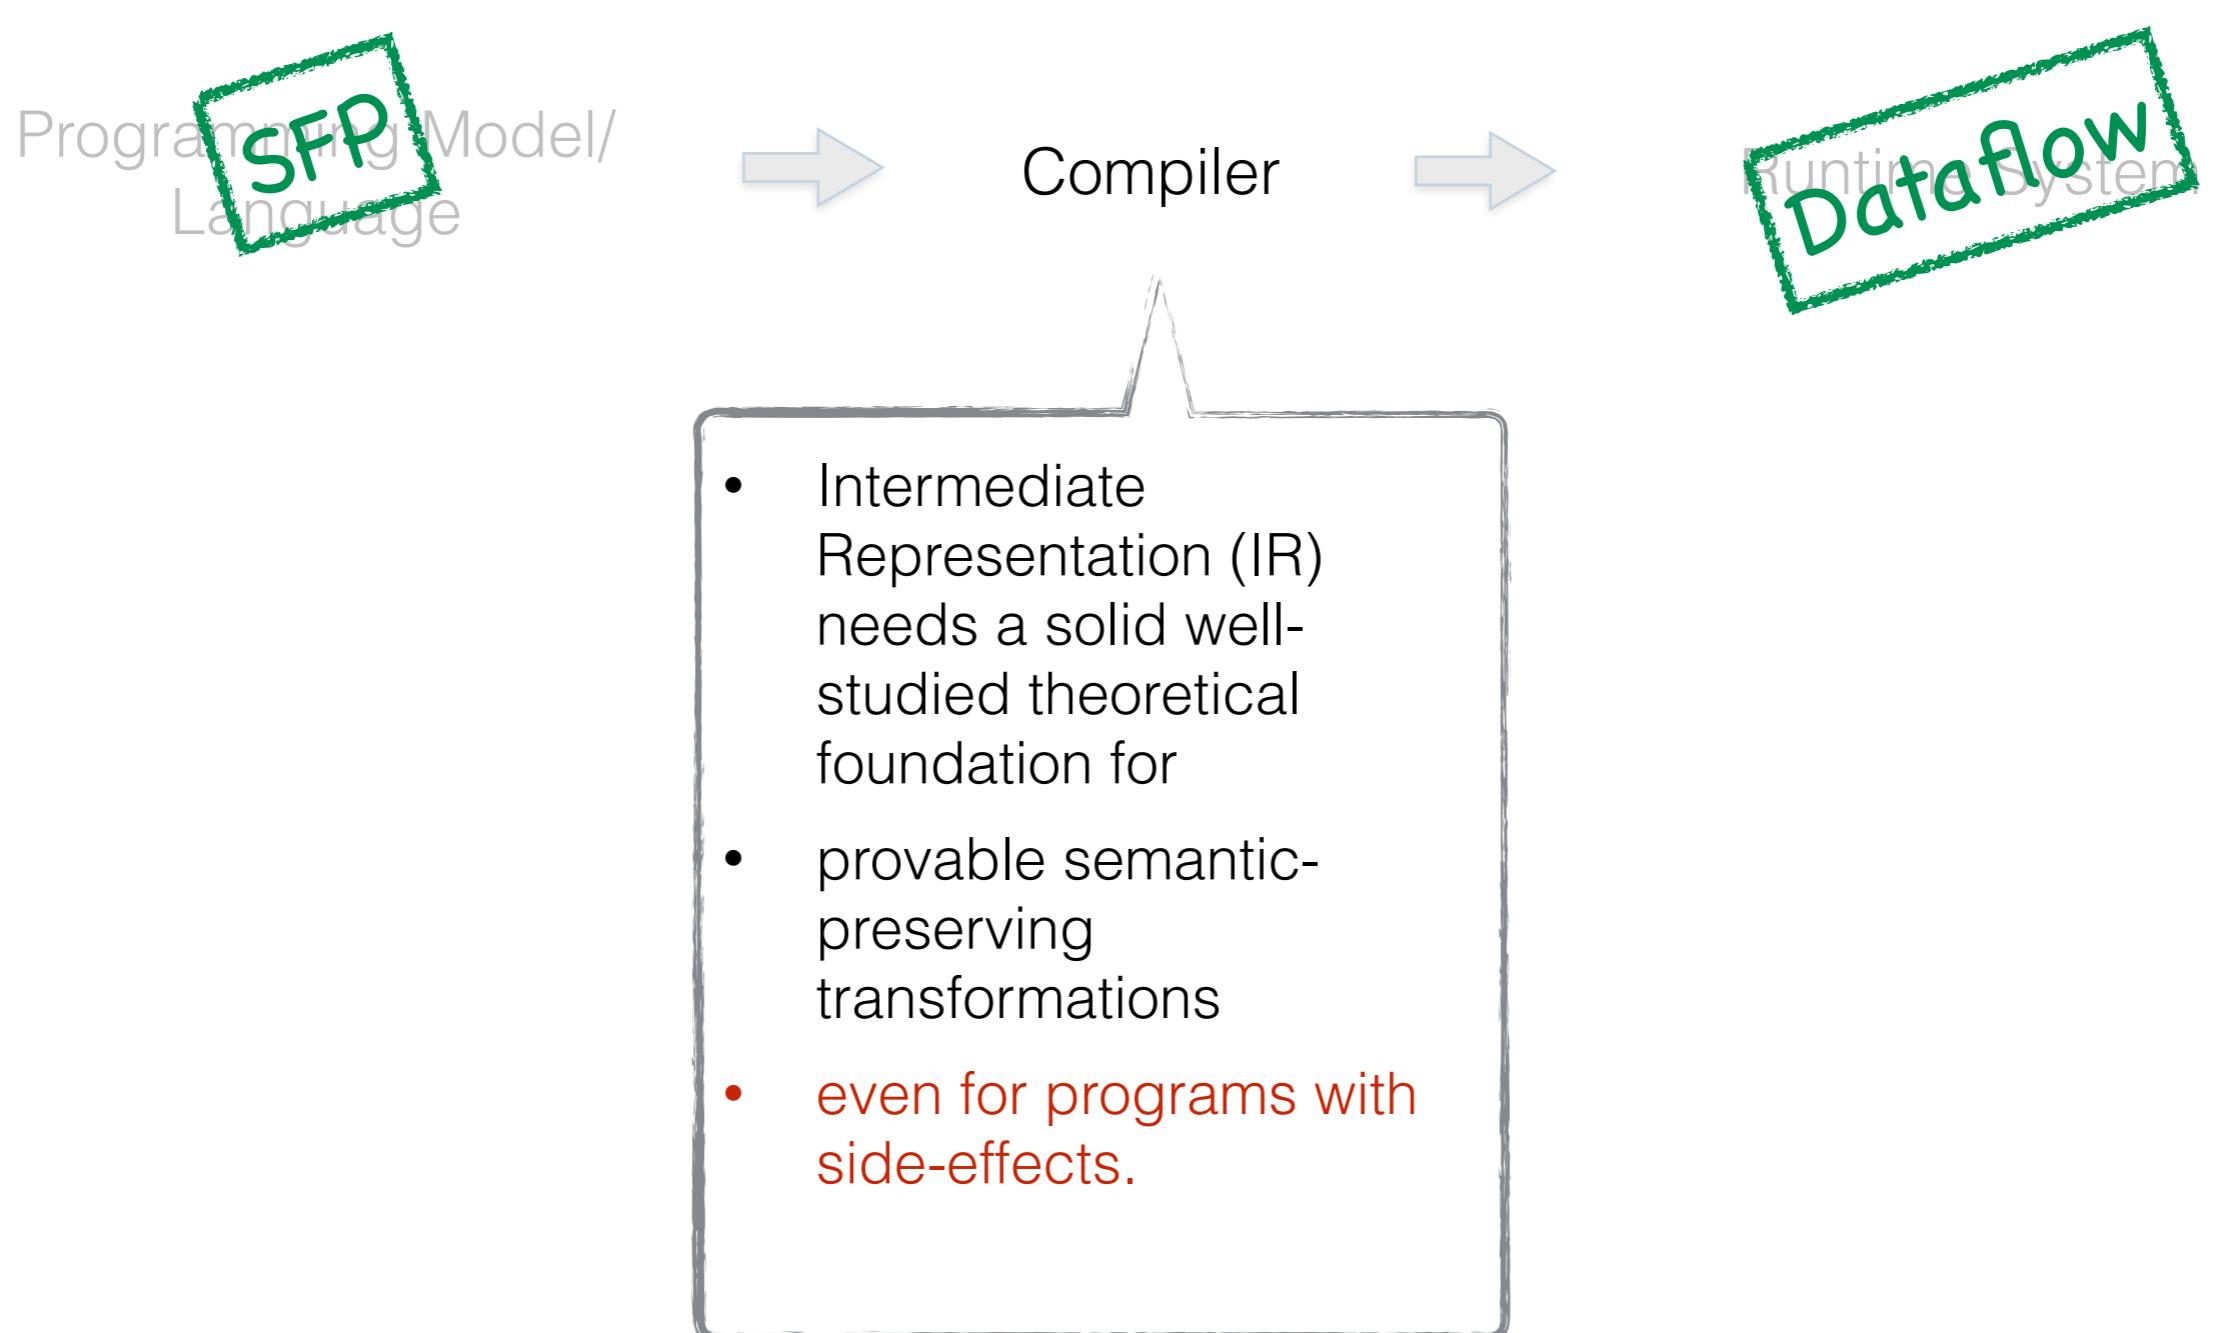


Compiler



Lambda calculus:  
- Variables  
- Abstraction  
- Application

Sebastian Ertel, Andrés Goens, Justus Adam, and Jeronimo Castrillon. 2018. Compiling for concise code and efficient I/O. In Proceedings of the 27th International Conference on Compiler Construction (CC 2018). ACM.



Programming Model/  
Language



Compiler



Dataflow  
Runtime System

SFP:

- Variables
- Stateful Functions
- Composition
- Algorithms
- Conditionals
- SMap

Lambda calculus:

- Variables
- Abstraction
- Application
- Combinators

# Stateful Functions



```
class FileLoad extends Filter {  
  
    private Channel<String> in, out;  
  
    public void work() {  
        // explicit channel control  
        String resource = in.take();  
        String contents = load(resource);  
        out.put(contents);  
    }  
  
    private Map<String, String> cache = new HashMap<>();  
  
    private String load(String resource){  
        String content = null;  
        if(!cache.containsKey(resource)){  
            content = new String(  
                Files.readAllBytes(Paths.get(resource))  
            );  
            cache.put(resource,content);  
            // cache eviction emitted for brevity  
        } else {  
            content = cache.get(resource);  
        }  
        return content;  
    }  
}
```

State encapsulation



```
class FileLoad {  
    Map<String, String> cache = new HashMap<>();  
  
    String load(String resource) {  
        String contents = null;  
        // load file data from disk or cache (omitted)  
        return contents;  
    }  
}  
  
struct FileLoad {  
    cache : HashMap,  
};  
  
impl FileLoad {  
    fn load(&self, resource:String) -> String {  
        let contents : String = {  
            // load file data from disk or cache (omitted)  
        };  
        contents  
    }  
}  
  
char* load(char* resource) {  
    static GHashTable* cache = g_hash_table_new();  
    char* contents = NULL;  
    // load file data from disk or cache (omitted)  
    return contents;  
}
```



# Foundations of Stateful Functions



# Foundations of Stateful Functions



State Threads:

```
load :: String -> State (HashMap String String)
          String
load :: String -> ST (MHashMap#(String, String))
```

# Foundations of Stateful Functions



State Threads:

```
load :: String -> State (HashMap String String)
          String
load :: String -> ST (MHashMap s String String)
          String
```

$$f_{st} = (a, s) \rightarrow (b, s)$$

Composition:

$$\begin{aligned} f &= a \rightarrow b \\ g &= b \rightarrow c \\ g \circ f &= a \rightarrow c \end{aligned}$$

$$f_{st} = (a, s_f) \rightarrow (b, s_f)$$

$$g_{st} = (a, s_g) \rightarrow (c, s_g)$$

$$g_{st} \circ f_{st} = ?$$

# Foundations of Stateful Functions



State Threads:

```
load :: String -> State (HashMap String String)
          String
load :: String -> ST (MHashMap#(String, String))
```

$$f_{st} = (a, s) \rightarrow (b, s)$$

Composition:

$$f = a \rightarrow b$$

$$g = b \rightarrow c$$

$$g \circ f = a \rightarrow c$$

$$f_{st} = (a, s_f) \rightarrow (b, s_f)$$

$$g_{st} = (a, s_g) \rightarrow (c, s_g)$$

$$g_{st} \circ f_{st} = ?$$

---

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. STCLang: state thread composition as a foundation for monadic dataflow parallelism. In Proceedings of the 12th ACM SIGPLAN International Symposium on Haskell (Haskell 2019). ACM.

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. Category-Theoretic Foundations of “STCLang: state thread composition as a foundation for monadic dataflow parallelism”. arXiv preprint arXiv:1906.12098

# Foundations of Stateful Functions



State Threads:

```
load :: String -> State (HashMap String String)
          String
load :: String -> ST (MHashMap#(String, String))
          String
```

$$f_{st} = (a, s) \rightarrow (b, s)$$

SMap:

- Formal analysis on determinism and deadlocks.
- Not only **maps** but also **folds**!
- Not only pipeline parallelism but also task-level and data parallelism!

Composition:

$$f = a \rightarrow b$$

$$g = b \rightarrow c$$

$$g \circ f = a \rightarrow c$$

$$f_{st} = (a, s_f) \rightarrow (b, s_f)$$

$$g_{st} = (a, s_g) \rightarrow (c, s_g)$$

$$g_{st} \circ f_{st} = ?$$

---

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. STCLang: state thread composition as a foundation for monadic dataflow parallelism. In Proceedings of the 12th ACM SIGPLAN International Symposium on Haskell (Haskell 2019). ACM.

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. Category-Theoretic Foundations of “STCLang: state thread composition as a foundation for monadic dataflow parallelism”. arXiv preprint arXiv:1906.12098

# Foundations of Stateful Functions



State Threads:

```
load :: String -> State (HashMap String String)
          String
load :: String -> ST (MHashMap#(String, String))
          String
```

$$f_{st} = (a, s) \rightarrow (b, s)$$

SMap:

- Formal analysis on determinism and deadlocks.
- Not only **maps** but also **folds**!
- Not only pipeline parallelism but also task-level and data parallelism!

Composition:

$$f = a \rightarrow b$$

$$g = b \rightarrow c$$

$$g \circ f = a \rightarrow c$$

$$f_{st} = (a, s_f) \rightarrow (b, s_f)$$

$$g_{st} = (a, s_g) \rightarrow (c, s_g)$$

$$g_{st} \circ f_{st} = ?$$

```
fn map_reduce() {
  let stateG = init_G();
  for data in list {
    let x = f(data);
    stateG.g(x)
  };
  stateG
}
```

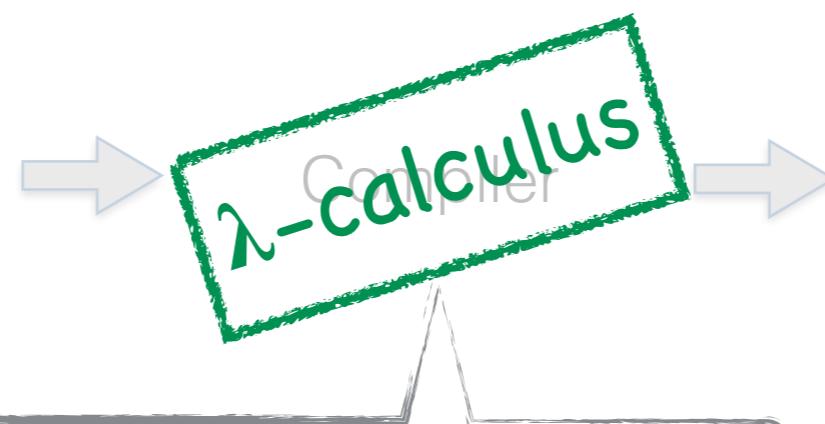
---

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. STCLang: state thread composition as a foundation for monadic dataflow parallelism. In Proceedings of the 12th ACM SIGPLAN International Symposium on Haskell (Haskell 2019). ACM.

Sebastian Ertel, Justus Adam, Norman A. Rink, Andrés Goens, and Jeronimo Castrillon. 2019. Category-Theoretic Foundations of “STCLang: state thread composition as a foundation for monadic dataflow parallelism”. arXiv preprint arXiv:1906.12098

Programming Model/  
Language

SFP

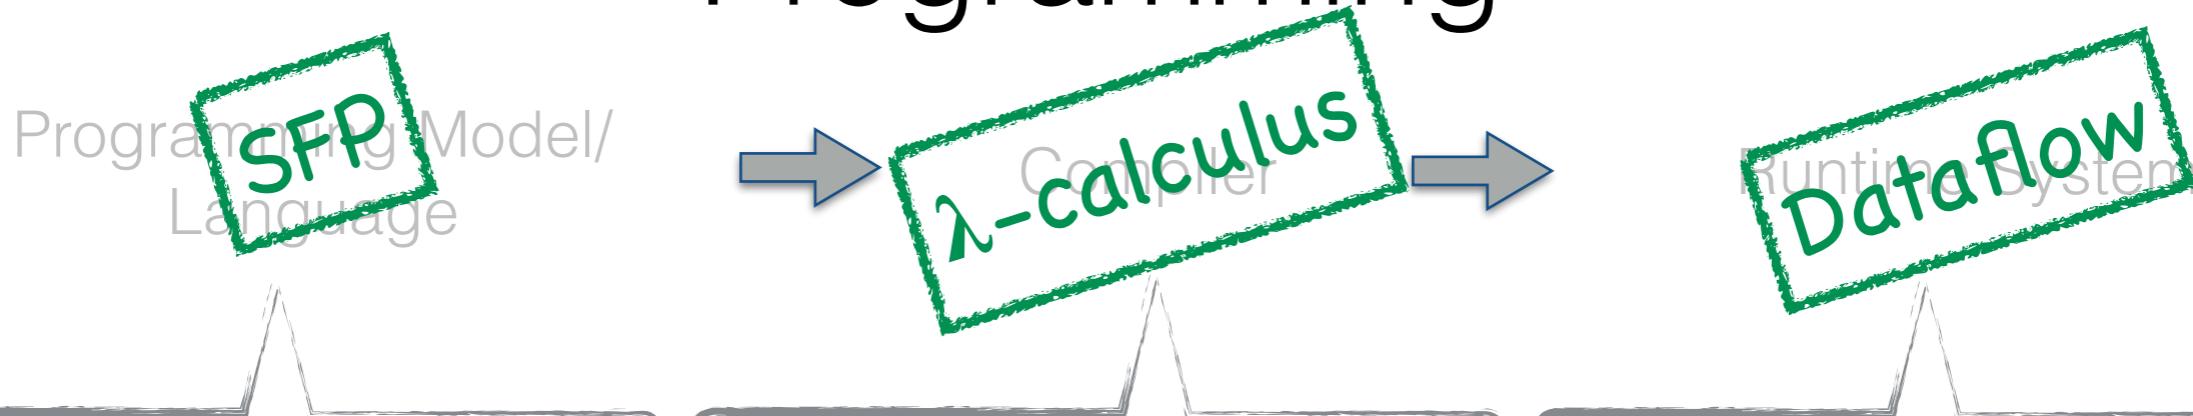


Runtime System

Dataflow

- Intermediate Representation (IR) needs a solid well-studied theoretical foundation for
- provable semantic-preserving transformations
- even for programs with side-effects.

# Ingredients for Implicit Parallel Programming

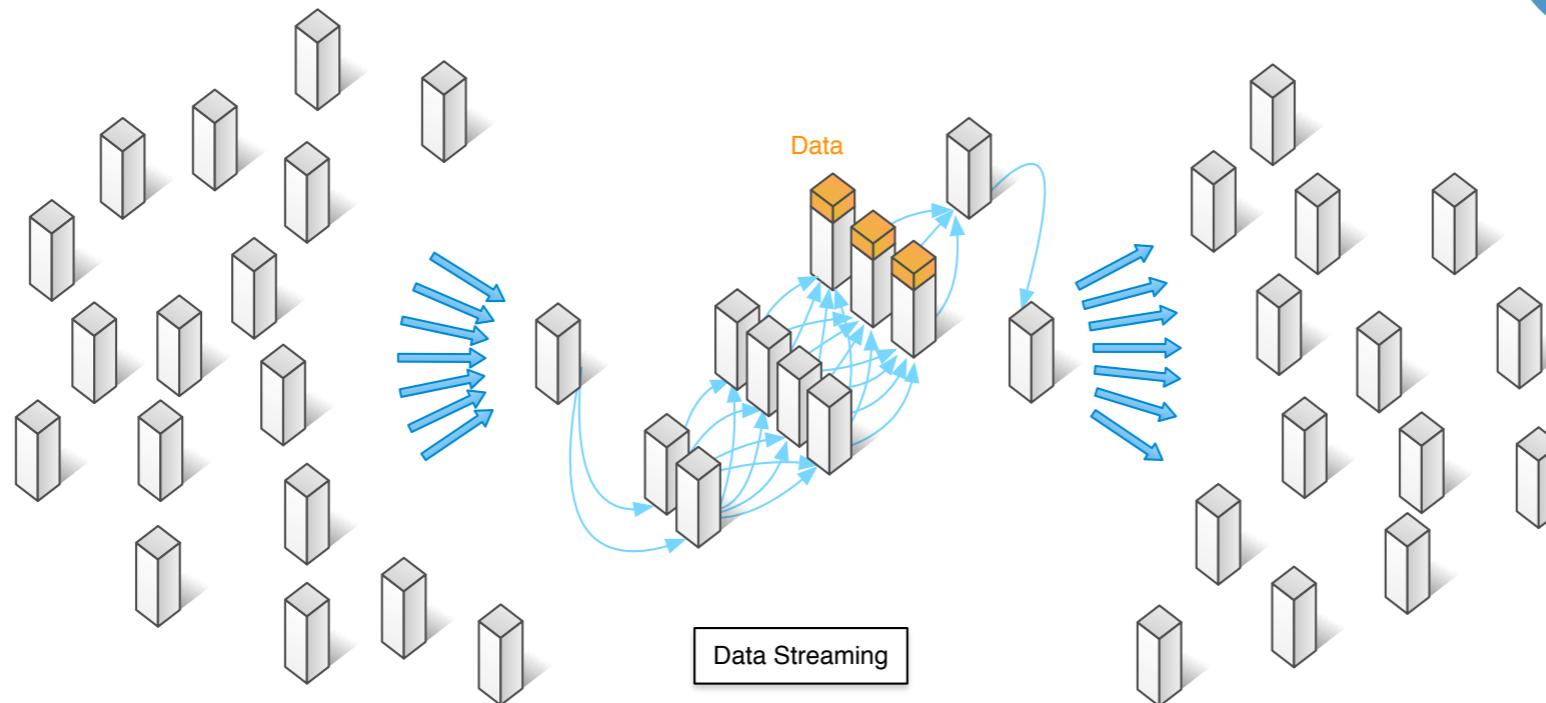


- No concurrency abstractions.
- Practicality:
  - Integration into existing programming models (e.g., OOP)/languages.
  - Gradual switch/reuse of existing code base.
- Intermediate Representation (IR) needs a solid well-studied theoretical foundation for
  - provable semantic-preserving transformations
  - even for programs with side-effects.
- Flexible/dynamic runtime representation.
- Efficient parallel execution of independent computations.
- Deterministic execution.
- Concurrent I/O.

# Changing the Game

- Functional Reactive Programming (FRP) based on stateful functions.
- FRP instead of streams for big data systems.

<https://github.com/ohua-dev/stc-lang>



- Current Main Focus: Rust integration ➡ Safety  
<https://github.com/ohua-dev/ohua-rust-runtime>
- STM alternative.
- User-defined functions and NoSQL (in Noria).



Thanks for your attention!



<https://ohua-dev.github.io/>