

Compiling for Concise Code and Efficient I/O

Sebastian Ertel, Andrés Goens, Justus Adam and Jeronimo Castrillon

Chair for Compiler Construction
TU Dresden

27th International Conference on Compiler Construction
Vienna, 25.2.2018

- Backend (server) system requirements: scalable, flexible, fault-tolerant ...

- Backend (server) system requirements: scalable, flexible, fault-tolerant ...



[<http://magnasoma.com/#/monolith-3>]

- Monolithic server program

- Backend (server) system requirements: scalable, flexible, fault-tolerant ...



[<http://magnasoma.com/#/monolith-3>]


- Monolithic server program
- Microservice: independent function
- Microservices compose to larger services

- Backend (server) system requirements: scalable, flexible, fault-tolerant ...



[<http://magnasoma.com/#/monolith-3>]

- Monolithic server program



amazon

<http://amazon.com>

[Source: I Love APIs 2015 by Chris Munns, licensed under CC BY 4.0, available at: [h p://bit.ly/2zboHTK](https://bit.ly/2zboHTK)]

- Microservice: independent function
- Microservices compose to larger services

- Backend (server) system requirements: scalable, flexible, fault-tolerant ...



[<http://magnasoma.com/#/monolith-3>]

- Monolithic server program




- Microservice: independent function
- Microservices compose to larger services

Efficient I/O

Concise Code

Efficient I/O

- Batching (& deduplication):
1 I/O call vs. n I/O calls




- Concurrency:
Computation—IO—Computation
- (Caching)

Concise Code

Efficient I/O

- Batching (& deduplication):
1 I/O call vs. n I/O calls



group I/O calls

Concise Code


- Breaks modularity!



- Concurrency:
Computation—IO—Computation
- (Caching)

Efficient I/O

- Batching (& deduplication):
1 I/O call vs. n I/O calls



→
group I/O calls

- Concurrency:
Computation—IO—Computation
- (Caching)

→
split up code
(e.g. into threads)

Concise Code

- Breaks modularity!




- Introduces optimization
aspect (concurrency)
into the algorithm!



Efficient I/O

- Batching (& deduplication):
1 I/O call vs. n I/O calls



- Concurrency:
Computation—IO—Computation
- (Caching)



group I/O calls

Concise Code

- Breaks modularity! ⚡

- Batching and concurrency
do not easily compose! ⚡

split up code
(e.g. into threads)



- Introduces optimization
aspect (concurrency)
into the algorithm! ⚡

Efficient I/O

- Batching (& deduplication):
1 I/O call vs. n I/O calls



Release the developer from this burden → Compiler-based Approach!

group I/O calls

Concise Code

- Breaks modularity!




split up code
(e.g. into threads)

- Concurrency:
Computation—IO—Computation
- (Caching)

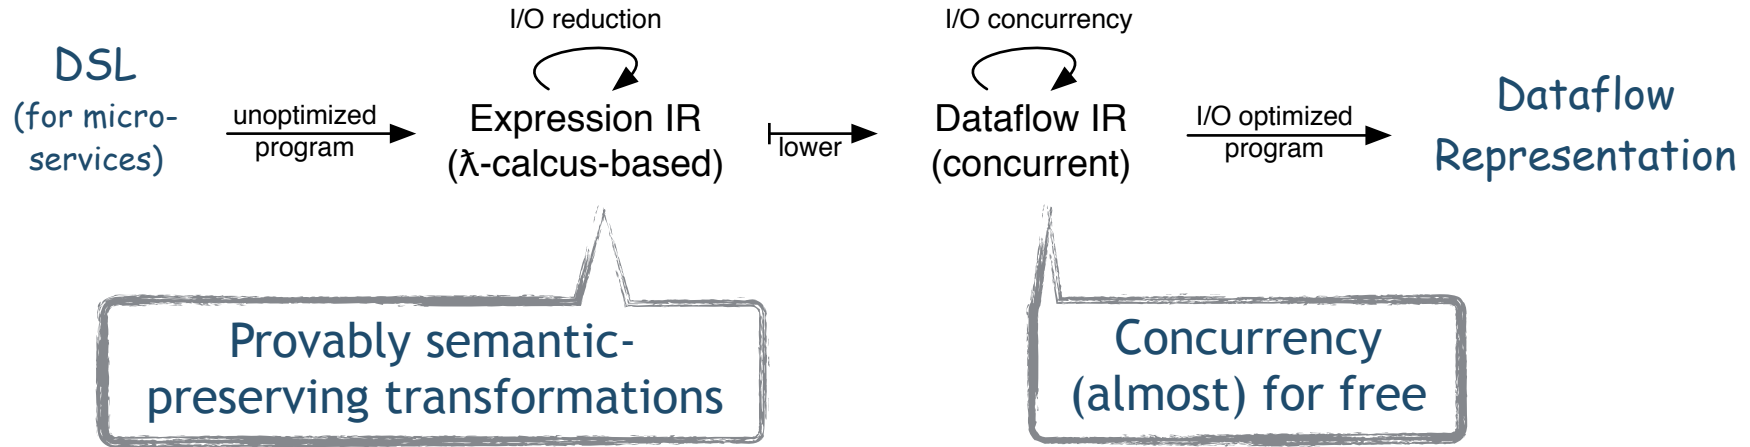
- Introduces optimization
aspect (concurrency)
into the algorithm!




Yauhau - A compiler framework for microservices



Yauhau - A compiler framework for microservices




A minimalistic DSL for microservices



```
† ::= v
| λv.†
| † †
| let v=† in †
| if († † †)
| fff(x1 ... xn)
| io(x)

map(λv.† [v1 ... vn])
```

Semantic-preserving Transformations




$\dagger ::= v$

**Disclaimer: Presentation diverges from paper for visualization purposes.
Data flow graphs before and after the transformations instead of
lambda expressions.**

$| \text{io}(x)$

$\text{map}(\lambda v. \dagger [v_1 \dots v_n])$


Semantic-preserving Transformations




```
t ::= v
    |  $\lambda v. t$ 
    | t t
    | let v = t in t
    | if (t t t)
    | fff(x1 ... xn)
    | io(x)
```

```
map( $\lambda v. t$  [v1 ... vn])
```

- Grouping independent I/O calls:




Semantic-preserving Transformations




```
† ::= v
| λv.t
| † †
| let v=† in †
| if († † †)
| ff_f(x1 ... xn)
| io(x)

map(λv.t [v1 ... vn])
```

- Grouping independent I/O calls:






```

t ::= v
  |  $\lambda v.t$ 
  | t t
  | let v=t in t
  | if (t t t)
  | fff(x1 ... xn)
  | io(x)
    
```

$\text{map}(\lambda v.t [v_1 \dots v_n])$


- Grouping independent I/O calls:



- Lifting I/O calls:



Semantic-preserving Transformations




```


t ::= v
  |  $\lambda v.t$ 
  | t t
  | let v=t in t
  | if (t t t)
  | fff(x1 ... xn)
  | io(x)
    
```

$\text{map}(\lambda v.t [v_1 \dots v_n])$


- Grouping independent I/O calls:



- Lifting I/O calls:



Semantic-preserving Transformations




```


t ::= v
  |  $\lambda v. t$ 
  | t t
  | let v=t in t
  | if (t t t)
  |  $ff_f(x_1 \dots x_n)$ 
  | io(x)
    
```

$map(\lambda v. t \ [v_1 \dots v_n])$


- Grouping independent I/O calls:



- Lifting I/O calls:



Semantic-preserving Transformations




```

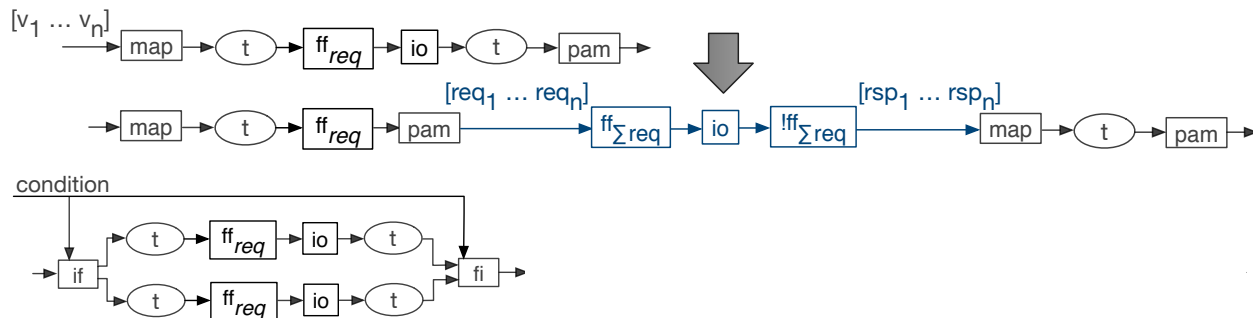
t ::= v
    |  $\lambda v. t$ 
    | t t
    | let v = t in t
    | if (t t t)
    | fff(x1 ... xn)
    | io(x)
    
```

map($\lambda v. t$ [v₁ ... v_n])


- Grouping independent I/O calls:



- Lifting I/O calls:



Semantic-preserving Transformations




```

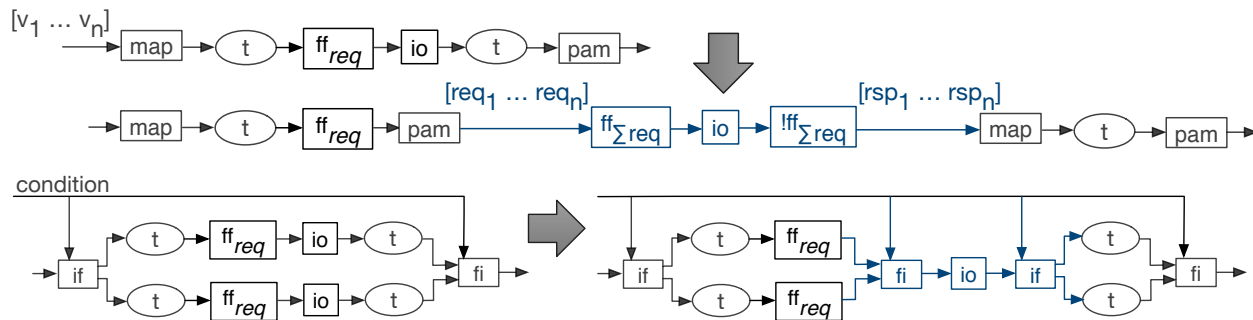
t ::= v
    |  $\lambda v. t$ 
    | t t
    | let v = t in t
    | if (t t t)
    | fff(x1 ... xn)
    | io(x)
    
```

map($\lambda v. t$ [v₁ ... v_n])


- Grouping independent I/O calls:



- Lifting I/O calls:



Bringing back Concurrency




- Computation—IO




- IO—IO

- IO—Computation

Bringing back Concurrency




- Computation—IO




- IO—IO

- IO—Computation


Bringing back Concurrency



- Computation—IO




- IO—IO




- IO—Computation


Bringing back Concurrency




- Computation—IO 



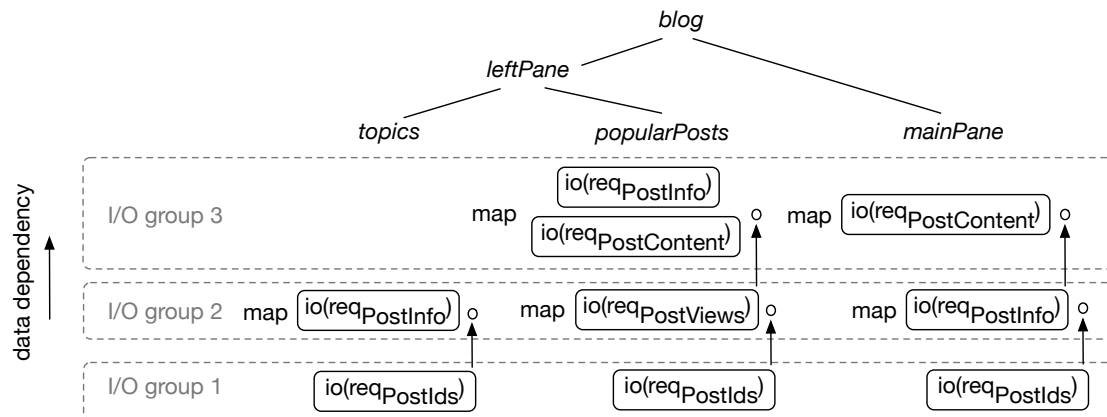
- IO—IO 



- IO—Computation 



Use Case: Blog



Version	seq	base	batch	full
Time [ms]	275 ± 25	292 ± 19	79 ± 1	66.5 ± 5.2

4x latency improvement!


- Haxl → Facebook's Haskell-based approach on the abstraction of applicative functors.
- Muse → Similar to Twitter's Stitch, Clojure-based implementation of an AST interpretation.

Simon Marlow, Louis Brandy, Jonathan Coens, and Jon Purdy. 2014. There is no fork: an abstraction for efficient, concurrent, and concise data access. ICFP '14.

Alexey Kachayev. 2015. Reinventing Haxl: Efficient, Concurrent and Concise Data Access. Presentation at EuroClojure 2015.

Got “microservices for evaluation”?


→ Level-Graphs:




Andrés Goens, Sebastian Ertel, Justus Adam, and Jeronimo Castrillon. 2018. Level Graphs: Generating Benchmarks for Concurrency Optimizations in Compilers. MULTIPROG’18.

Got “microservices for evaluation”?

→ Level-Graphs:




→ Programs:





Andrés Goens, Sebastian Ertel, Justus Adam, and Jeronimo Castrillon. 2018. Level Graphs: Generating Benchmarks for Concurrency Optimizations in Compilers. MULTIPROG’18.

Yauhau outperforms State-of-the-art

→ Level-Graphs:




→ Programs:



Andrés Goens, Sebastian Ertel, Justus Adam, and Jeronimo Castrillon. 2018. Level Graphs: Generating Benchmarks for Concurrency Optimizations in Compilers. MULTIPROG'18.

Yauhau outperforms State-of-the-art




→ Level-Graphs:





→ Programs:



serialize



Andrés Goens, Sebastian Ertel, Justus Adam, and Jeronimo Castrillon. 2018. Level Graphs: Generating Benchmarks for Concurrency Optimizations in Compilers. MULTIPROG'18.

- I/O optimizations performed at the compiler level
 - Efficient I/O 
 - Concise Code 
- Use lambda calculus to provide semantic-preserving transformations!
- Use dataflow for concurrency/parallel execution!

Try it out:

<https://tud-ccc.github.io/yauhau-doc/>

```
(defn ohua-batching []  
  (<-ohua  
    (time-end (blog (time-begin three)))  
    :compile-with-config {:df-transformations yauhau.ir-transform/transformations}))  
  
  (defalgo blog [pid]  
    (let [pop1 (popular-posts pid)  
          top (topics pid)  
          spane (render-side-pane pop1 top)  
          mpane (main-pane pid)]  
      (render-page spane mpane)))  
  
  (defalgo topics [pid]  
    (let [ids (fetch-ids pid)  
          topics2 (smap (algo [pid] (let [info (fetch-info pid)] (:topic info))) ids)  
          concatenated (concat topics2)  
          freqs (frequencies concatenated)]  
      (render-topics freqs)))
```

```
do
  -- level 3
  (local4, local5, local6, local7) <- (,,,)
  <$> getData "source" [100] <*> compute [100]
  <*> getData "source" [100] <*> compute [100]
  -- level 2
  local3 <- subfun3 [100, local4, local5, local7]
  -- level 1
  (local1, local2) <- (,)
  <$> compute [100, local4]
  <*> getData "source" [100, local3]
  -- level 0
  getData "source" [100, local1, local2, local6]
```

Haxl (< GHC 8)

```
(cats/mlet
  [; level 3
  [local-4 local-5 local-6 local-7]
  (cats/<$>
    clojure.core/vector
    (get-data "source" 100)
    (cats/<$> (compute (cats/return 100)))
    (get-data "source" 100)
    (cats/<$> (compute (cats/return 100))))
  ; level 2
  local-3
  (subfun-3 local-4 local-5 local-7)
  ; level 1
  [local-1 local-2]
  (cats/<$>
    clojure.core/vector
    (cats/<$>
      (compute (cats/return 100)
        (cats/return local-4)))
    (get-data "source" 100 local-3))]
  ; level 0
  (get-data "source" 100 local-1 local-2
    local-6)))
```

Muse (with Cats)

```
(let
  [; level 3
  [local-4 local-5 local-6 local-7]
  (vector (get-data "source" 100)
    (compute 100)
    (get-data "source" 100)
    (compute 100))]
  ; level 2
  local-3 (subfun-3 local-4 local-5 local-7)
  ; level 1
  [local-1 local-2]
  (vector (compute 100 local-4)
    (get-data "source" 100 local-3))]
  ; level 0
  (get-data "source" 100 local-1 local-2
    local-6)))
```

Yauhau